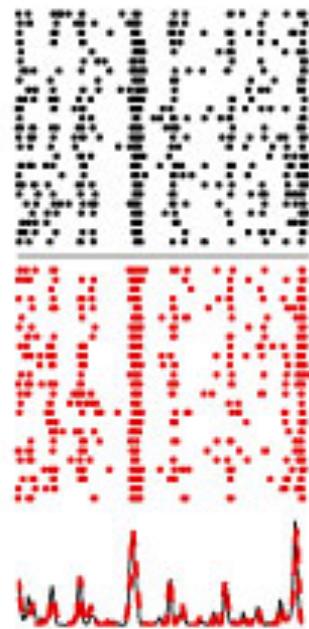


Statistical models for neural coding: GLMs & beyond



Jonathan Pillow

Princeton Neuroscience Institute, Psychology,
Center for Statistics & Machine Learning
Princeton University

CSHL summer course 2024

scientific influences



Rich Zemel

University of
Arizona

(Now Columbia)



Eero Simoncelli

NYU



Peter Dayan



Peter Latham

Gatsby Computational
Neuroscience Unit

scientific influences



Alex Huk

UT Austin



Mala Murthy

Princeton University

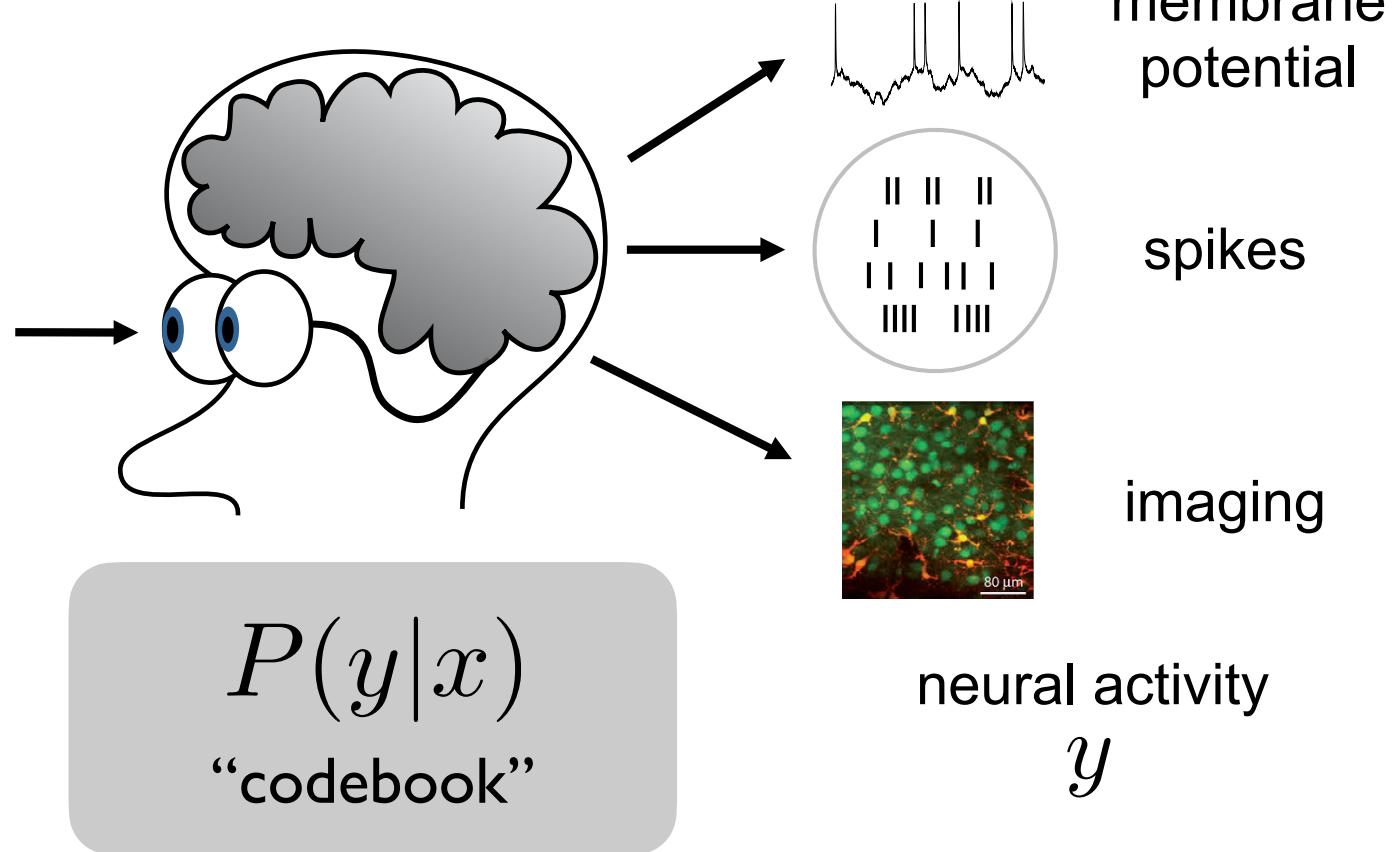


Ilana Witten

neural coding problem



stimulus
 x

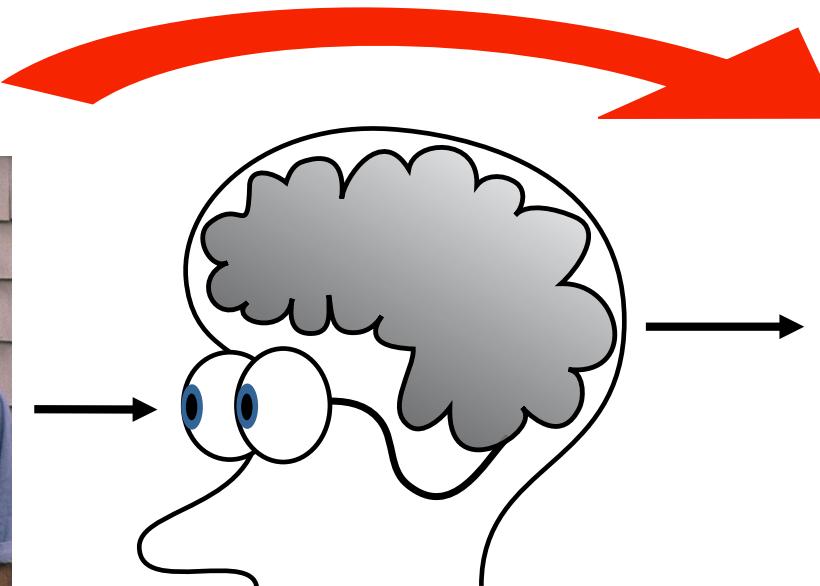


Q: What is the probabilistic relationship between stimuli and neural activity?

encoding



x



$$P(y|x)$$

“codebook”

neural activity
 y

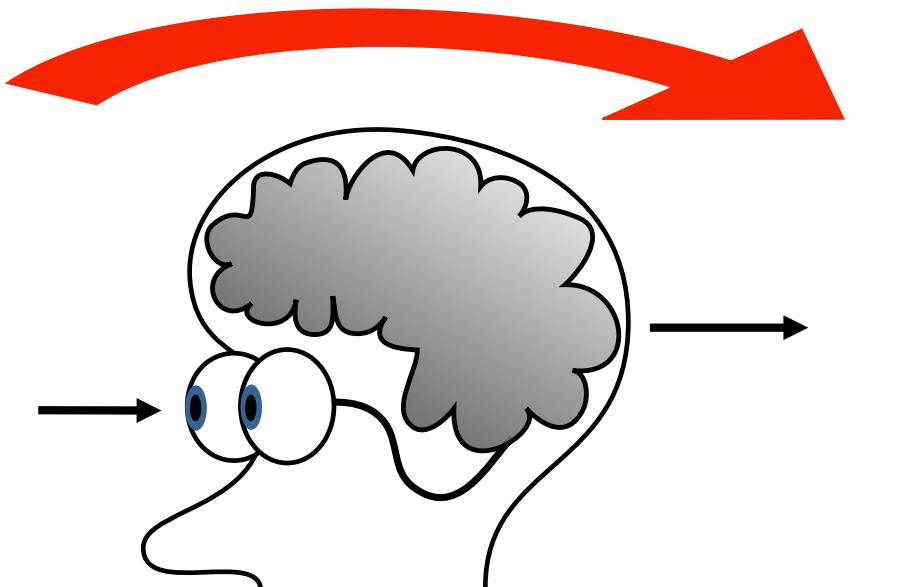
uses:

- insight into sensory processing
- sensory prosthetics

encoding



x



$$P(y|x)$$

“codebook”

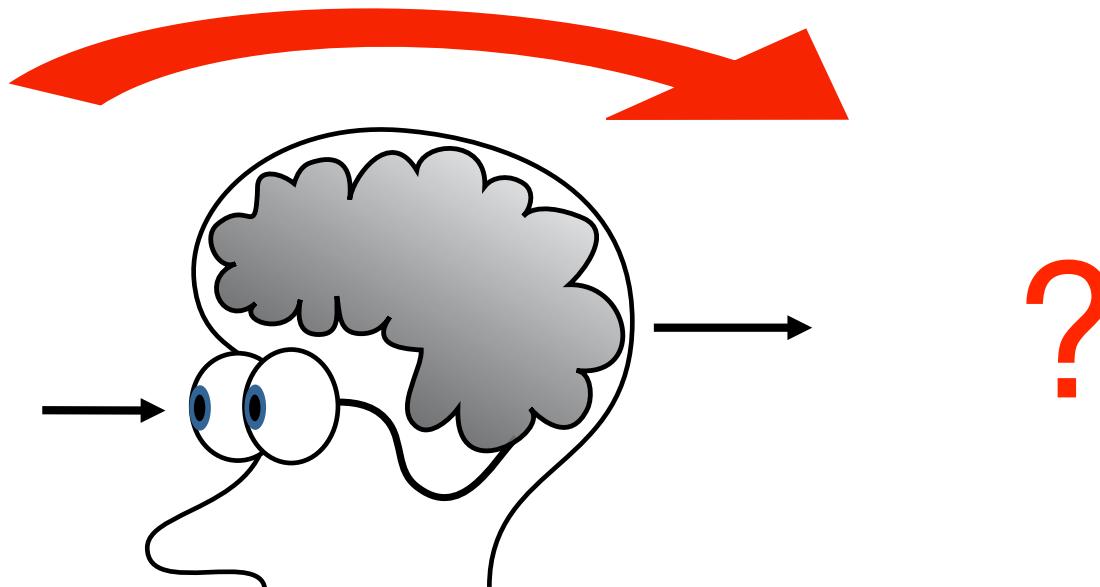
neural activity
 y

- uses:
- insight into sensory processing
 - sensory prosthetics

encoding



x



$$P(y|x)$$

“codebook”

neural activity
 y

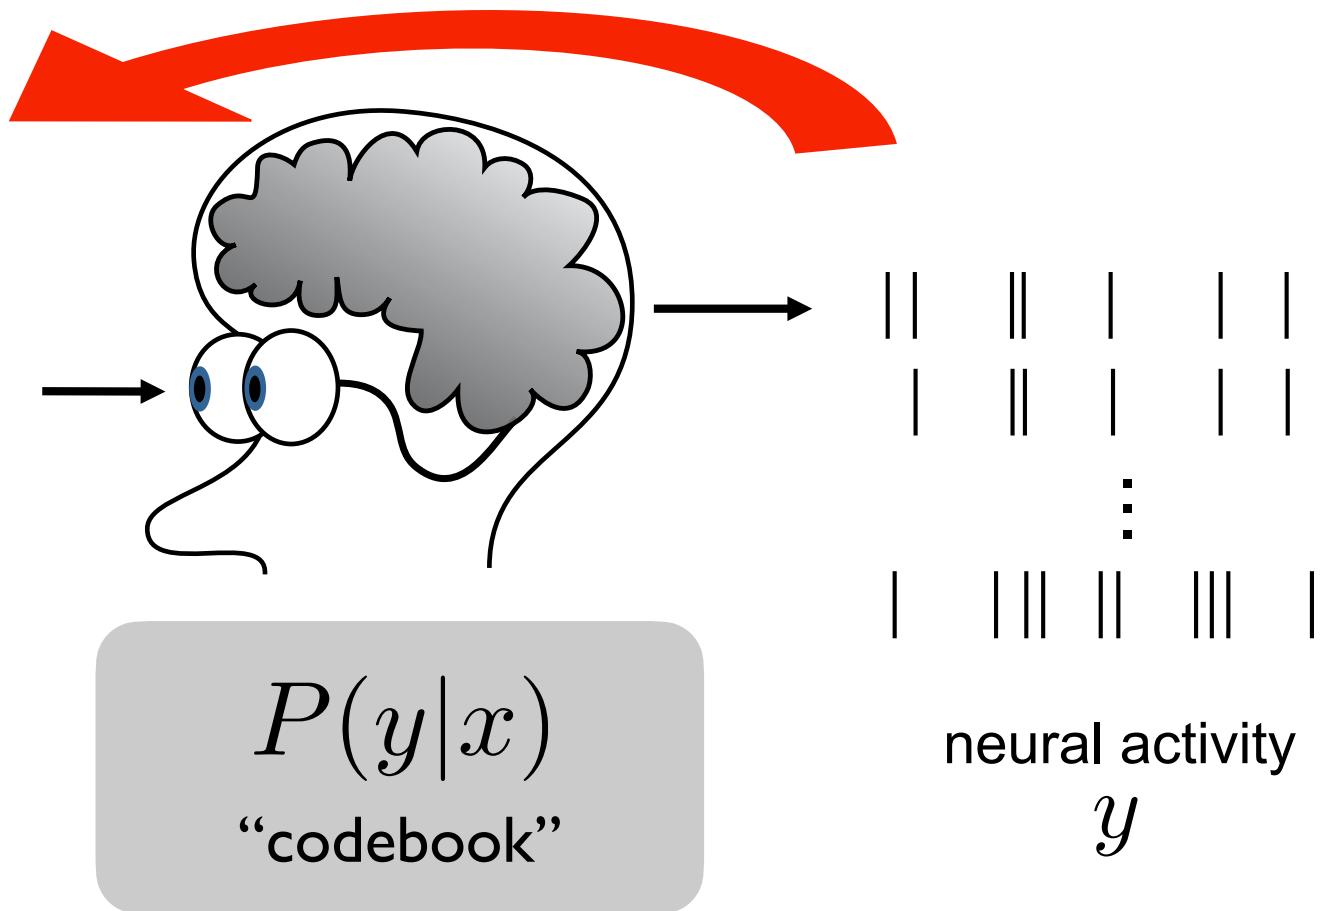
uses:

- insight into sensory processing
- sensory prosthetics

decoding

?

x



Bayes' Rule:

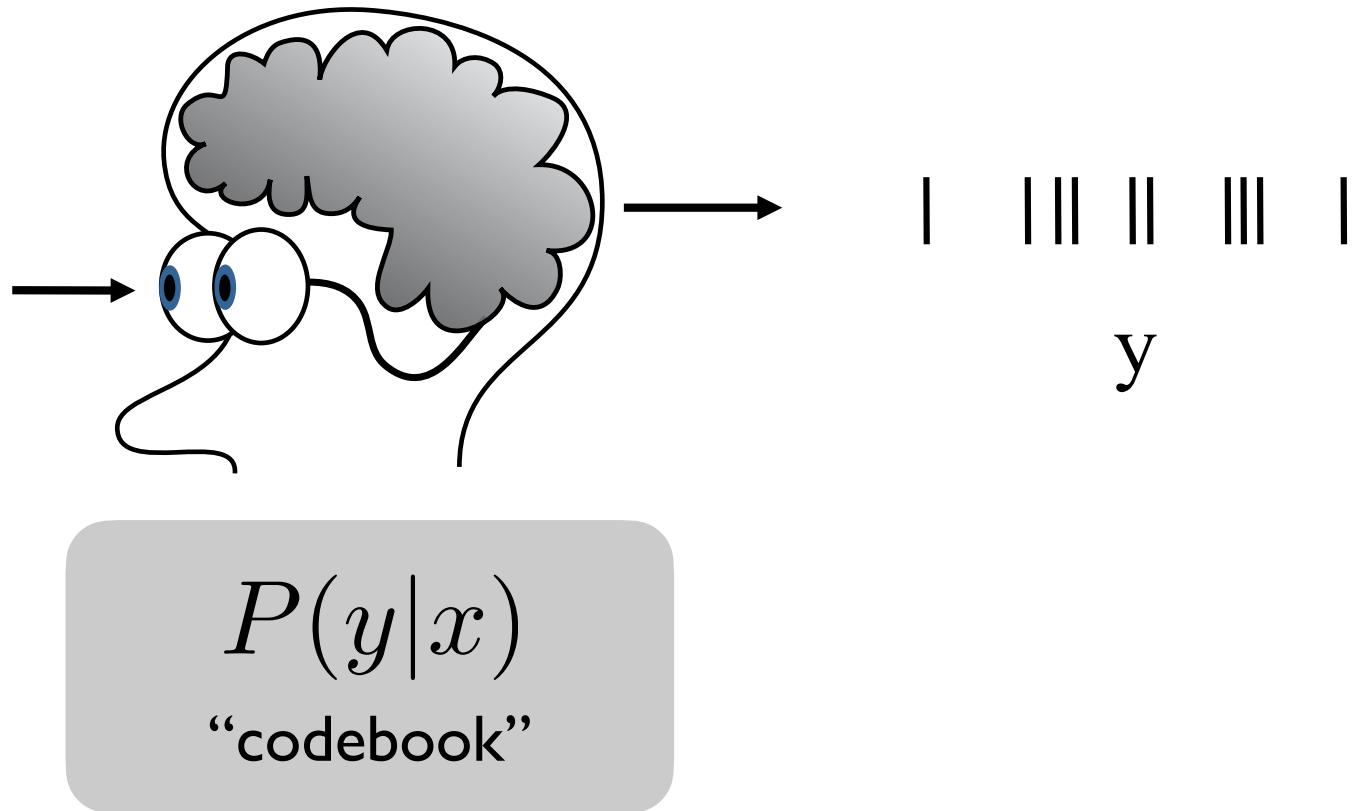
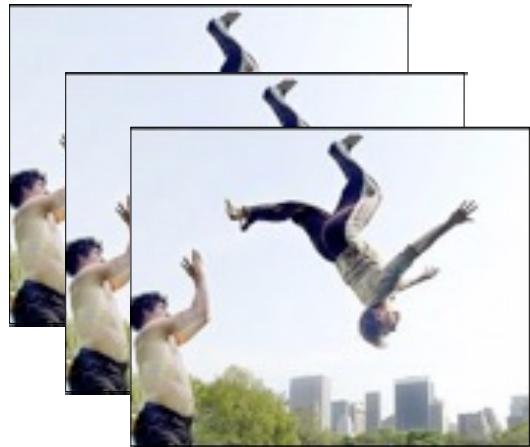
$$P(x|y) \propto P(y|x)P(x)$$

posterior likelihood prior

uses:

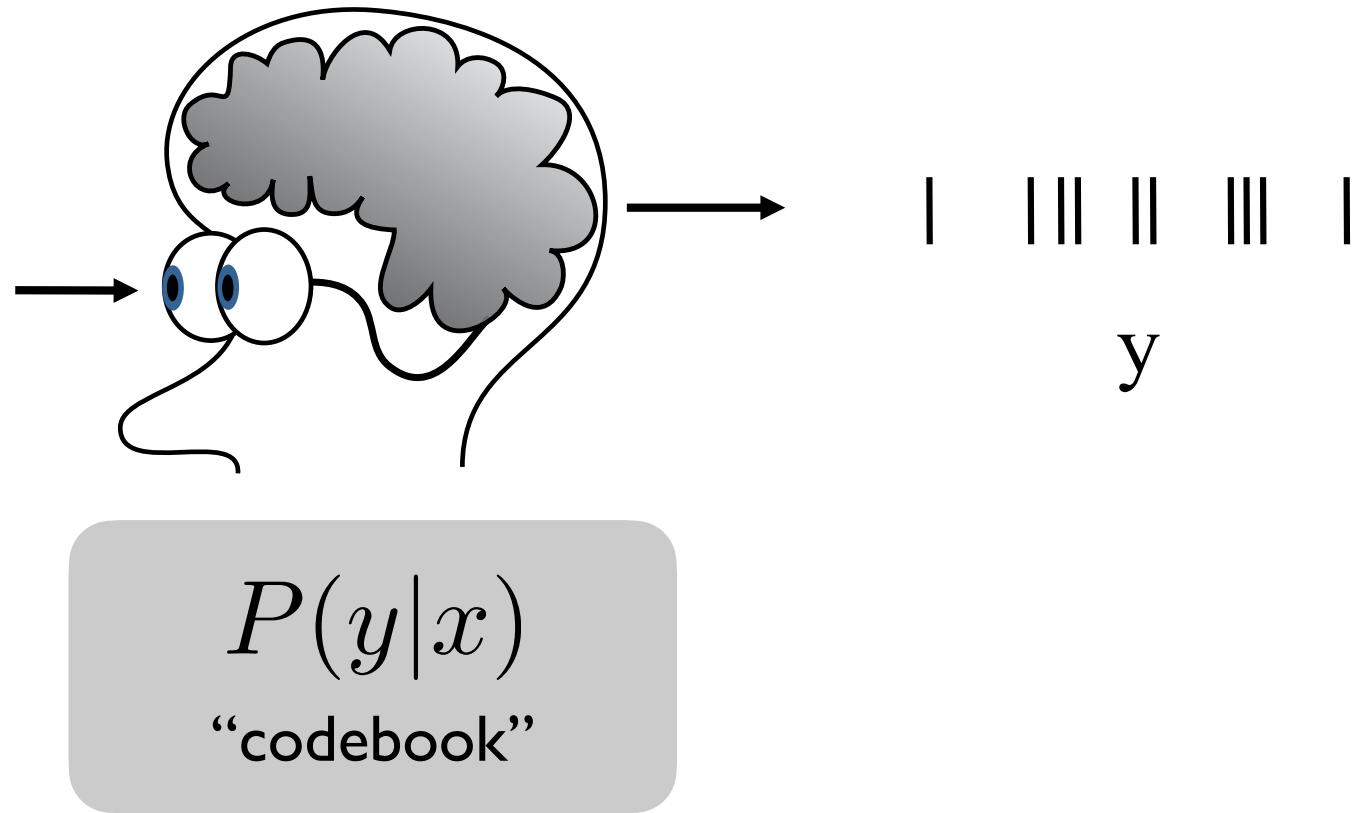
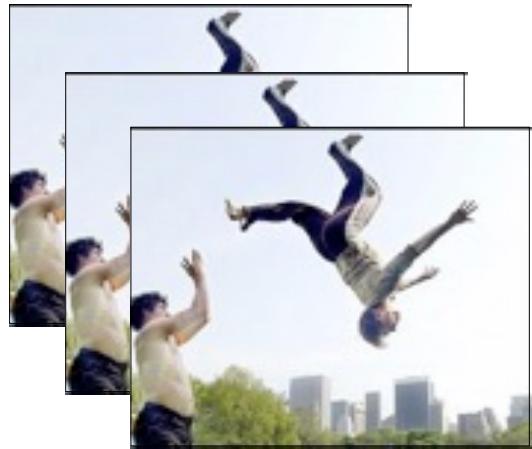
- constraints on neural readout of pop codes
- motor prosthetics

general name: “regression models”



- compare to “latent variable models”, which focus only on $P(y)$ (e.g., dimensionality reduction, clustering, state space models)
- Lea to present these tomorrow!

model-based approach

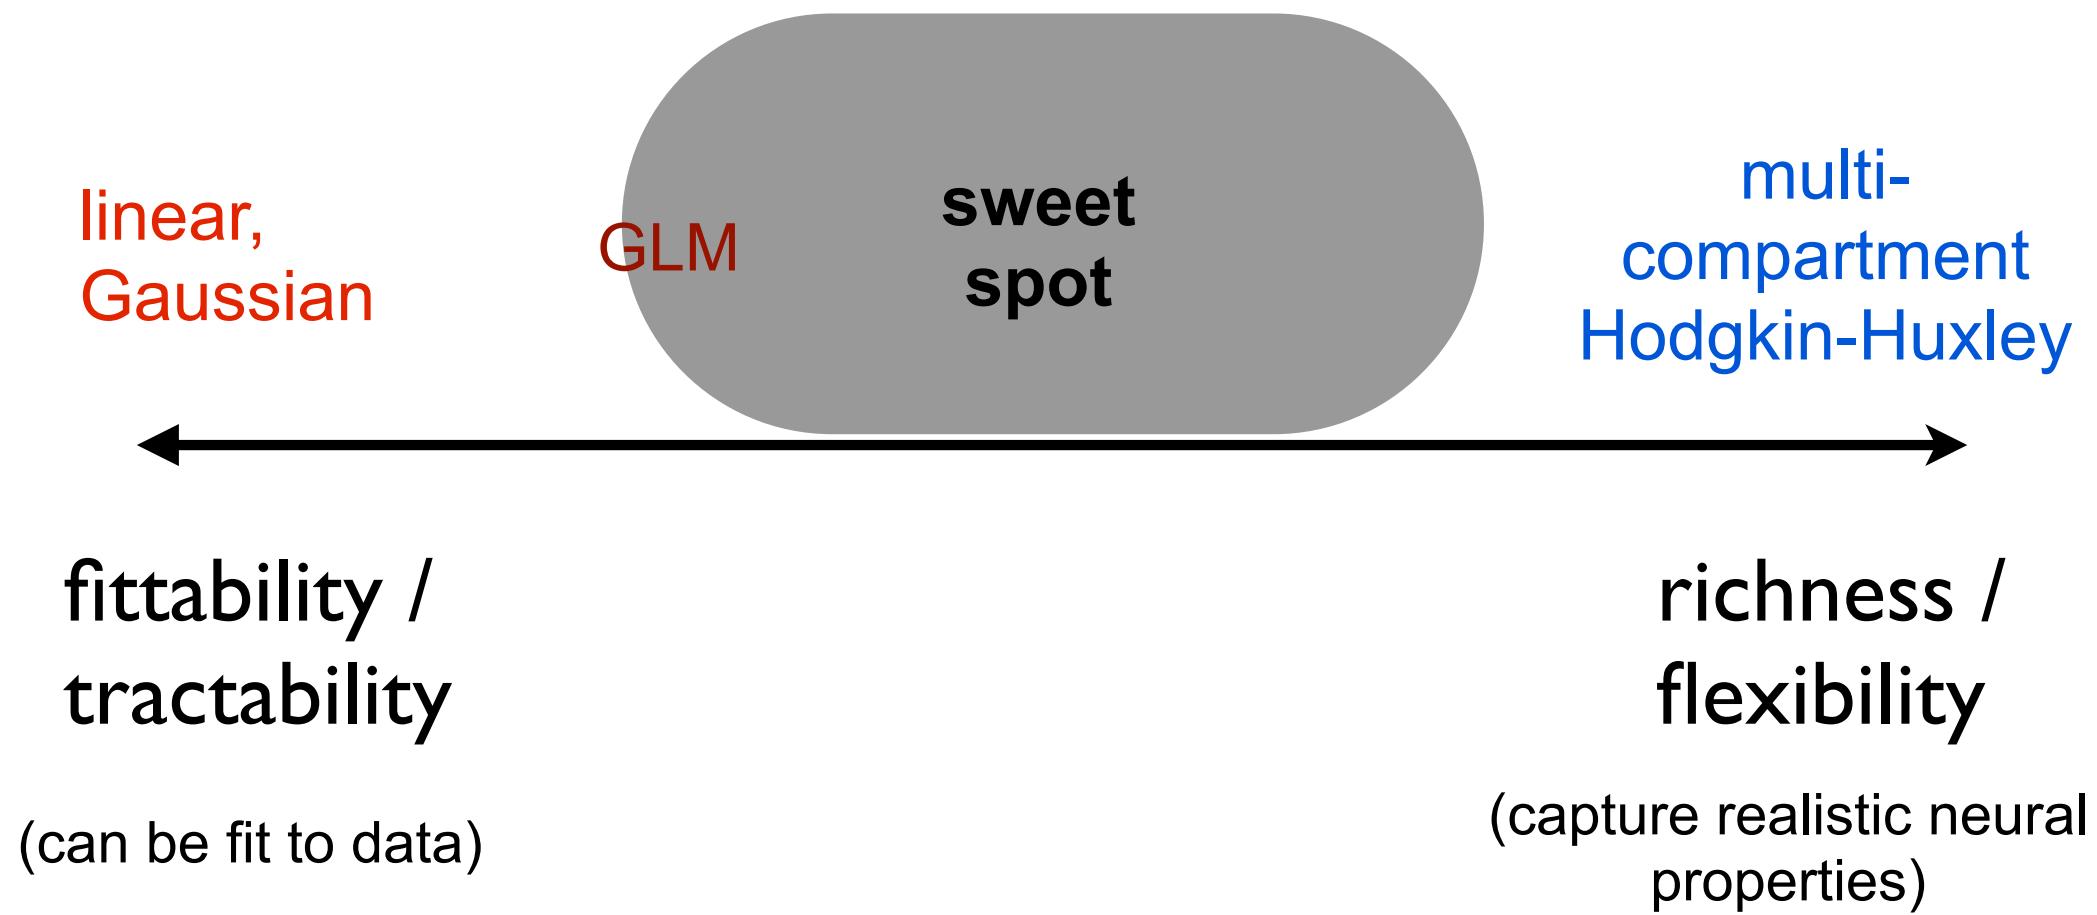


$$\text{model: } P_{\theta}(y|x) \approx P(y|x)$$

Goal: find model that closely approximates true encoding distribution

Question: what criteria for picking a model?

model desiderata



Outline

I. maximum likelihood for simple models

- spike-count codes

$$P(y|x, \theta)$$

2. Multi-neuron models

- Generalize linear model (GLM)

$$P(y_1, y_2, \dots, y_n | x, \theta)$$

Example 1: linear Poisson neuron

spike count

$$y \sim Poiss(\lambda)$$

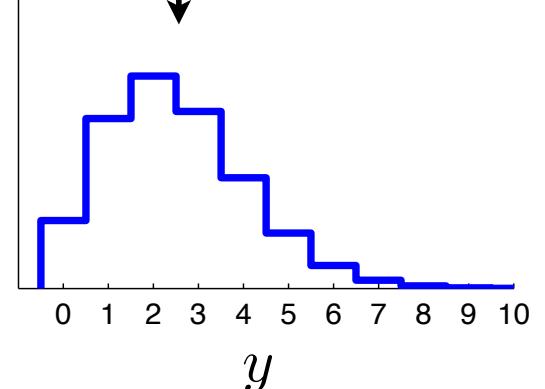
spike rate

$$\lambda = \theta x$$

parameter

stimulus

λ = mean = variance

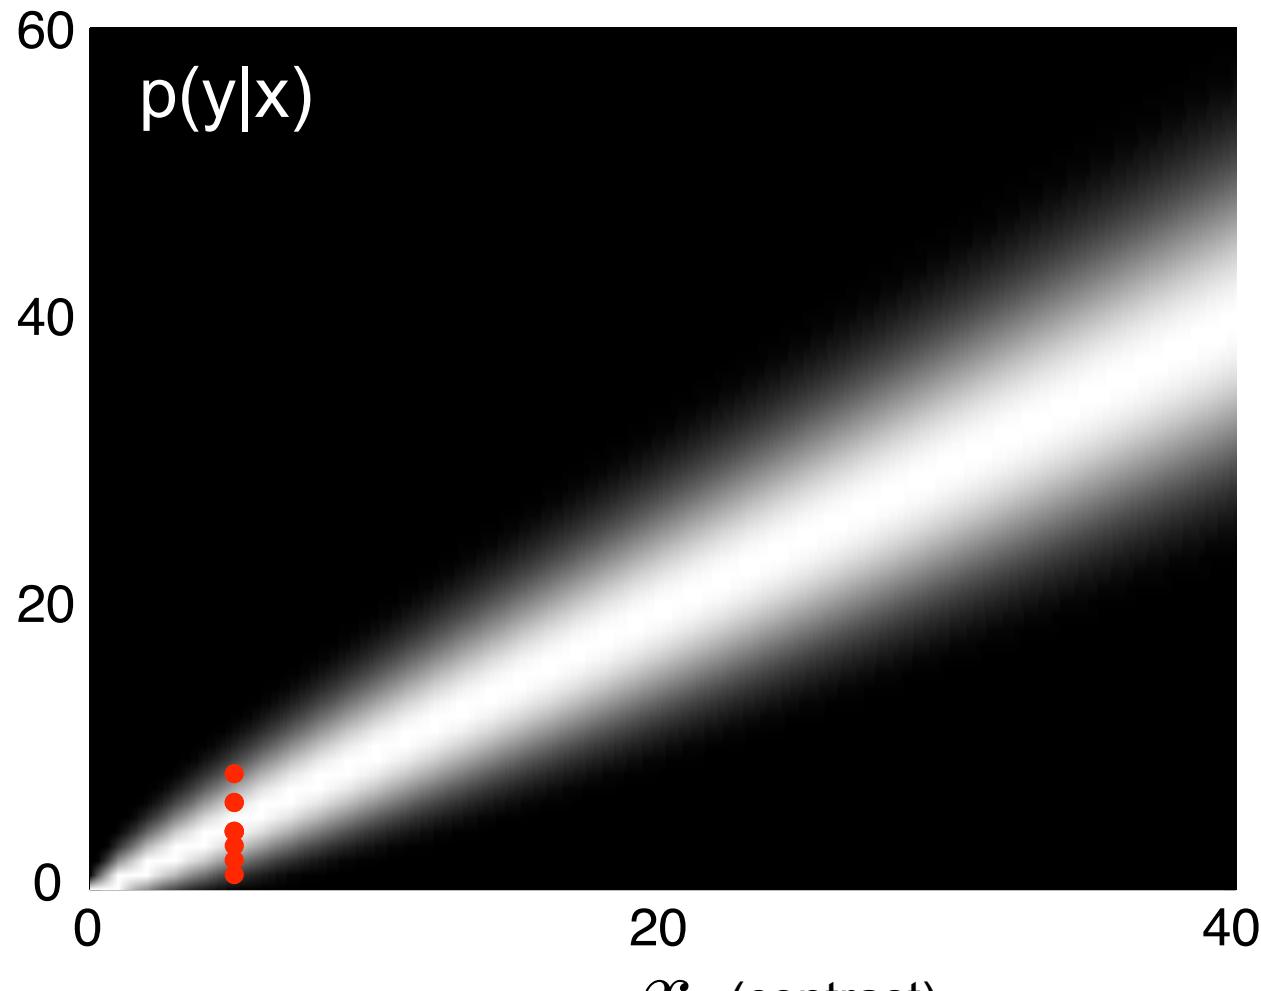


encoding model:

$$\begin{aligned} P(y|x, \theta) &= \frac{1}{y!} \lambda^y e^{-\lambda} \\ &= \frac{1}{y!} (\theta x)^y e^{-(\theta x)} \end{aligned}$$

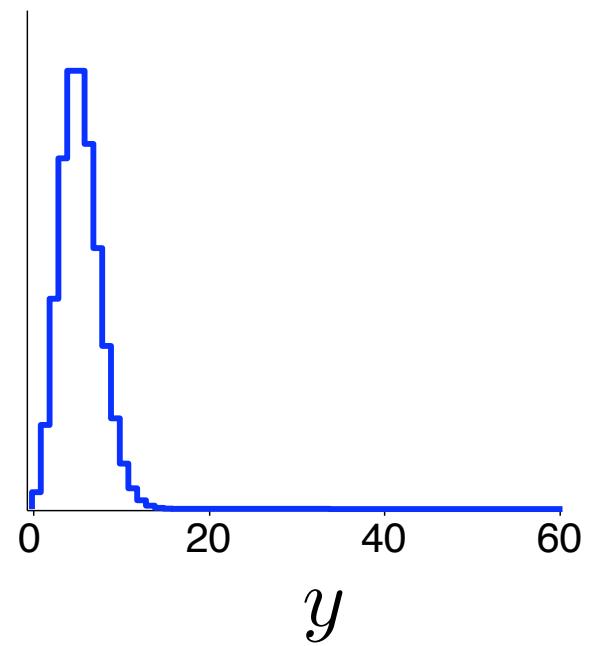
$$\text{mean}(y) = \theta x$$

$$\text{var}(y) = \theta x$$



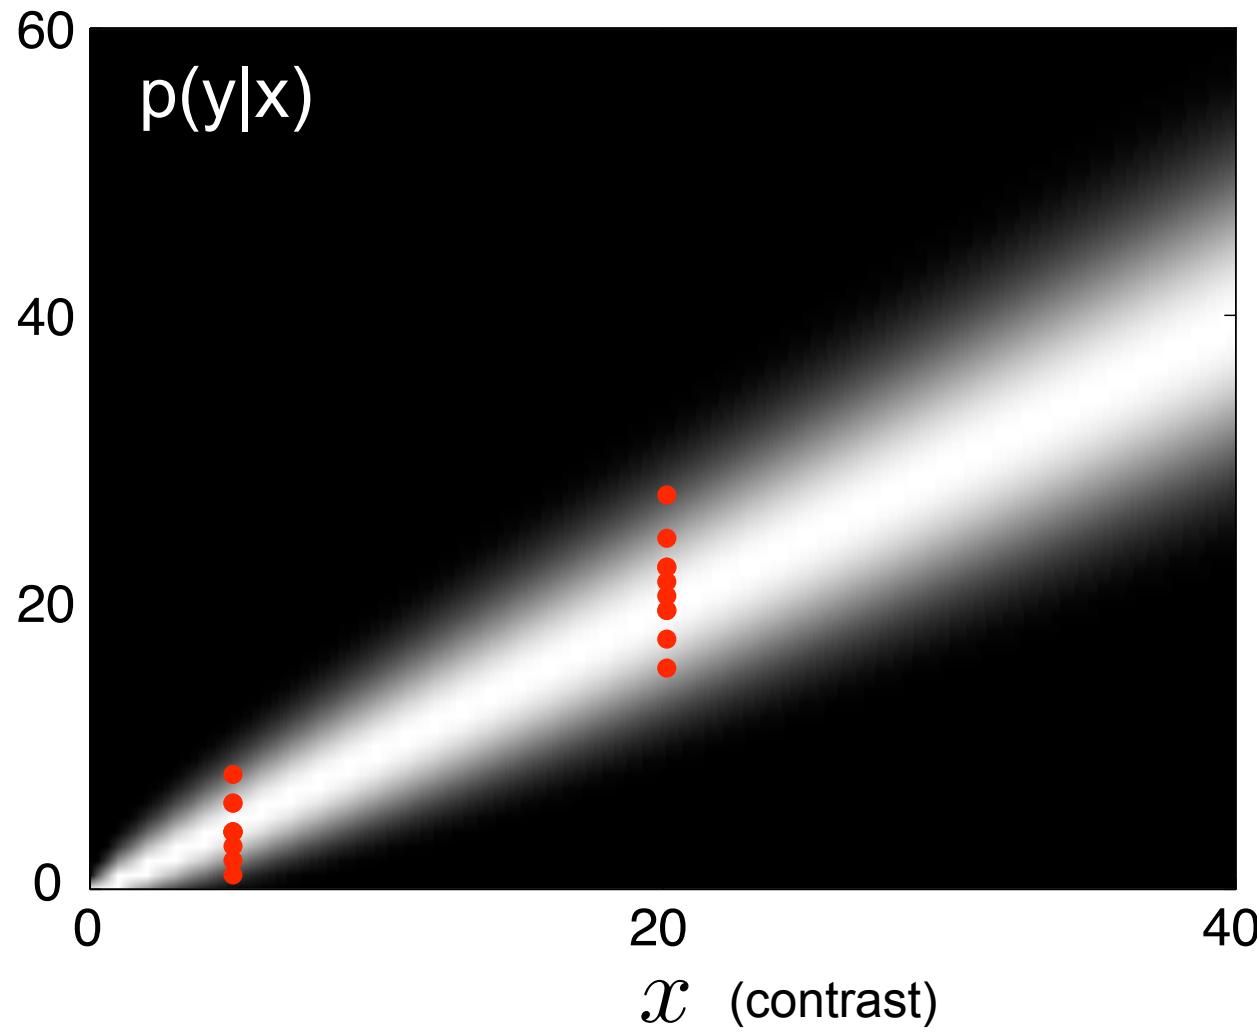
conditional distribution

$$p(y|x = 5)$$



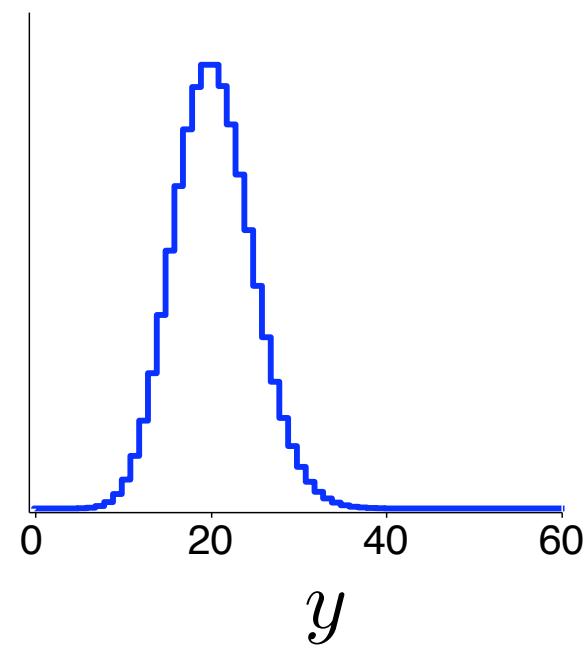
$$\text{mean}(y) = \theta x$$

$$\text{var}(y) = \theta x$$



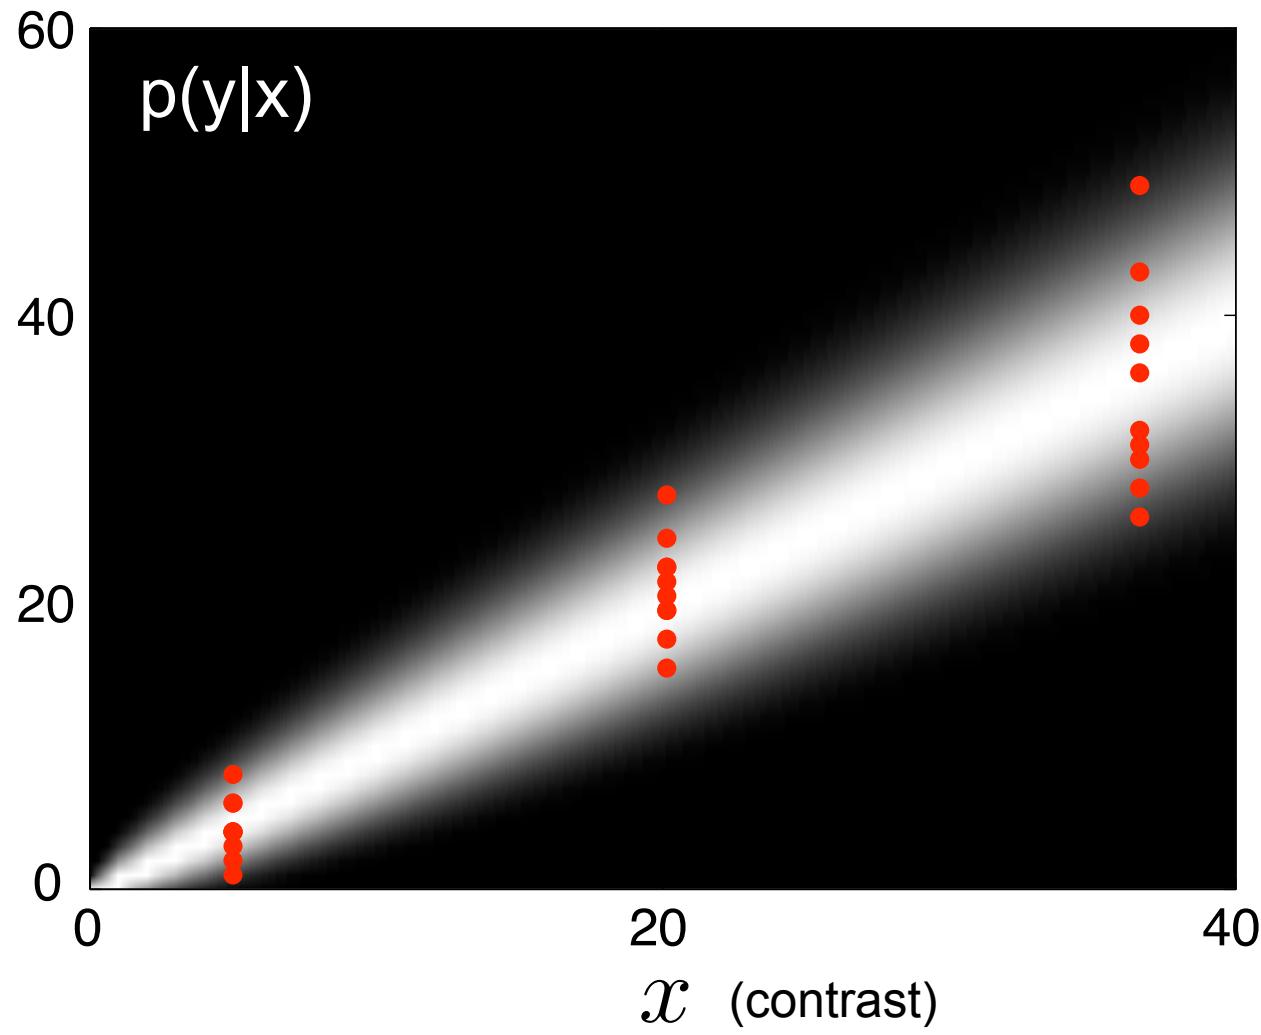
conditional distribution

$$p(y|x = 20)$$



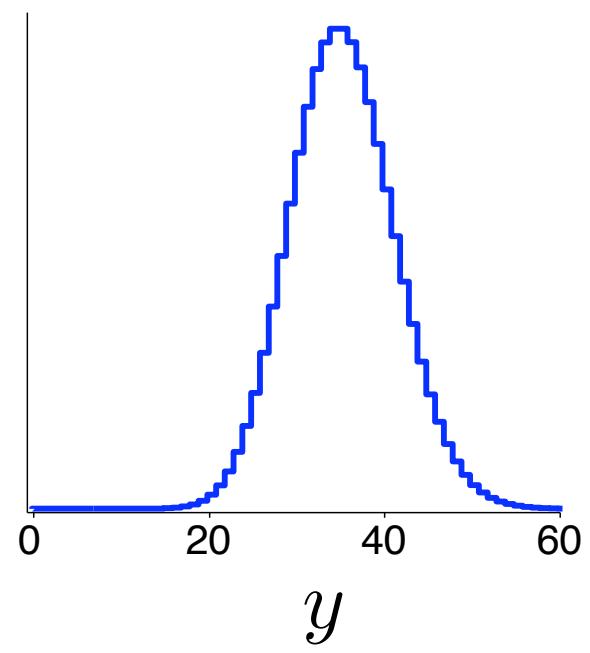
$$\text{mean}(y) = \theta x$$

$$\text{var}(y) = \theta x$$



conditional distribution

$$p(y|x = 35)$$



Maximum Likelihood Estimation:

- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$

↙ ↓ ↓
all spike all parameters
counts stimuli

$$P(Y|X, \theta) = \prod_{i=1}^N \underbrace{P(y_i|x_i, \theta)}_{\text{single-trial probability}}$$

Q: what assumption are we making about the responses?

A: conditional independence across trials!

Maximum Likelihood Estimation:

- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$


all spike counts all stimuli parameters

$$P(Y|X, \theta) = \prod_{i=1}^N \underbrace{P(y_i|x_i, \theta)}_{\text{single-trial probability}}$$

Q: what assumption are we making about the responses?

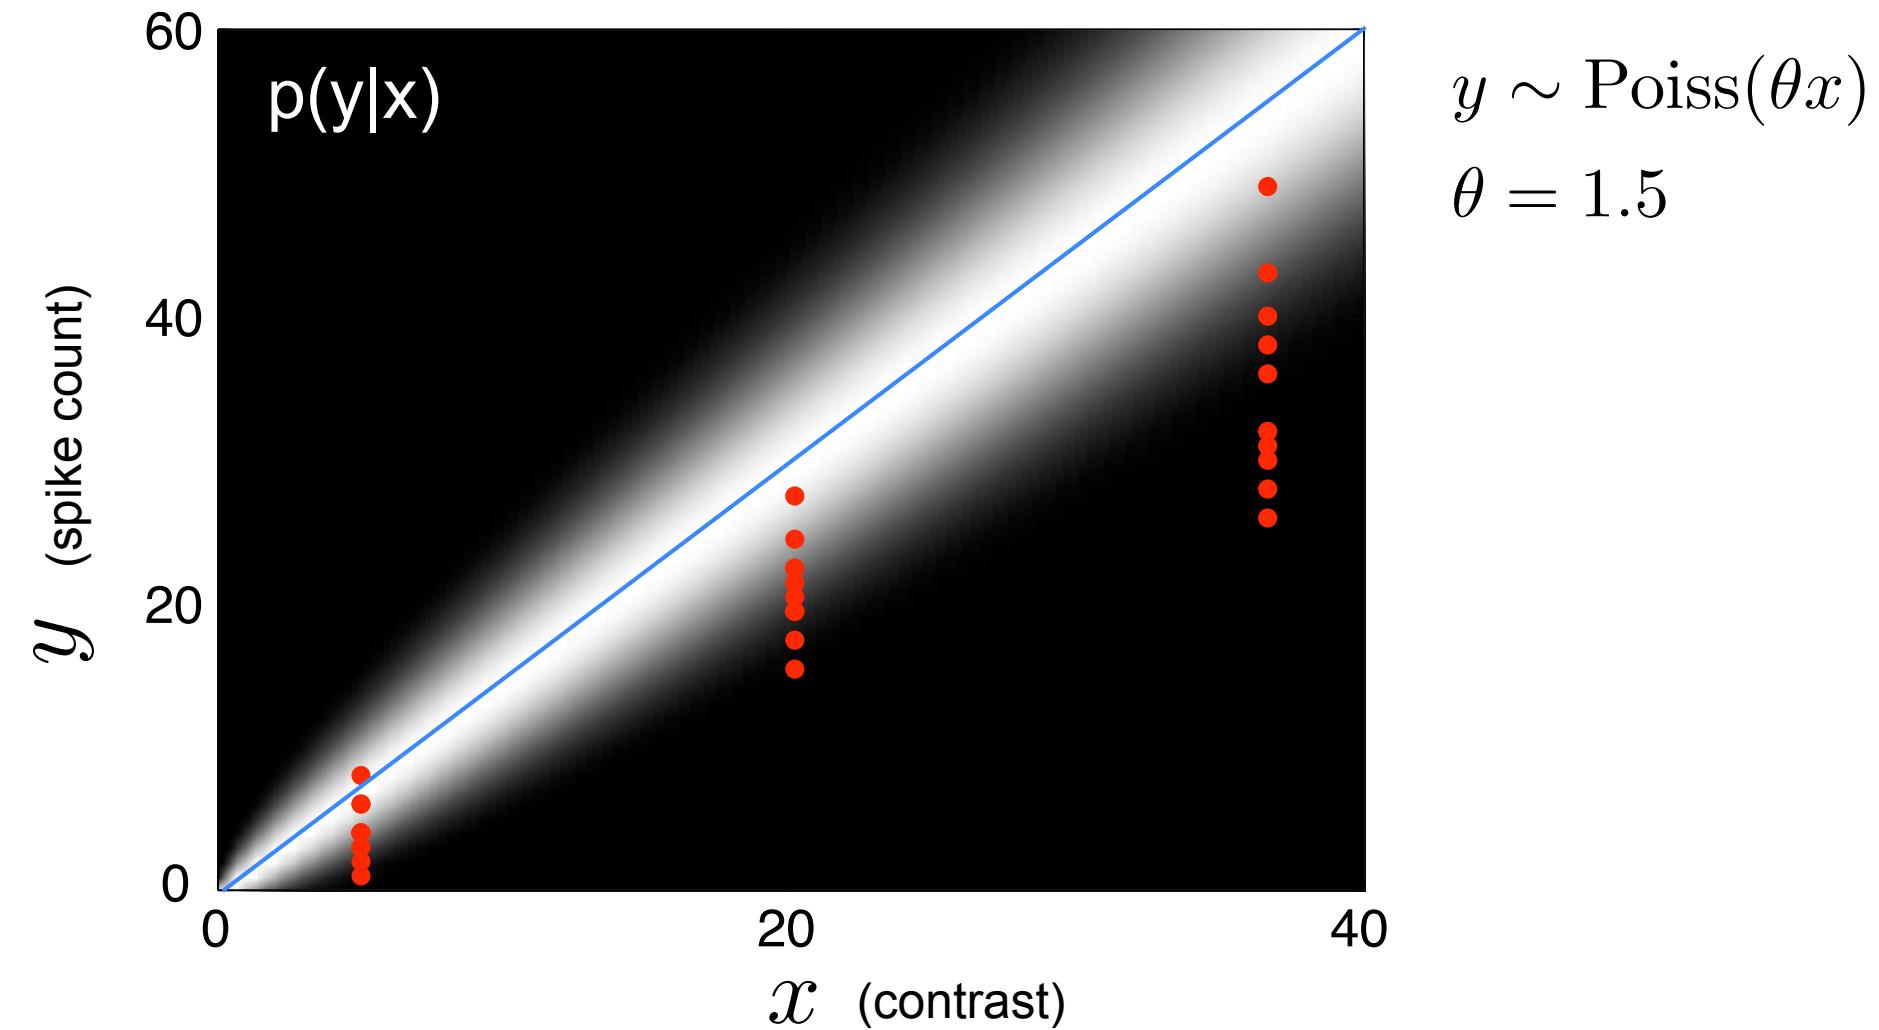
A: conditional independence across trials!

Q: when do we call $P(Y|X, \theta)$ a *likelihood*?

A: when considering it as a function of θ !

Maximum Likelihood Estimation:

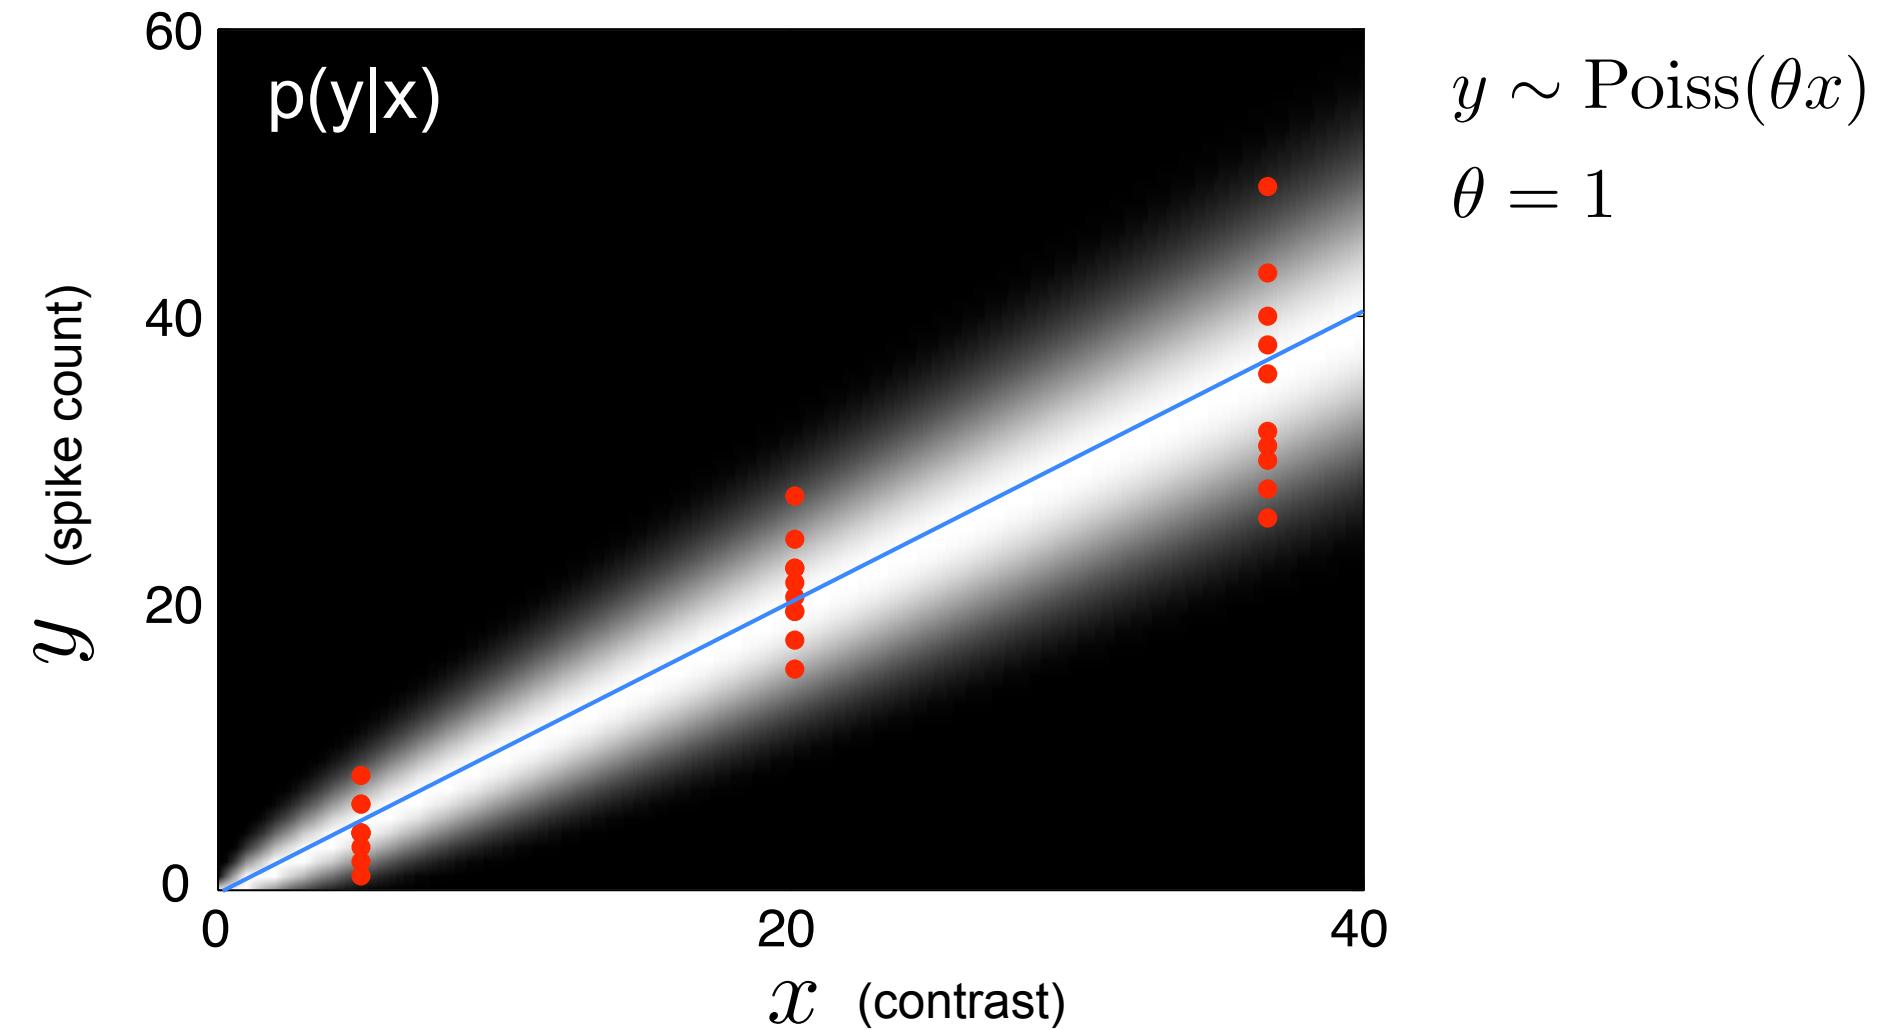
- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$



- could in theory do this by turning a knob

Maximum Likelihood Estimation:

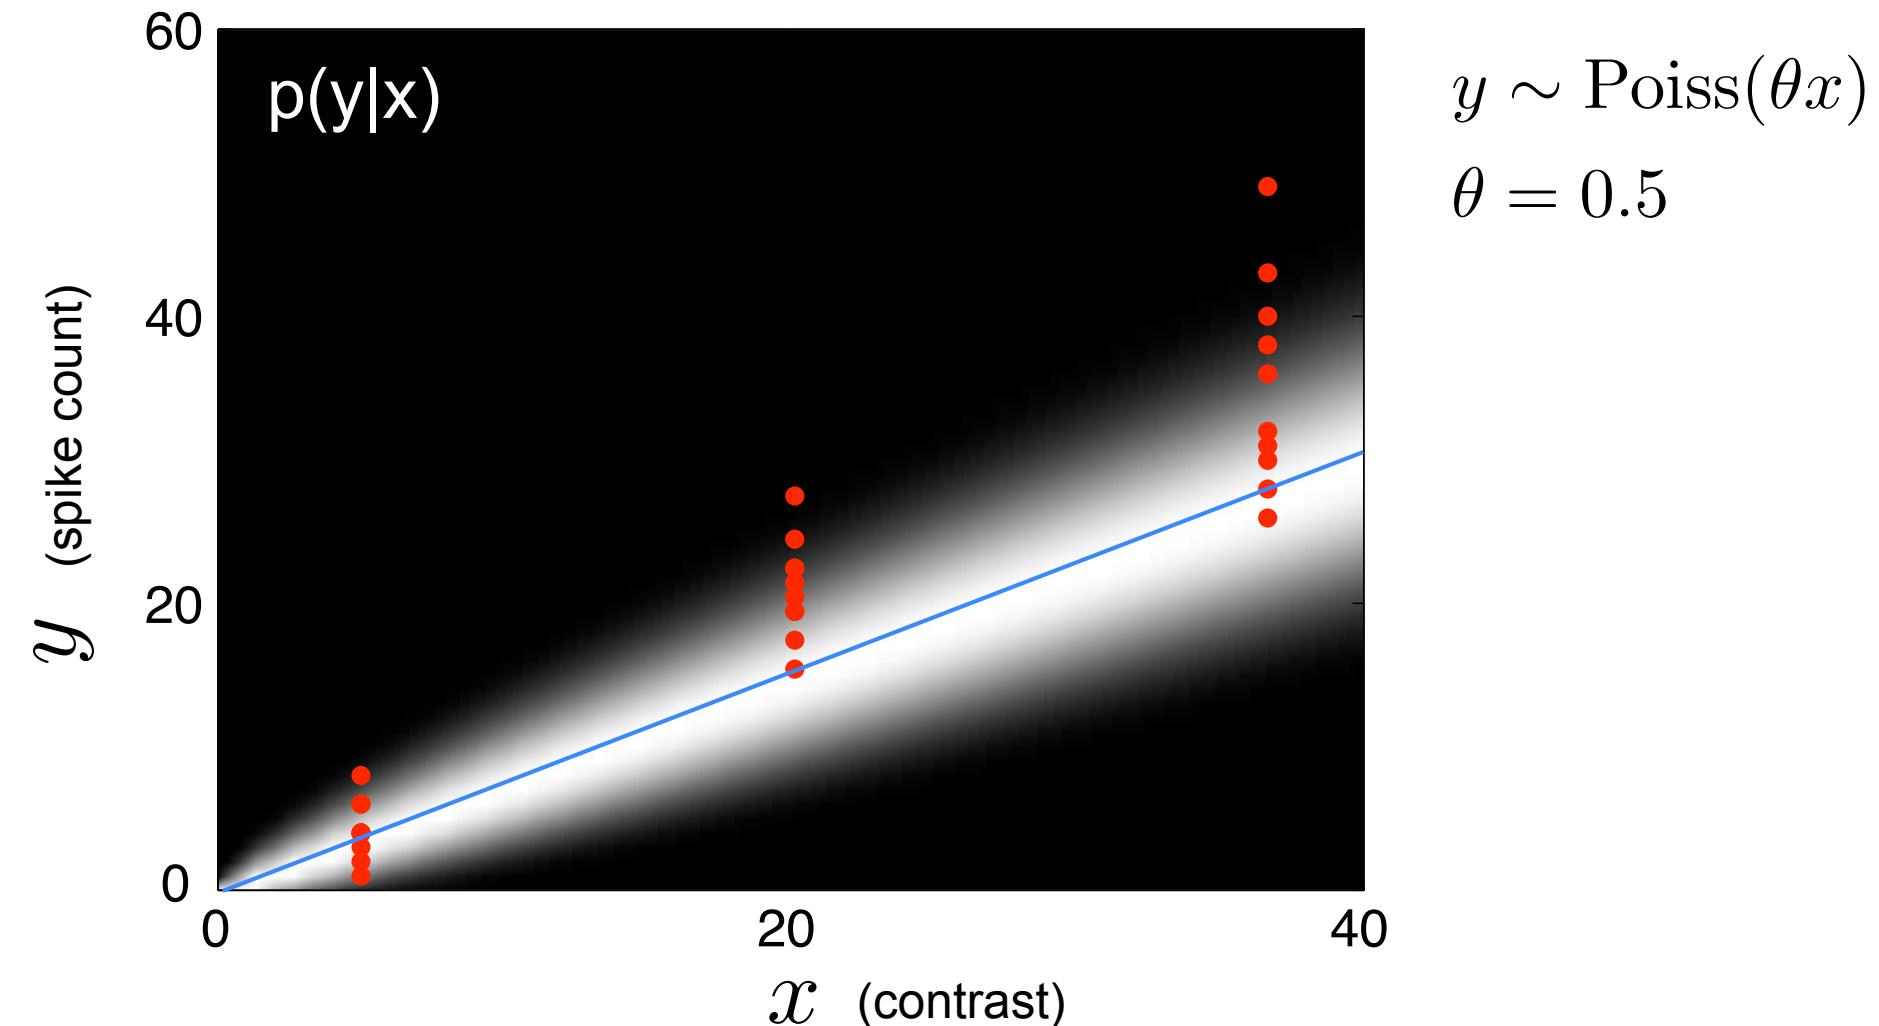
- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$



- could in theory do this by turning a knob

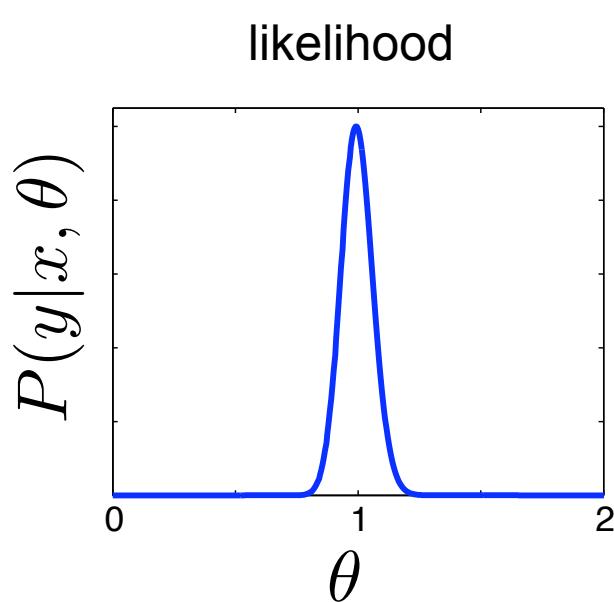
Maximum Likelihood Estimation:

- given observed data (Y, X) , find θ that maximizes $P(Y|X, \theta)$



- could in theory do this by turning a knob

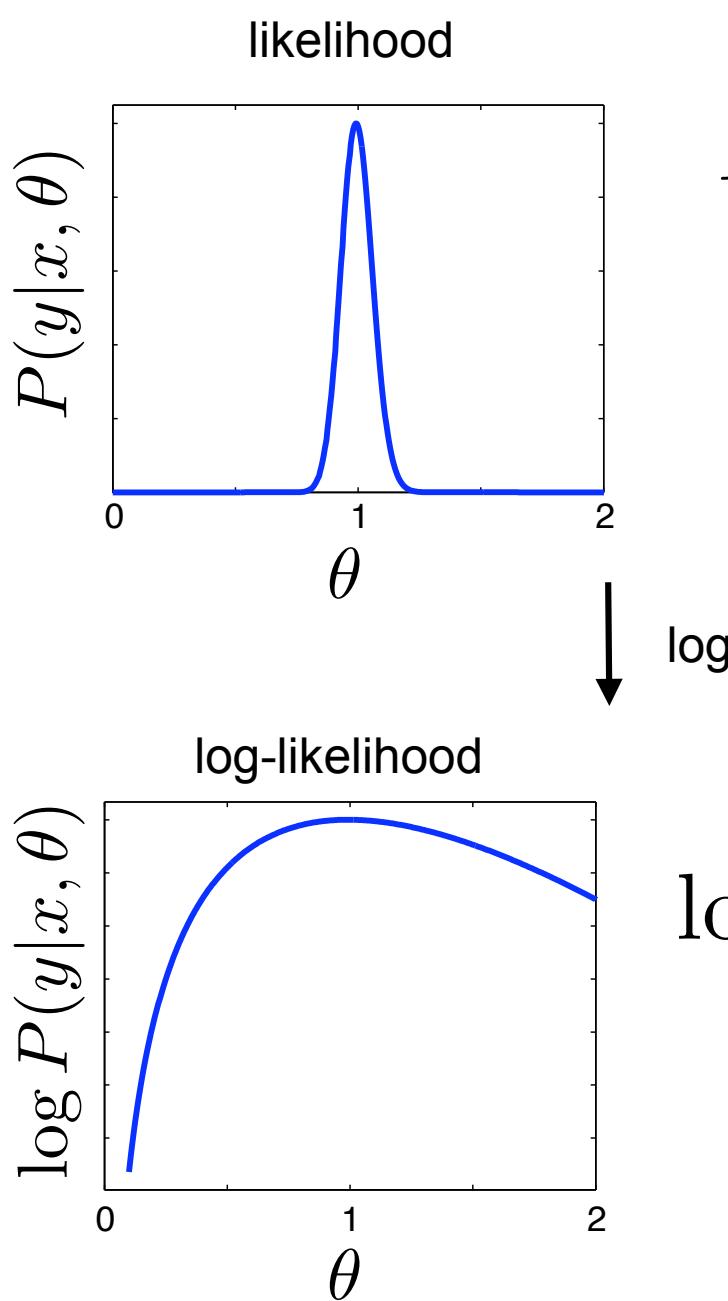
Likelihood function: $P(Y|X, \theta)$ as a function of θ



Because data are independent:

$$\begin{aligned}P(Y|X, \theta) &= \prod_i P(y_i|x_i, \theta) \\&= \prod_i \frac{1}{y_i!} (\theta x_i)^{y_i} e^{-(\theta x_i)}\end{aligned}$$

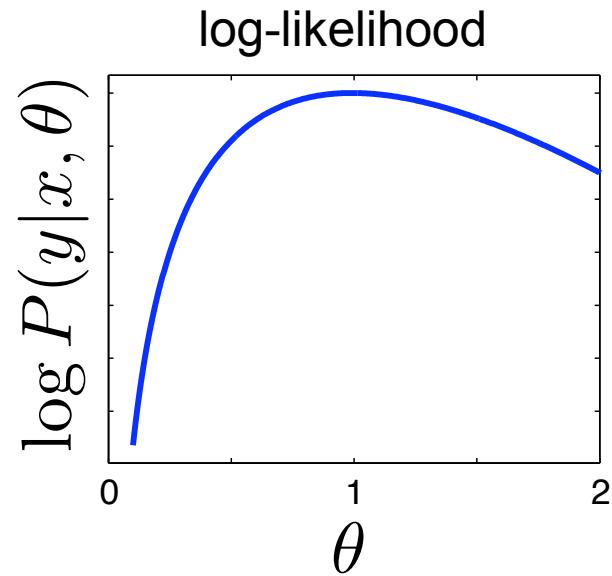
Likelihood function: $P(Y|X, \theta)$ as a function of θ



Because data are independent:

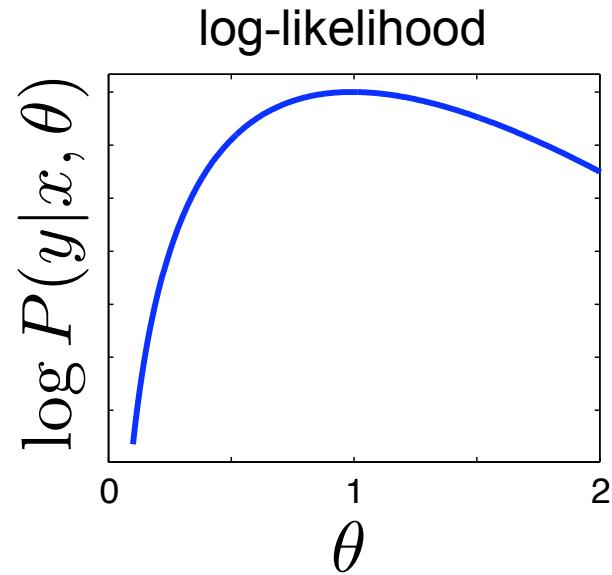
$$\begin{aligned}P(Y|X, \theta) &= \prod_i P(y_i|x_i, \theta) \\&= \prod_i \frac{1}{y_i!} (\theta x_i)^{y_i} e^{-(\theta x_i)}\end{aligned}$$

$$\begin{aligned}\log P(Y|X, \theta) &= \sum_i \log P(y_i|x_i, \theta) \\&= \sum y_i \log \theta - \theta x_i + c\end{aligned}$$



$$\begin{aligned}\log P(Y|X, \theta) &= \sum_i \log P(y_i|x_i, \theta) \\ &= \sum y_i \log \theta - \theta x_i + c \\ &= \log \theta(\sum y_i) - \theta(\sum x_i)\end{aligned}$$

Do it: solve for θ



$$\begin{aligned}
 \log P(Y|X, \theta) &= \sum_i \log P(y_i|x_i, \theta) \\
 &= \sum y_i \log \theta - \theta x_i + c \\
 &= \log \theta (\sum y_i) - \theta (\sum x_i)
 \end{aligned}$$

- Closed-form solution when model in “exponential family”

$$\frac{d}{d\theta} \log P(Y|X, \theta) = \frac{1}{\theta} \sum y_i - \sum x_i = 0$$

$$\implies \hat{\theta}_{ML} = \frac{\sum y_i}{\sum x_i}$$

Properties of the MLE

(maximum likelihood estimator)

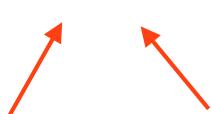
- consistent
(converges to true θ in limit of infinite data)
- efficient
(converges as quickly as possible,
i.e., achieves minimum possible asymptotic error)

Example 2: linear Gaussian neuron

spike count $y \sim \mathcal{N}(\mu, \sigma^2)$

spike rate $\mu = \theta x$

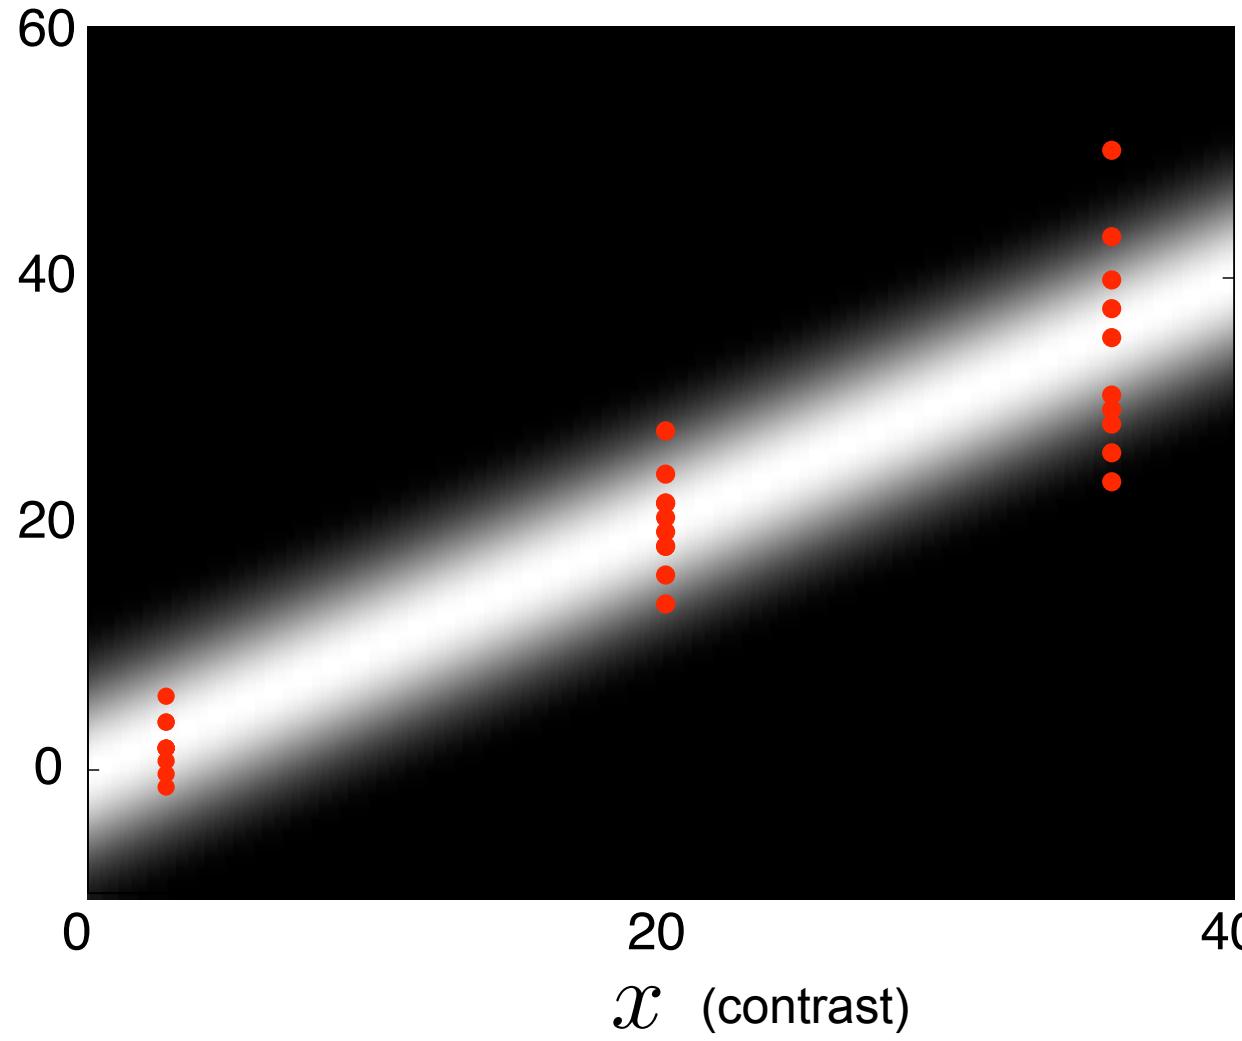
parameter stimulus



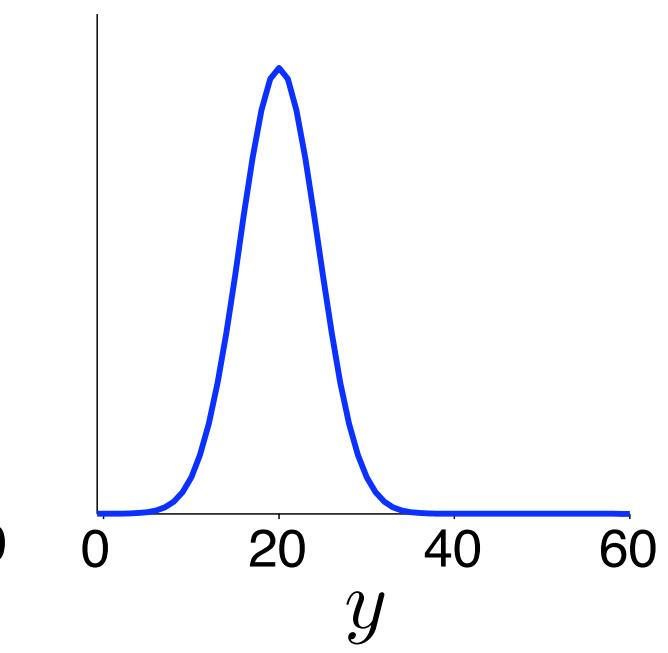
encoding model:

$$P(y|x, \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y-\theta x)^2}{2\sigma^2}}$$

$$\text{mean}(y) = \theta x$$
$$\text{var}(y) = \sigma^2$$



encoding distribution
 $p(y|x = 20)$



All slices have same width

$$P(y|x, \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y-\theta x)^2}{2\sigma^2}}$$

Log-Likelihood $\log P(Y|X, \theta) = -\sum \frac{(y_i - \theta x_i)^2}{2\sigma^2} + c$

Do it: differentiate, set to zero, and solve.

$$P(y|x, \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} e^{\frac{-(y-\theta x)^2}{2\sigma^2}}$$

Log-Likelihood $\log P(Y|X, \theta) = -\sum \frac{(y_i - \theta x_i)^2}{2\sigma^2} + c$

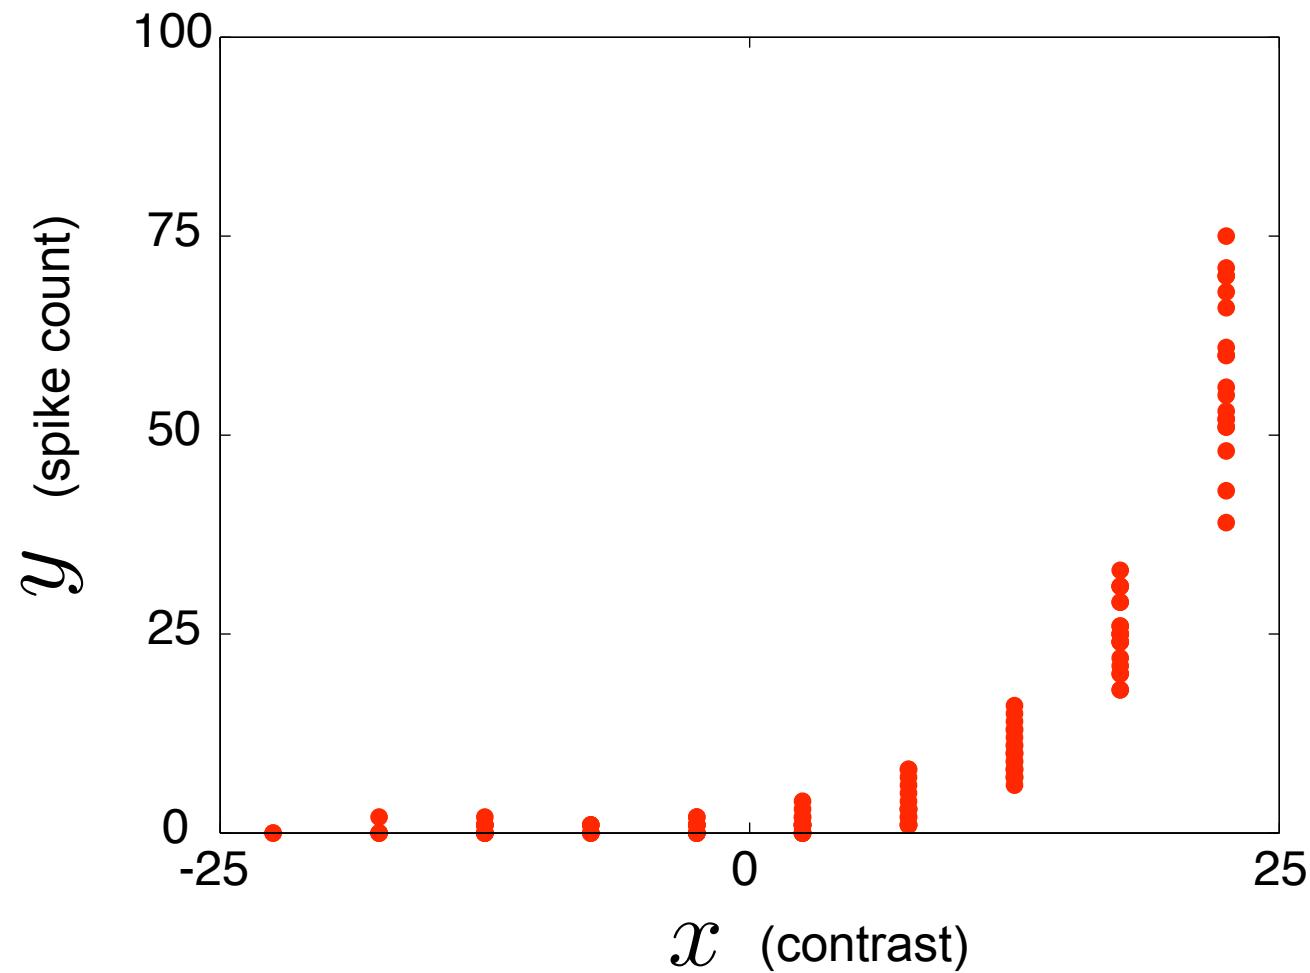
$$\frac{d}{d\theta} \log P(Y|X, \theta) = -\sum \frac{(y_i - \theta x_i)x_i}{\sigma^2} = 0$$

Maximum-Likelihood Estimator: $\hat{\theta}_{ML} = \frac{\sum y_i x_i}{\sum x_i^2}$

(“Least squares regression” solution)

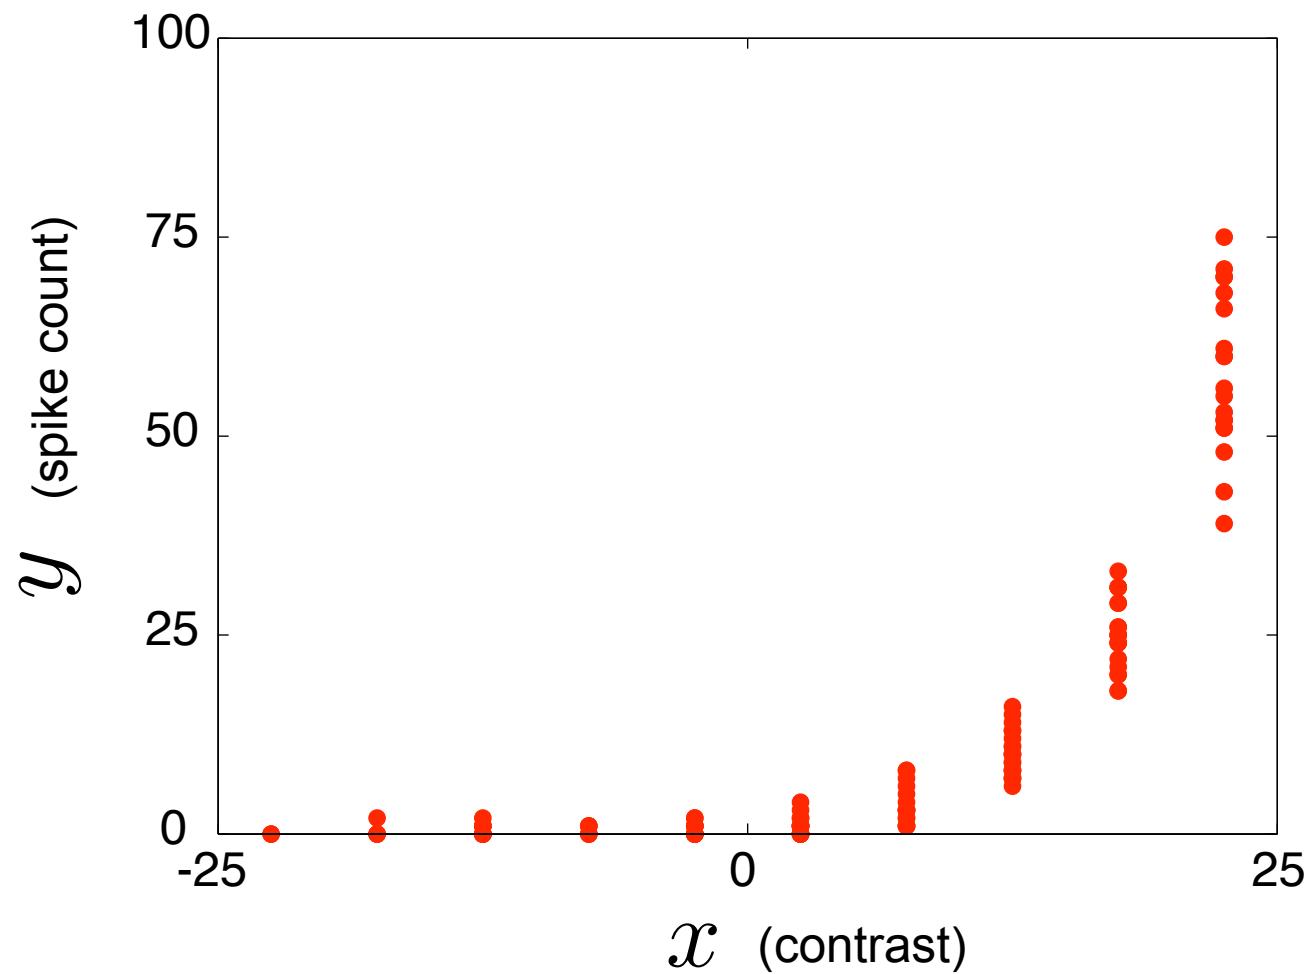
(Recall that for Poisson, $\hat{\theta}_{ML} = \frac{\sum y_i}{\sum x_i}$)

Example 3: unknown neuron



Be the computational neuroscientist: what model would you use?

Example 3: unknown neuron



More general setup: $y \sim Poiss(\lambda)$

This is a GLM!

$$\lambda = f(\theta x)$$

for some nonlinear function f

GLMs

- Be careful about terminology:

GLM

\neq

GLM

General Linear Model

Generalized Linear Model

(Nelder 1972)

Linear

Linear

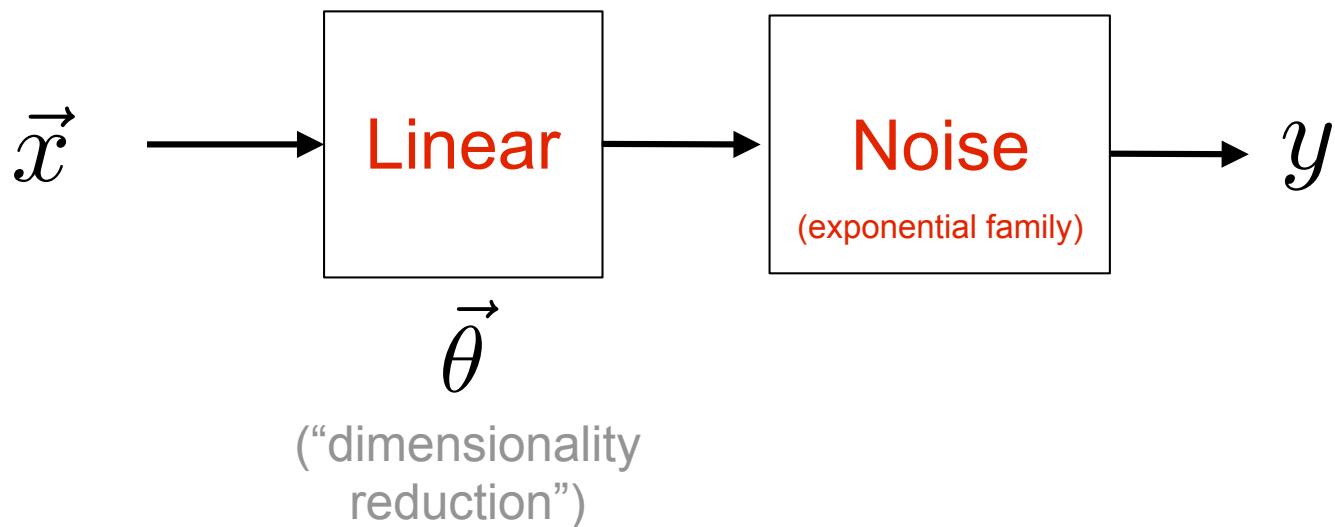
2003 interview with John Nelder...

Stephen Senn: I must confess to having some confusion when I was a young statistician between general linear models and generalized linear models. Do you regret the terminology?

John Nelder: I think probably I do. I suspect we should have found some more fancy name for it that would have stuck and not been confused with the general linear model, although general and generalized are not quite the same. I can see why it might have been better to have thought of something else.

Moral:
Be careful when naming your model!

1. General Linear Model (GLM)

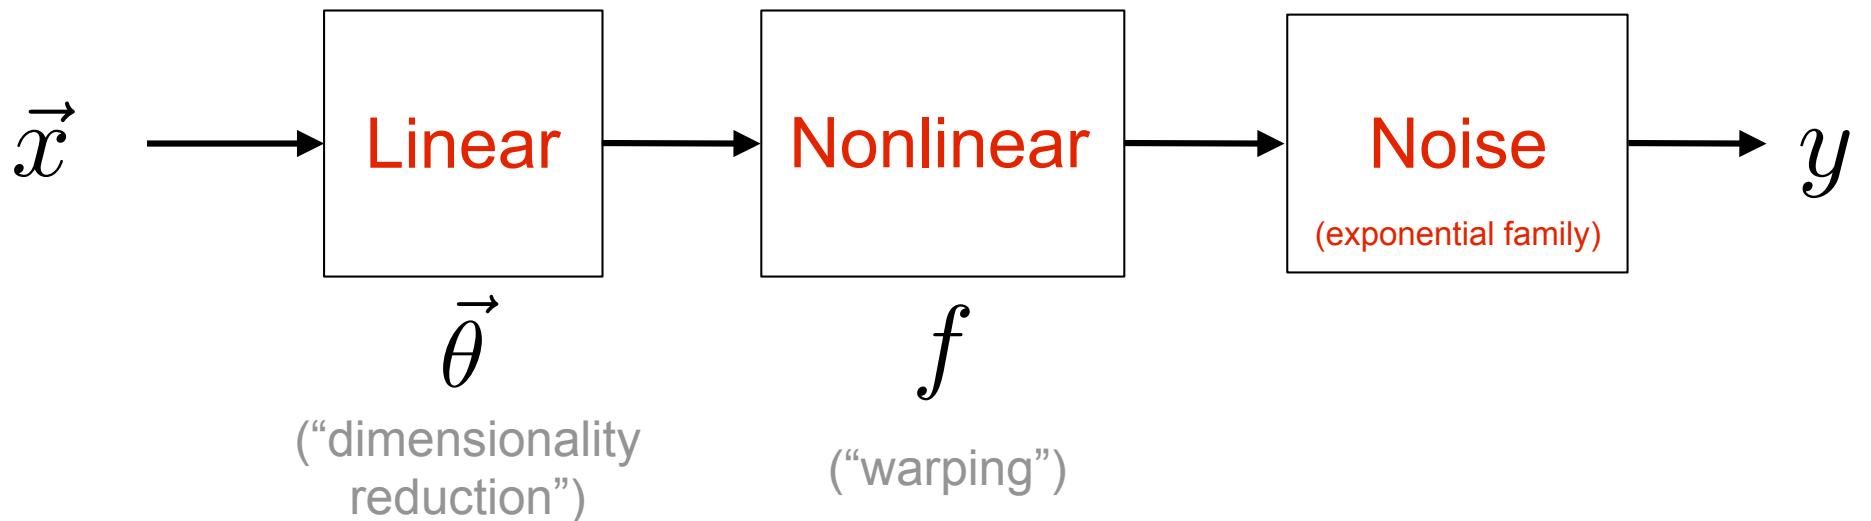


Examples:

- 1. Gaussian $y = \vec{\theta} \cdot \vec{x} + \epsilon$
- 2. Poisson $y \sim \text{Poiss}(\vec{\theta} \cdot \vec{x})$

where $\epsilon \sim \mathcal{N}(0, \sigma^2)$

2. Generalized Linear Model



- Examples:
1. Gaussian $y = f(\vec{\theta} \cdot \vec{x}) + \epsilon$
 2. Poisson $y \sim \text{Poiss}(f(\vec{\theta} \cdot \vec{x}))$
 3. Bernoulli $y \sim \text{Ber}(f(\vec{\theta} \cdot \vec{x}))$
- $\mathcal{N}(0, \sigma^2)$

2. Generalized Linear Model

Question: what restrictions (if any) must be placed on the nonlinearity f in each of the GLMs below?

Examples:

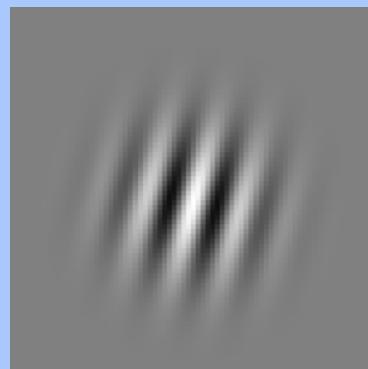
1. Gaussian $y = f(\vec{\theta} \cdot \vec{x}) + \epsilon$ $\mathcal{N}(0, \sigma^2)$
2. Poisson $y \sim \text{Poiss}(f(\vec{\theta} \cdot \vec{x}))$
3. Bernoulli $y \sim \text{Ber}(f(\vec{\theta} \cdot \vec{x}))$

2. Generalized Linear Model

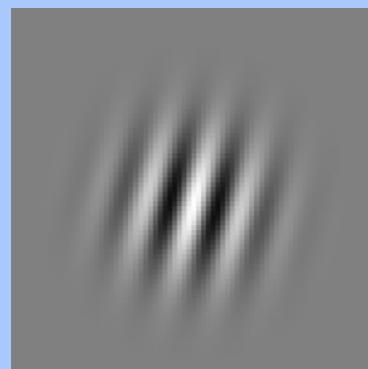
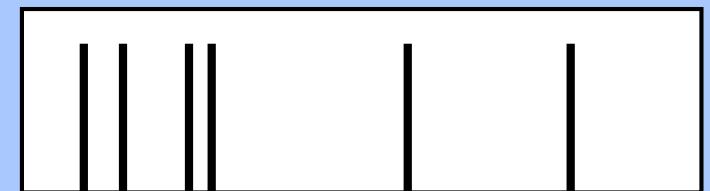
Question: what restrictions (if any) must be placed on the nonlinearity f in each of the GLMs below?

		$\mathcal{N}(0, \sigma^2)$
none	1. Gaussian	$y = f(\vec{\theta} \cdot \vec{x}) + \epsilon$
non-negative	2. Poisson	$y \sim \text{Poiss}(f(\vec{\theta} \cdot \vec{x}))$
bounded between 0 and 1	3. Bernoulli	$y \sim \text{Ber}(f(\vec{\theta} \cdot \vec{x}))$

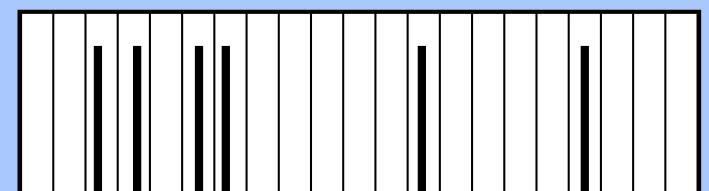
Part II: from spike counts to spike trains!



“6 spikes”



spike train



001101100000100001000

spike count / bin

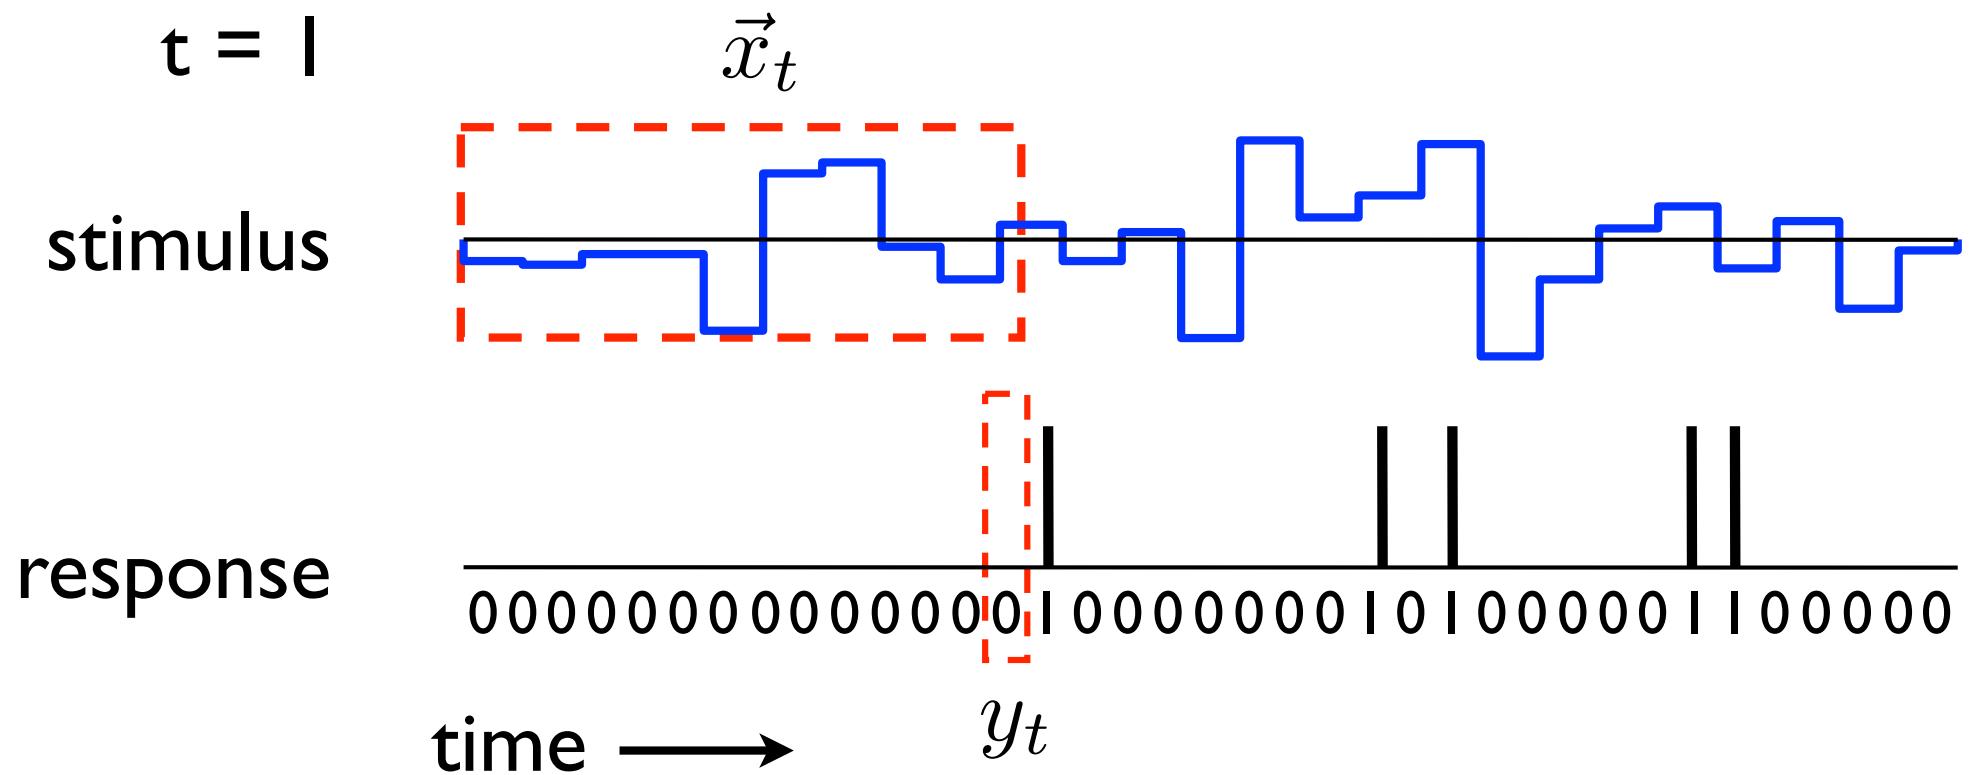
response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

linear filter vector stimulus at time t

walk through the data
one time bin at a time

$$t = 1$$



response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

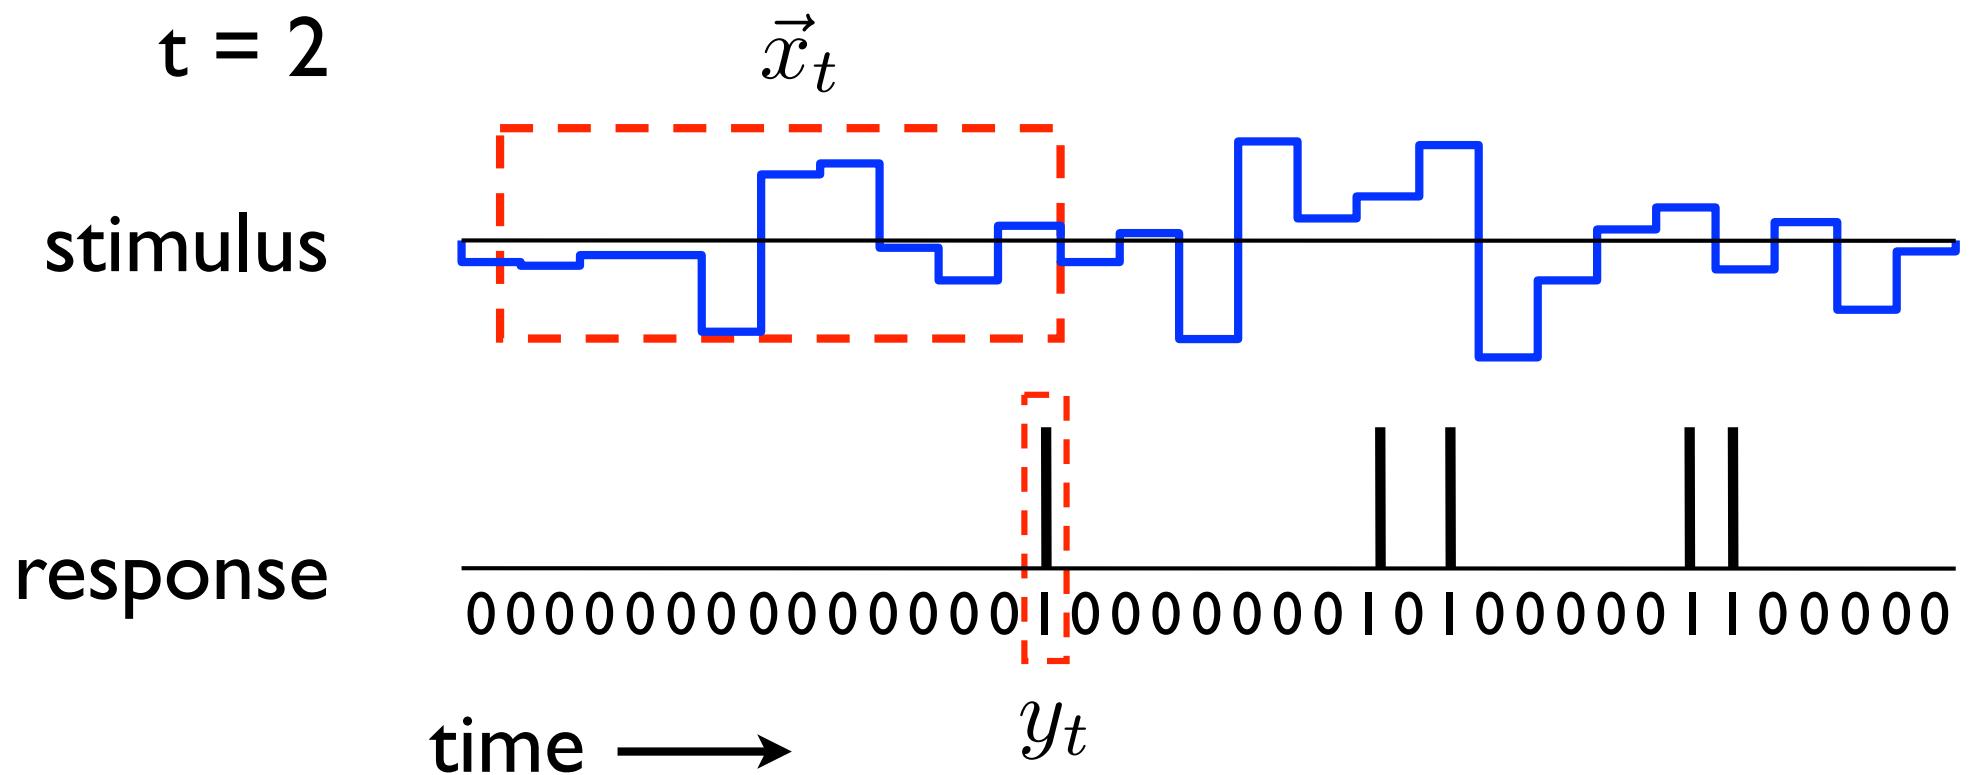
linear
filter

vector stimulus
at time t

$N(0, \sigma^2)$

walk through the data
one time bin at a time

$$t = 2$$



response at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

↑
linear
filter

↑
vector stimulus
at time t

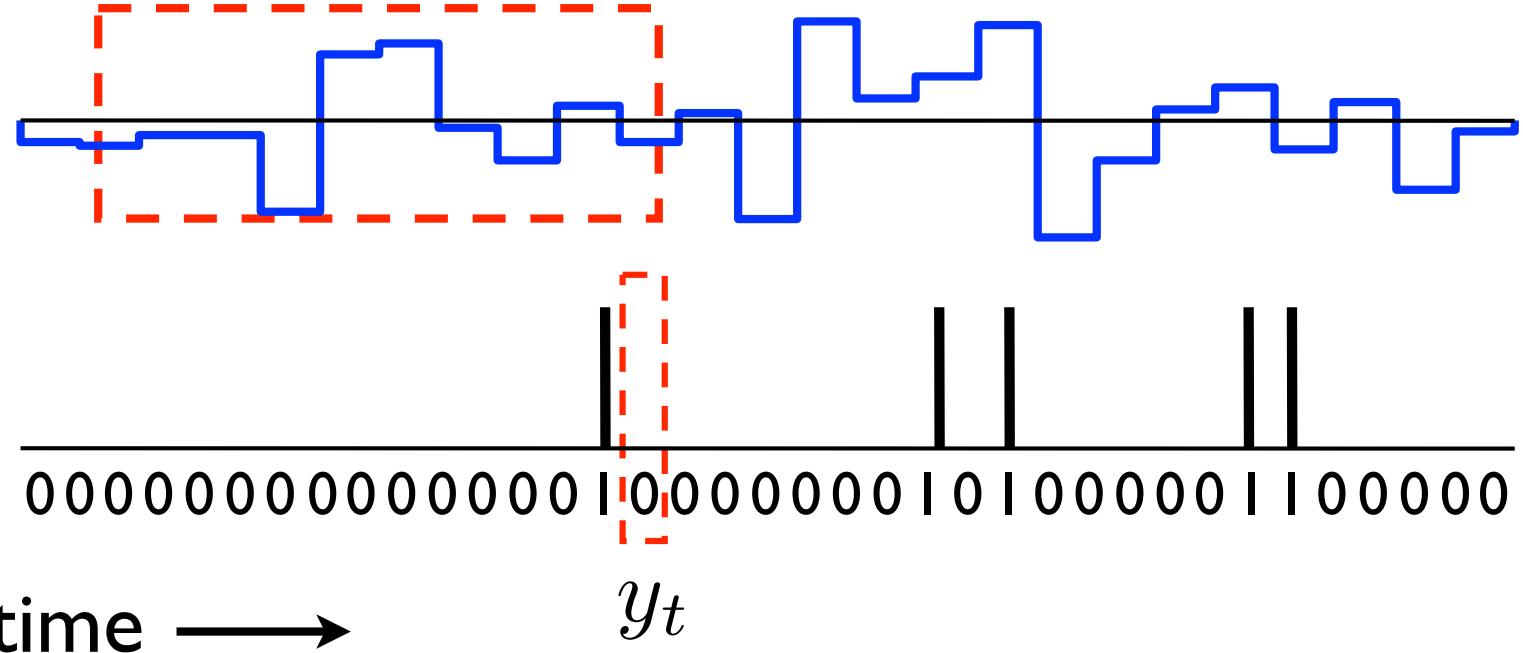
↗
 $N(0, \sigma^2)$

walk through the data
one time bin at a time

t = 3

\vec{x}_t

stimulus



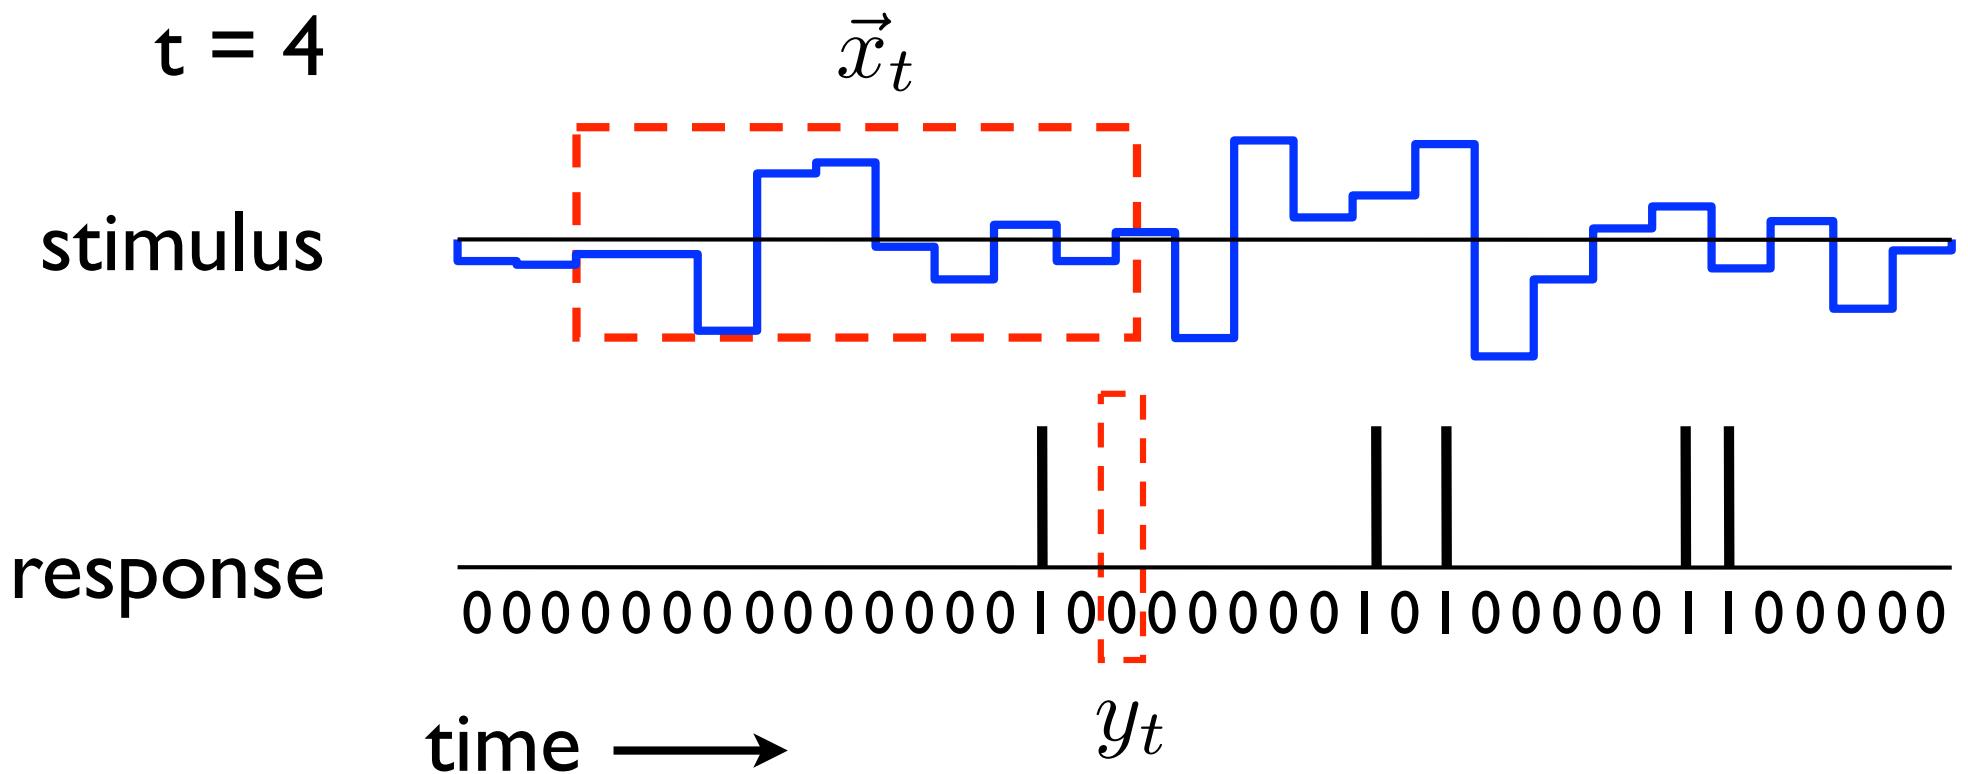
response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

↑ ↑
linear filter vector stimulus
at time t

walk through the data
one time bin at a time

$t = 4$



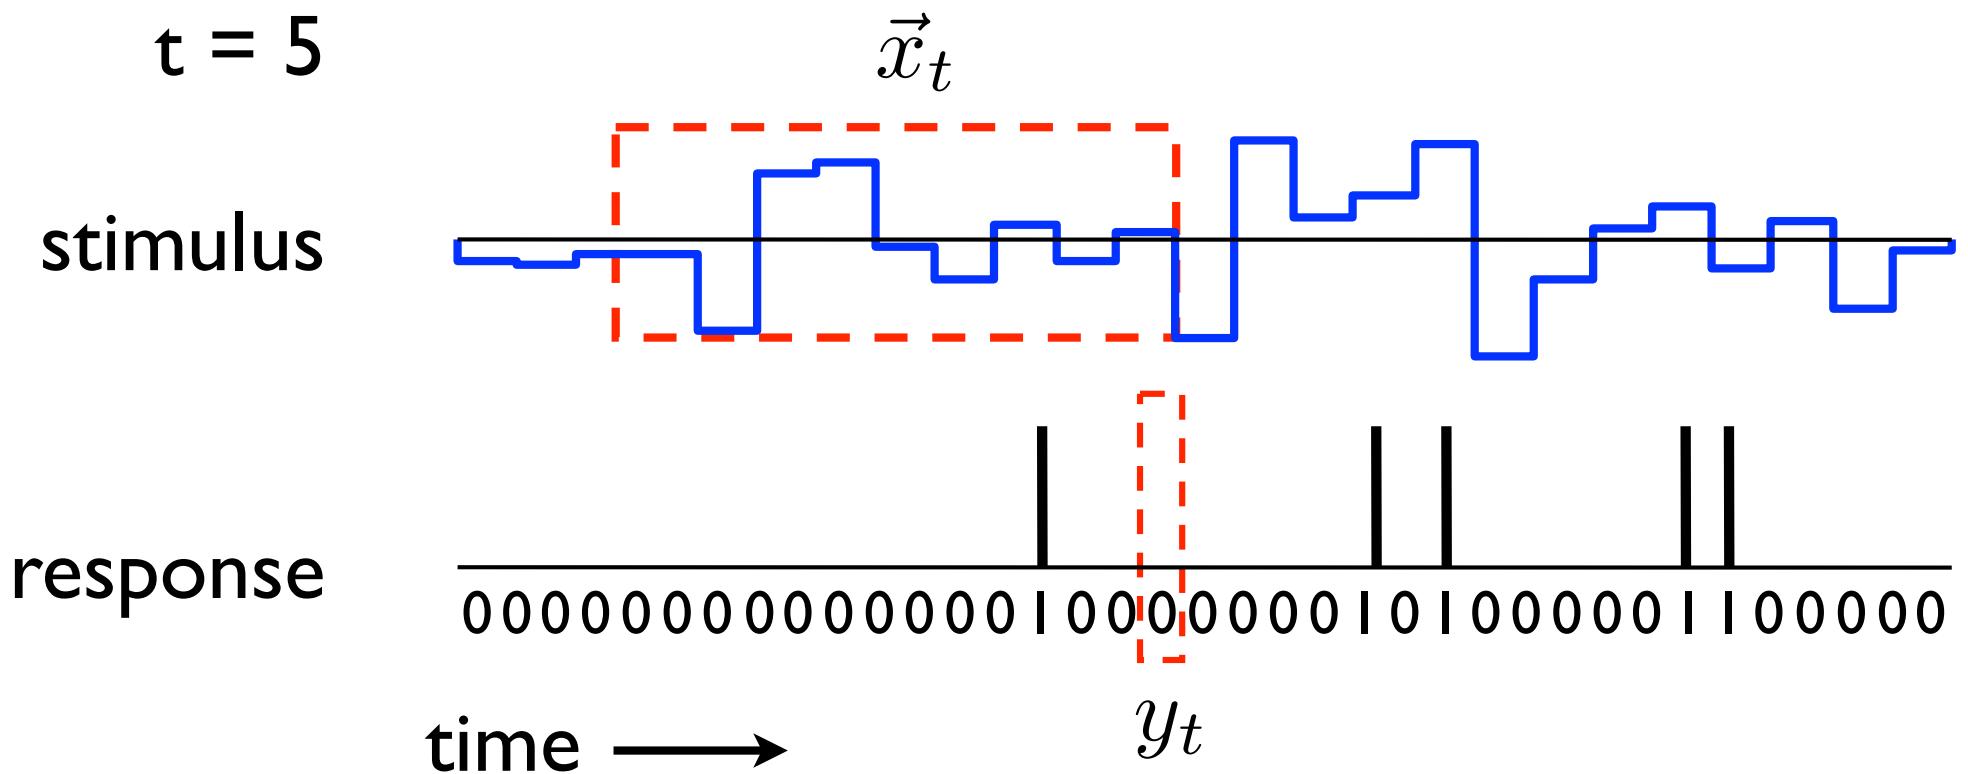
response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

↑ ↑
linear filter vector stimulus
at time t

walk through the data
one time bin at a time

$t = 5$



response
at time t

$$y_t = \vec{k} \cdot \vec{x}_t + \text{noise}$$

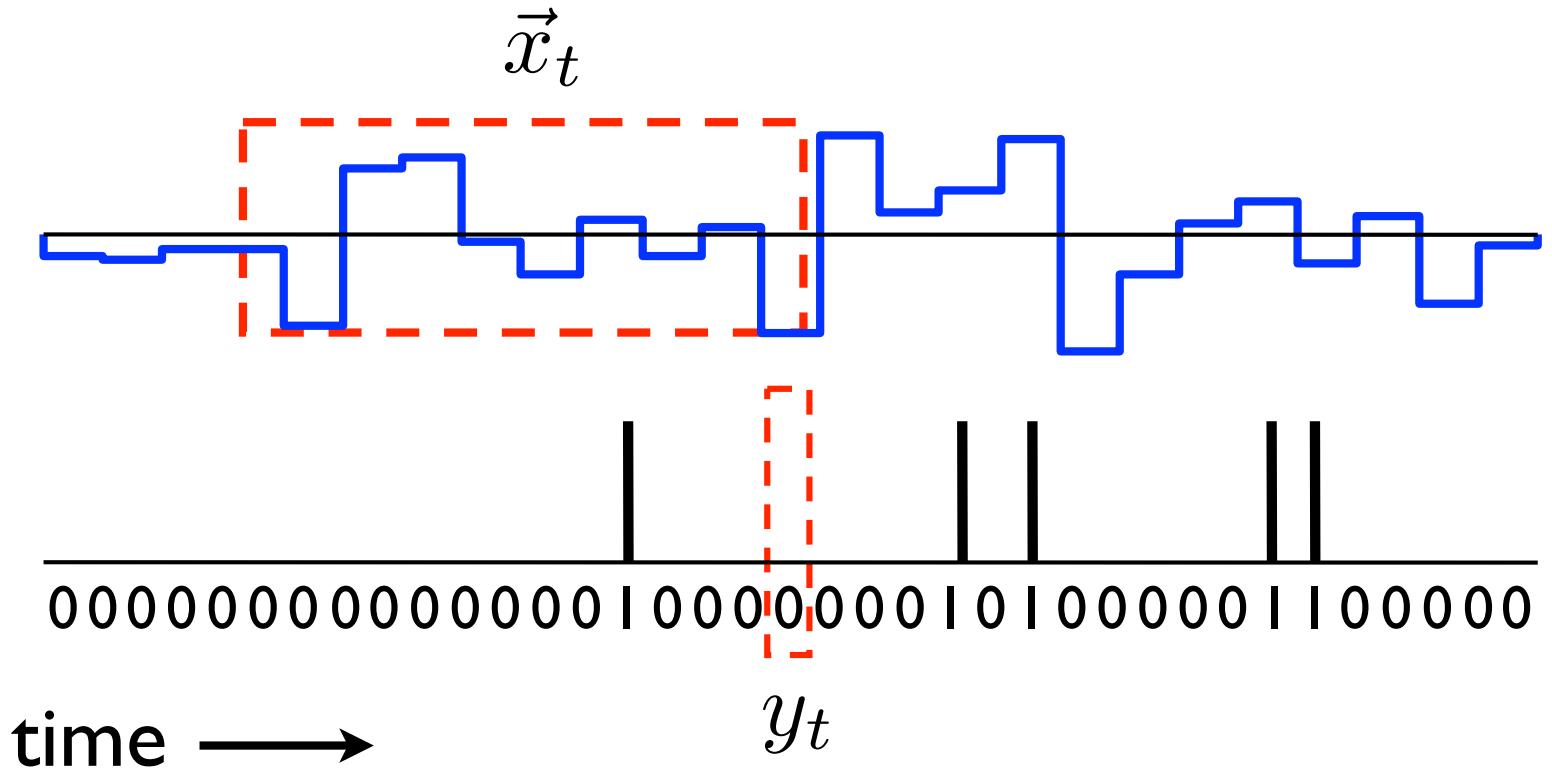
linear filter vector stimulus
at time t

walk through the data
one time bin at a time

$t = 6$

stimulus

response



Build up to following matrix version:

Y = $X\vec{k}$ + *noise*

↑
time ↓

=

$\begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \end{bmatrix}$

$\begin{bmatrix} \text{blue step functions} & \text{blue step functions} & \text{blue step functions} \end{bmatrix}$

$\begin{bmatrix} \vec{k} \\ \vdots \end{bmatrix}$

design matrix

“Hankel matrix”
(every row is a shifted copy of the row above)

(see also: Toeplitz matrix)

Build up to following matrix version:

$$Y = X\vec{k} + \text{noise}$$

“Linear-Gaussian” GLM:

(aka “least-squares regression”)

$$\hat{k} = \underbrace{(X^T X)^{-1}}_{\text{stimulus covariance}} \underbrace{X^T Y}_{\text{spike-triggered avg (STA)}}$$

Formally:

$$Y = X\vec{k} + \vec{\epsilon}$$

↑ ↑ ↗

N(0, \sigma^2 I)
iid Gaussian noise vector

time ↓

$$= \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \text{blue step function} \\ \text{blue step function} \\ \text{blue step function} \\ \vdots \end{bmatrix} \begin{bmatrix} \vec{k} \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \vdots \end{bmatrix}$$

equivalent to writing:

$$Y|X, \vec{k} \sim \mathcal{N}(X\vec{k}, \sigma^2 I)$$

or

Take log,
differentiate and
set to zero.

$$P(Y|X, \vec{k}) = \frac{1}{|2\pi\sigma^2 I|^{\frac{T}{2}}} \exp \left(-\frac{1}{2\sigma^2} (Y - X\vec{k})^\top (Y - X\vec{k}) \right)$$

Formally:

$$Y = X\vec{k} + \vec{\epsilon}$$

↑ ↑ ↗

$\begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \end{bmatrix} = \begin{bmatrix} \text{blue step functions} \\ \vdots \end{bmatrix} \begin{bmatrix} \vec{k} \end{bmatrix} + \begin{bmatrix} \epsilon_1 \\ \epsilon_2 \\ \epsilon_3 \\ \vdots \end{bmatrix}$

time ↓

$N(0, \sigma^2 I)$
 iid Gaussian noise vector

equivalent to w

$P(Y|X,$

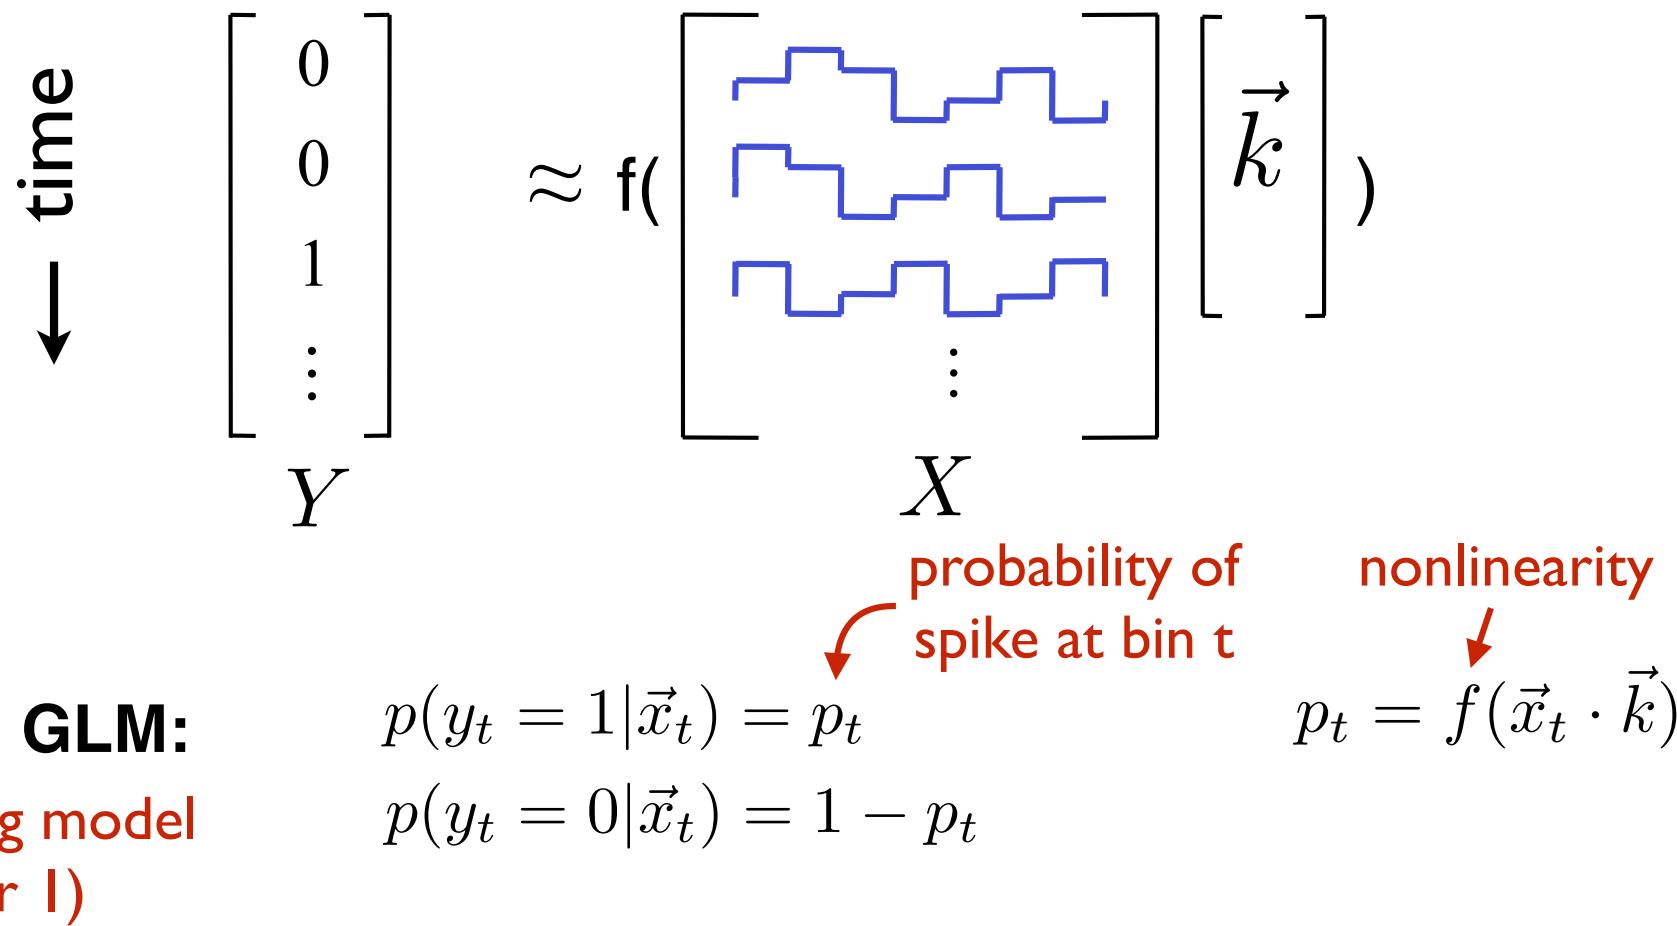
General point: minimizing a sum of squares is *always* equivalent to maximizing likelihood under a Gaussian noise model!

$\sigma^2 I)$

Take log,
differentiate and
set to zero.

$X\vec{k})^\top (Y - X\vec{k})$)

But noise is *not* Gaussian: binary responses



But noise is *not* Gaussian: binary responses

$$\begin{array}{c} \text{time} \\ \downarrow \\ \left[\begin{array}{c} 0 \\ 0 \\ 1 \\ \vdots \\ Y \end{array} \right] \approx f(\left[\begin{array}{c} \text{ } \\ \text{ } \\ \text{ } \\ \text{ } \\ X \end{array} \right] \left[\begin{array}{c} \vec{k} \\ \vec{k} \\ \vec{k} \\ \vdots \\ \vec{k} \end{array} \right]) \end{array}$$

Bernoulli GLM:
(coin flipping model
 $y = 0$ or 1)

probability of spike at bin t

$$p(y_t = 1 | \vec{x}_t) = p_t$$

nonlinearity

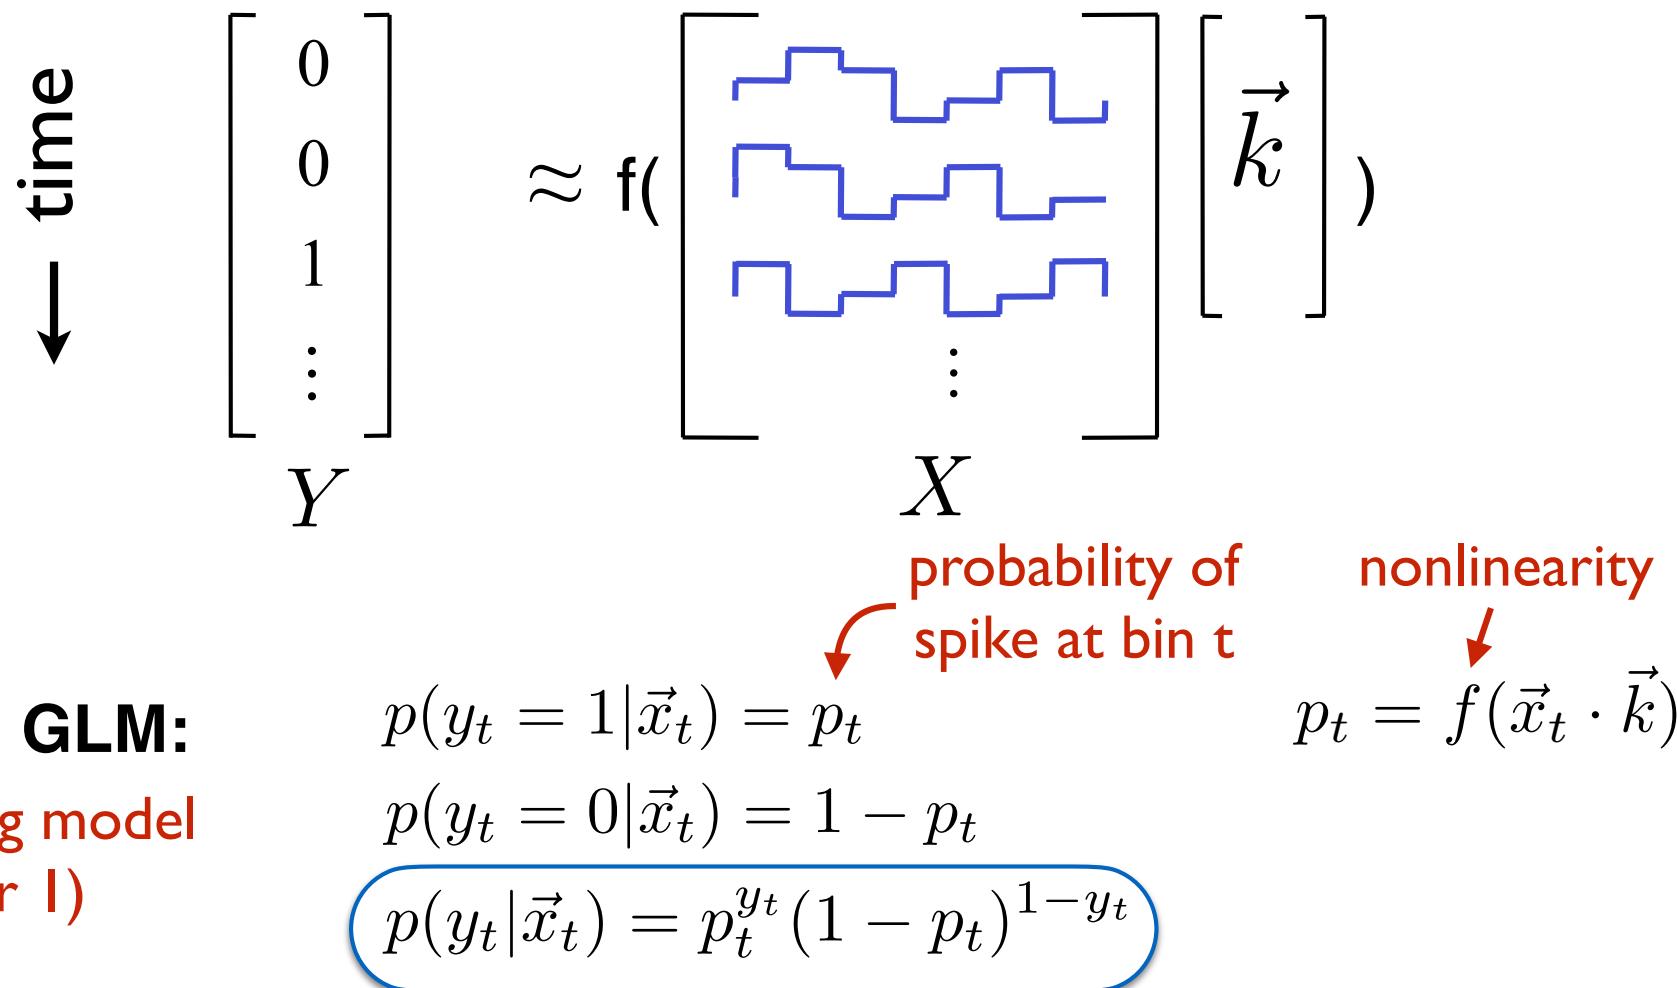
$$p_t = f(\vec{x}_t \cdot \vec{k})$$

$p(y_t = 0 | \vec{x}_t) = 1 - p_t$

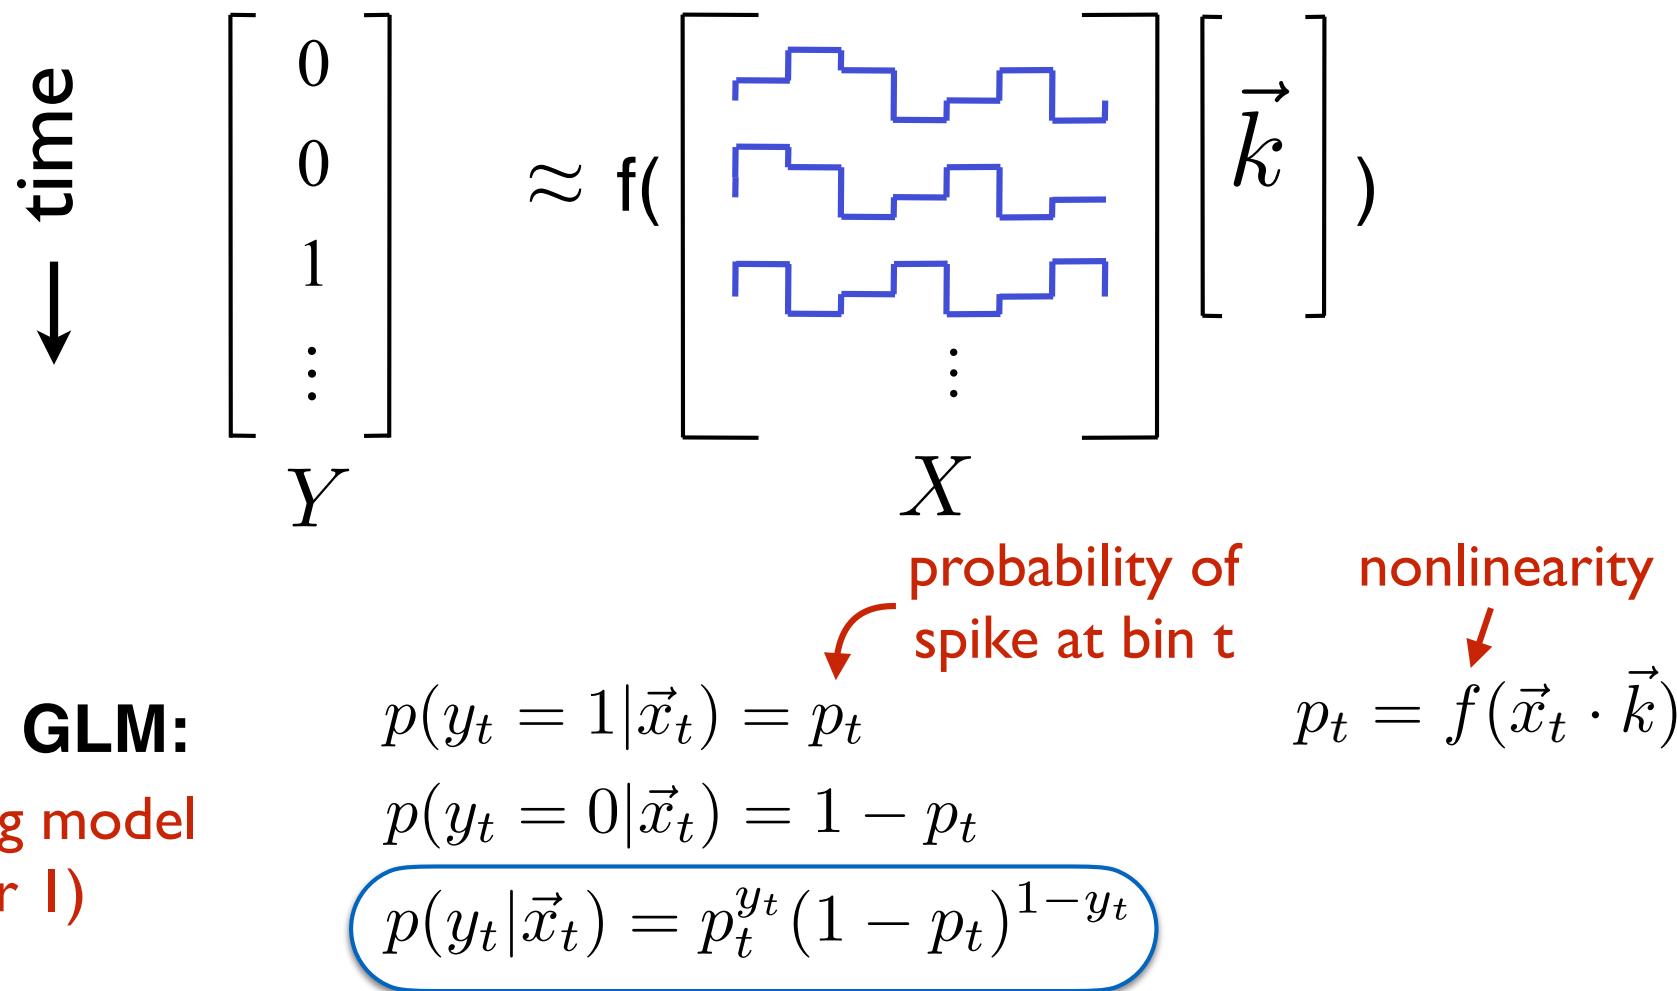
logistic regression: $f(z) = \frac{1}{1 + e^{-z}}$ logistic function

- so logistic regression is a special case of a Bernoulli GLM!

But noise is *not* Gaussian: binary responses



But noise is *not* Gaussian: binary responses



log-likelihood: $\mathcal{L} = \sum_{t=1}^T \left(y_t \log f(\vec{x}_t \cdot \vec{k}) + (1 - y_t) \log(1 - f(\vec{x}_t \cdot \vec{k})) \right)$

But noise is *not* Gaussian: binary responses

$$\begin{array}{c} \text{time} \\ \downarrow \\ \left[\begin{array}{c} 0 \\ 0 \\ 1 \\ \vdots \\ Y \end{array} \right] \end{array} \approx f(\left[\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \vdots \\ X \end{array} \right]) \left[\begin{array}{c} \vec{k} \\ \vec{k} \\ \vec{k} \\ \vdots \\ \vec{k} \end{array} \right]$$

Matrix-vector version:

$$\mathcal{L} = Y^\top \log f(X \vec{k}) + (1 - Y)^\top \log(1 - f(X \vec{k}))$$

in python:

```
z = f(X @ k) # compute P(heads) for each time bin
L = np.dot(y, np.log(z)) + np.dot(1-y, np.log(1-z))
```

log-likelihood: $\mathcal{L} = \sum_{t=1}^T \left(y_t \log f(\vec{x}_t \cdot \vec{k}) + (1 - y_t) \log(1 - f(\vec{x}_t \cdot \vec{k})) \right)$

non-negative integer responses ("Poisson regression")

Poisson GLM:
(integer $y \geq 0$)

$$\begin{aligned} \text{firing rate} \quad \lambda_t &= f(\vec{x}_t \cdot \vec{k}) \\ \text{(aka "conditional intensity")} \end{aligned}$$
$$y_t | \vec{x}_t, \vec{k} \sim \text{Poiss}(\Delta \lambda_t)$$

nonlinearity

time bin size

encoding distribution:

$$p(y_t | \vec{x}_t, \vec{k}) = \frac{(\Delta \lambda_t)^{y_t}}{y_t!} e^{-\Delta \lambda_t}$$

log-likelihood:

$$\mathcal{L} = \log p(Y|X, \vec{k}) = \sum_t \left(y_t \log f(\vec{x}_t \cdot \vec{k}) - \Delta f(\vec{x}_t \cdot \vec{k}) \right) + \text{const}$$

python:

```
z = f(x @ k) # conditional intensity
L = np.dot(y, np.log(z)) - dt*np.sum(z)
```

GLM: Summary:

I. “Linear-Gaussian” GLM: $Y|X, \vec{k} \sim \mathcal{N}(X\vec{k}, \sigma^2 I)$

(note: assuming “identity”
nonlinearity)

$$\hat{k} = (X^T X)^{-1} X^T Y$$

```
k_hat = np.linalg.solve(X.T@X, X.T@Y)
```

2. Bernoulli GLM: $y_t|\vec{x}_t, \vec{k} \sim \text{Ber}(f(\vec{x}_t \cdot \vec{k}))$

log-likelihood: $\mathcal{L} = Y^\top \log f(X\vec{k}) - (1 - Y)^\top \log(1 - f(X\vec{k}))$

```
z = f(X @ k) # compute P(heads) for each time bin  
L = np.dot(y, np.log(z)) + np.dot(1-y, np.log(1-z))
```

3. Poisson GLM: $y_t|\vec{x}_t, \vec{k} \sim \text{Poiss}(\Delta\lambda_t)$

log-likelihood: $\mathcal{L} = Y^\top \log f(X\vec{k}) - \Delta\mathbf{1}^\top f(X\vec{k})$

```
z = f(X @ k) # conditional intensity  
L = np.dot(y, np.log(z)) - dt*np.sum(z)
```

GLM tutorial:

MATLAB: <https://github.com/pillowlab/GLMspiketraintutorial>

python: https://github.com/pillowlab/GLMspiketraintutorial_python

- **tutorial1_PoissonGLM.m** - fitting of a linear-Gaussian GLM and Poisson GLM (aka LNP model) to RGC neurons stimulated with temporal white noise stimulus.
- **tutorial2_spikehистcoupledGLM.m** - fitting of a Poisson GLM with spike-history and coupling between neurons.
- **tutorial3_regularization_linGauss.m** - regularizing linear-Gaussian model parameters using maximum a posteriori (MAP) estimation under two kinds of priors:
 - (1) ridge regression (aka "L2 penalty");
 - (2) L2 smoothing prior (aka "graph Laplacian").
- **tutorial4_regularization_PoissonGLM.m** - MAP estimation of Poisson-GLM parameters using same two priors as in tutorial3.

Do it: tutorial1_PoissonGLM.m

1. cd into correct directory ('GLMspiketraintutorial')
2. open tutorial: use ctrl-enter, ctrl-down

Dataset:

- retinal ganglion cell spike trains [Uzzell & Chichilnisky 2004]
- 2 OFF cells, 2 on cells
- full-field, binary white noise stimuli

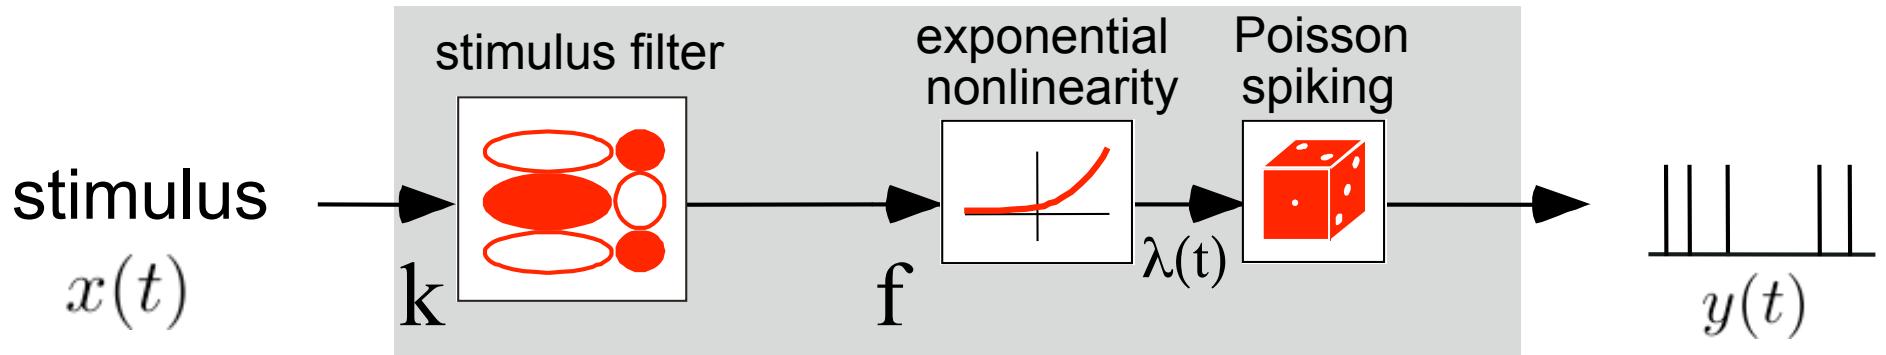
Do it: tutorial1_PoissonGLM.m

1. cd into correct directory ('GLMspiketraintutorial')
2. open tutorial: use ctrl-enter, ctrl-down

Topics:

- 1-2. load data, bin spike trains
3. build design matrix (slow and fast versions)
- 4a. compute STA (for visualization)
- 4b. linear-Gaussian GLM ("least-squares regression")
5. Poisson GLM: fit assuming exponential nonlinearity
6. Poisson GLM: non-parametric estimate of nonlinearity
7. model comparison: log-likelihood & mutual information
8. model comparison: AIC
9. simulate from fitted model

GLM (aka Linear-Nonlinear-Poisson)

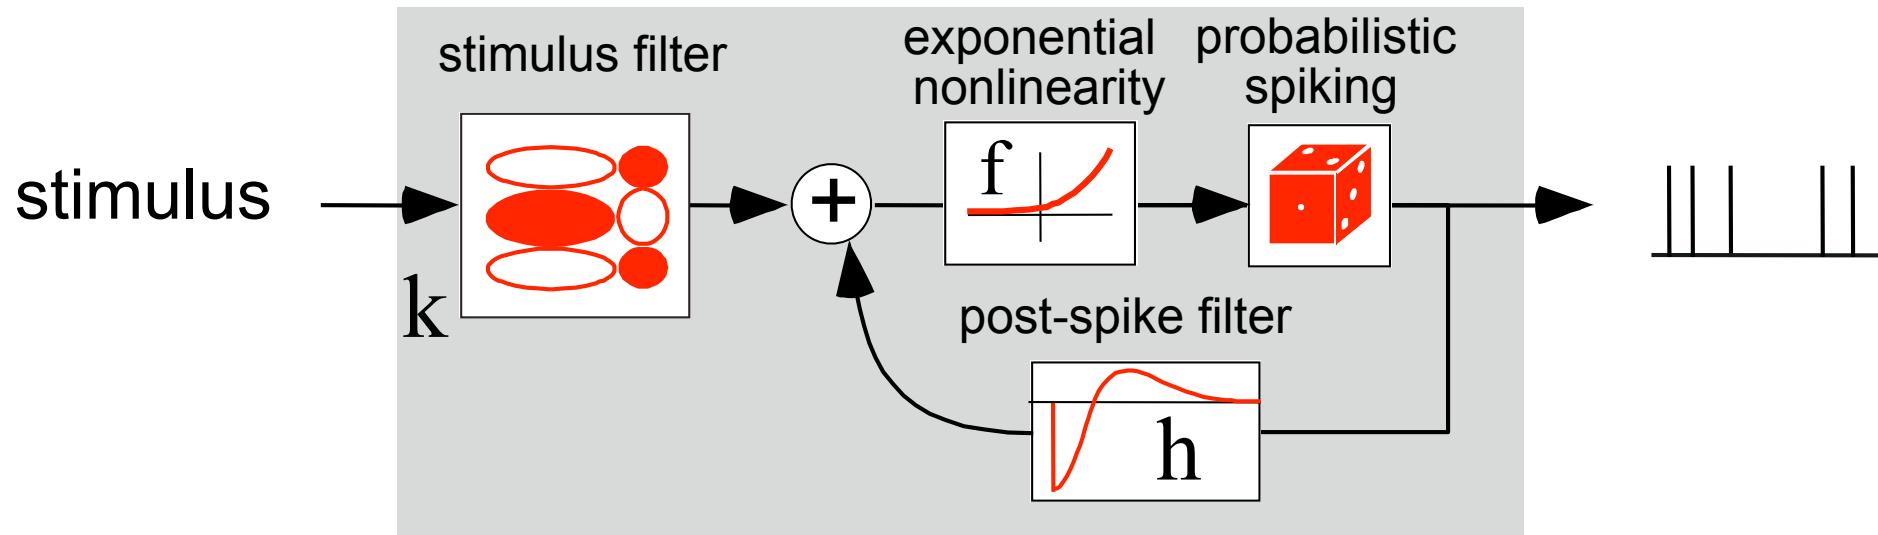


conditional intensity
(spike rate)

$$\lambda(t) = f(k \cdot x(t))$$

- output: Poisson process
- problem: assumes spiking depends only on stimulus!

Poisson GLM with spike-history dependence



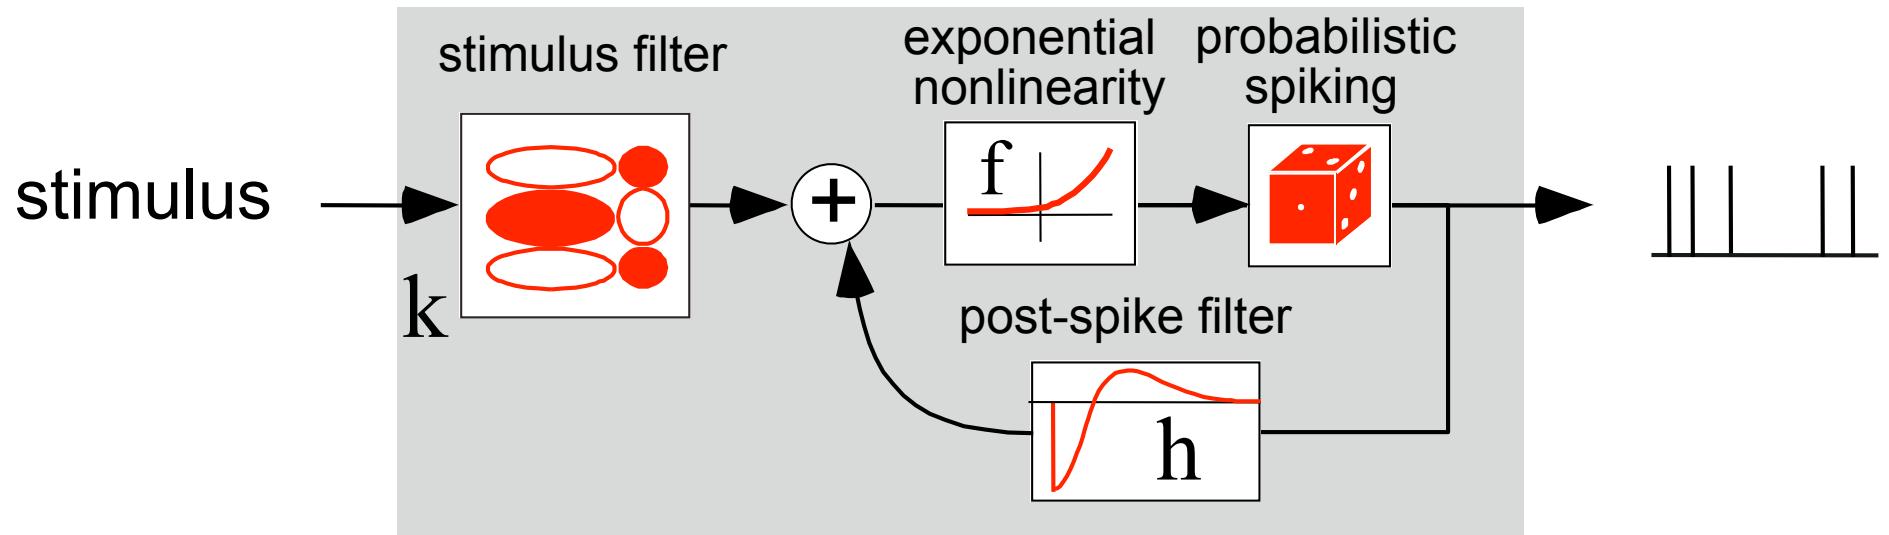
conditional intensity
(spike rate)

$$\lambda(t) = f(\vec{k} \cdot \vec{x}(t) + \vec{h} \cdot \vec{y}_{hist}(t))$$

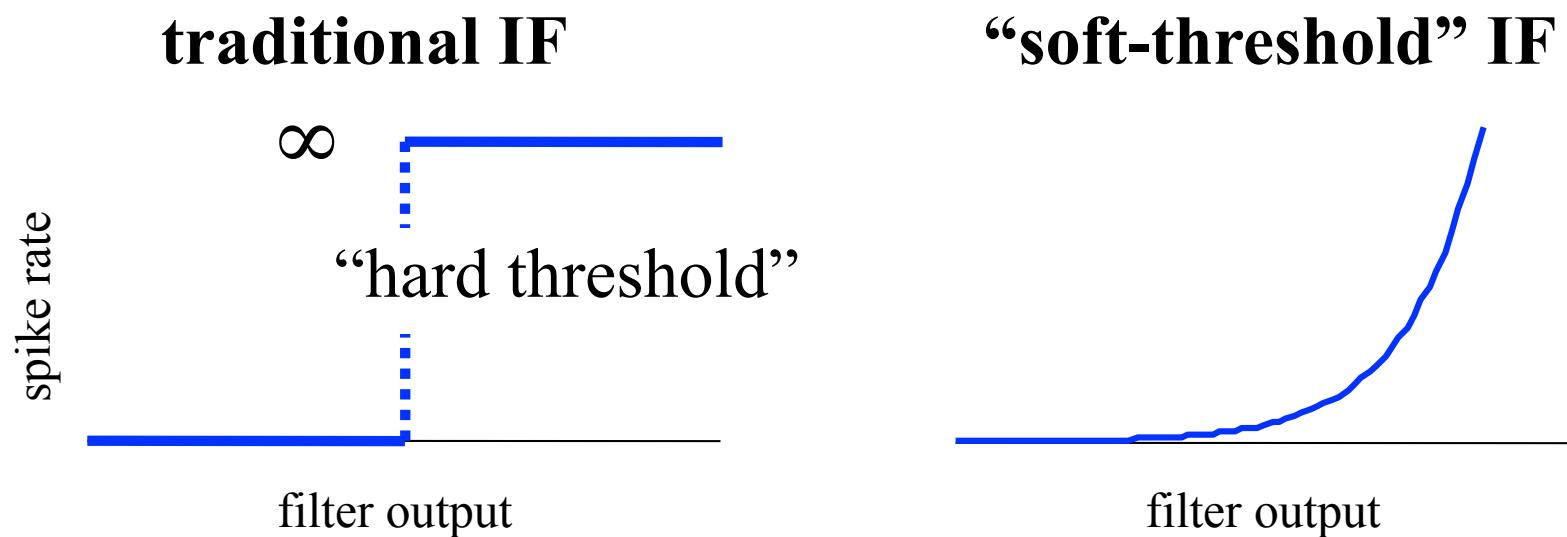
$$= e^{\vec{k} \cdot \vec{x}(t)} \cdot e^{\vec{h} \cdot \vec{y}_{hist}(t)}$$

- output: no longer a Poisson process

Poisson GLM with spike-history dependence

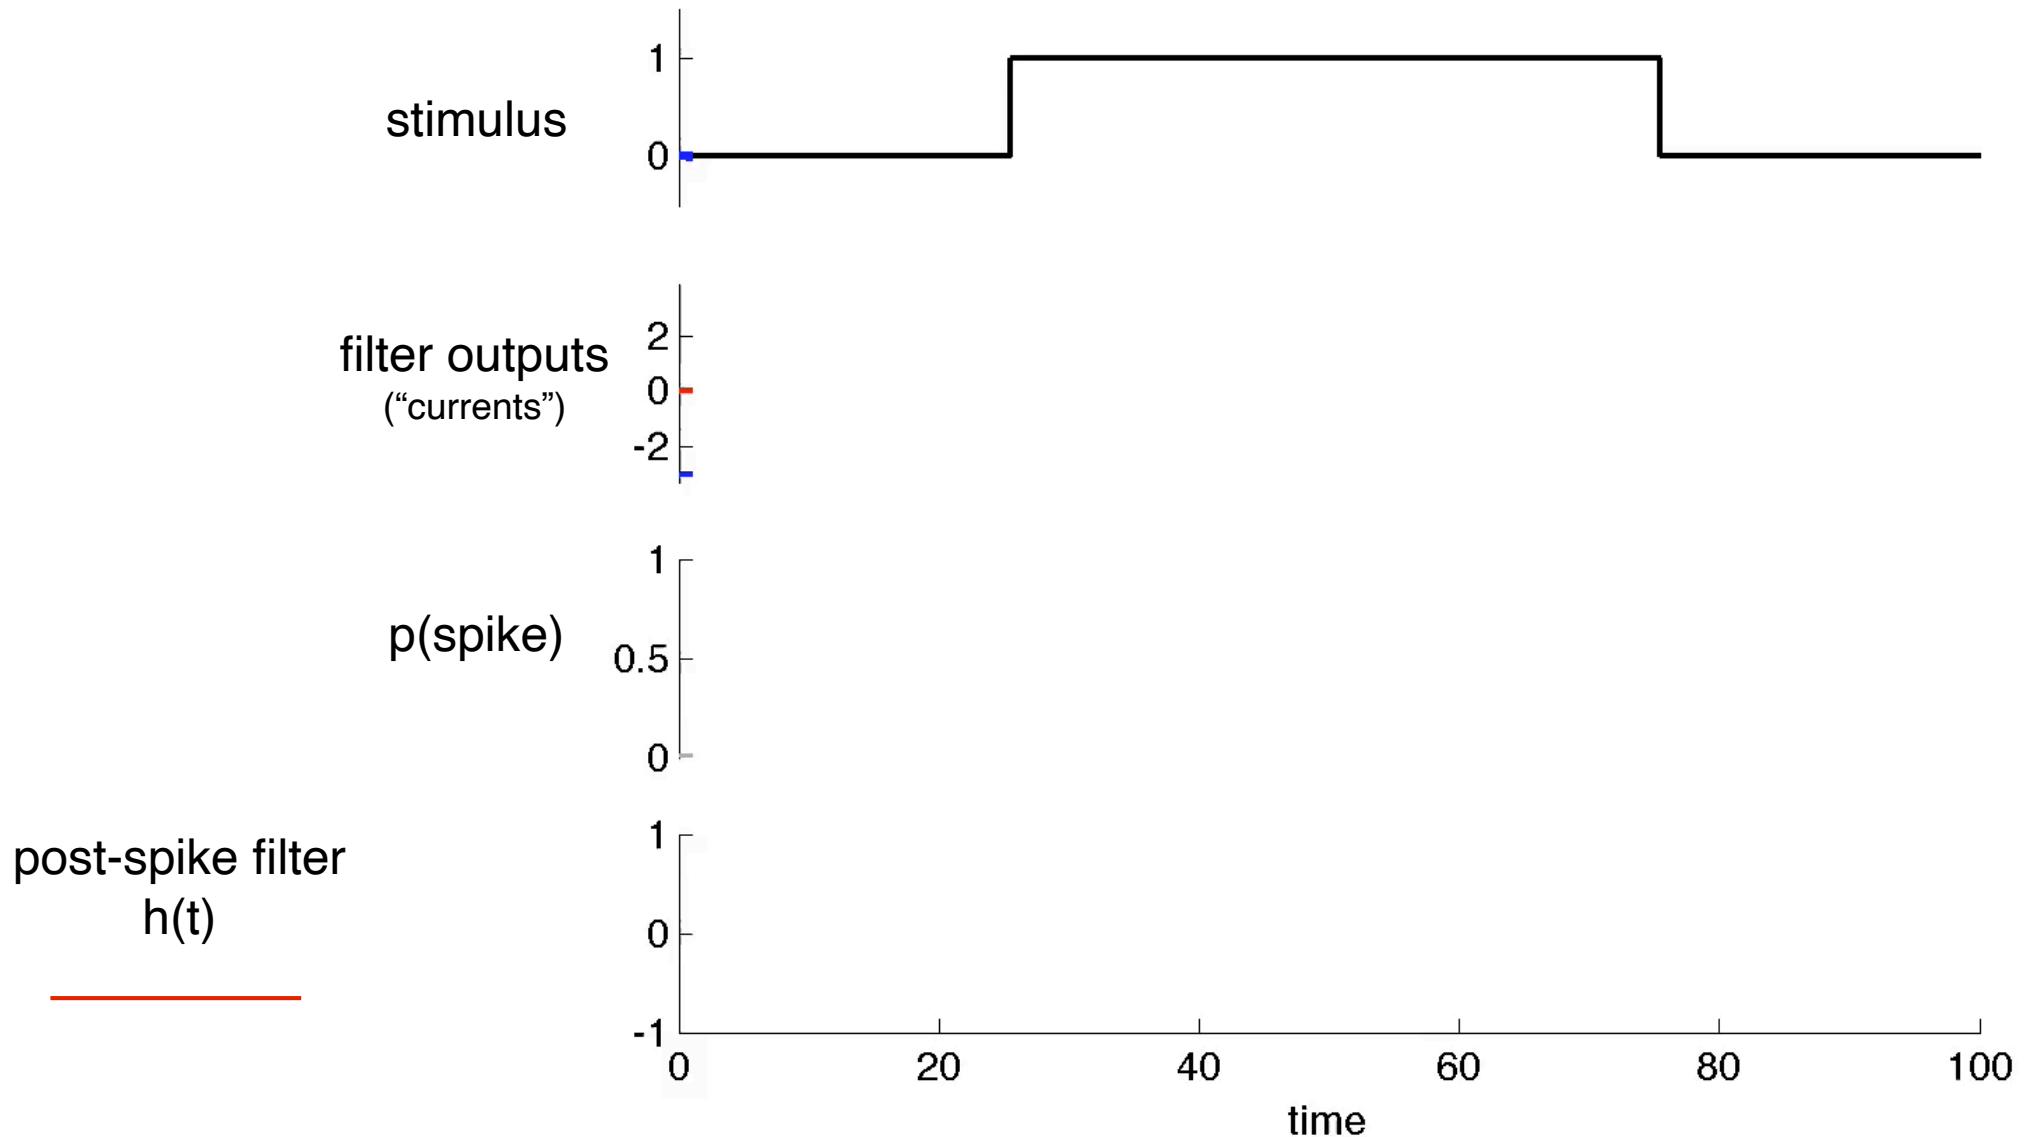


- interpretation: “soft-threshold” integrate-and-fire model



GLM dynamic behaviors

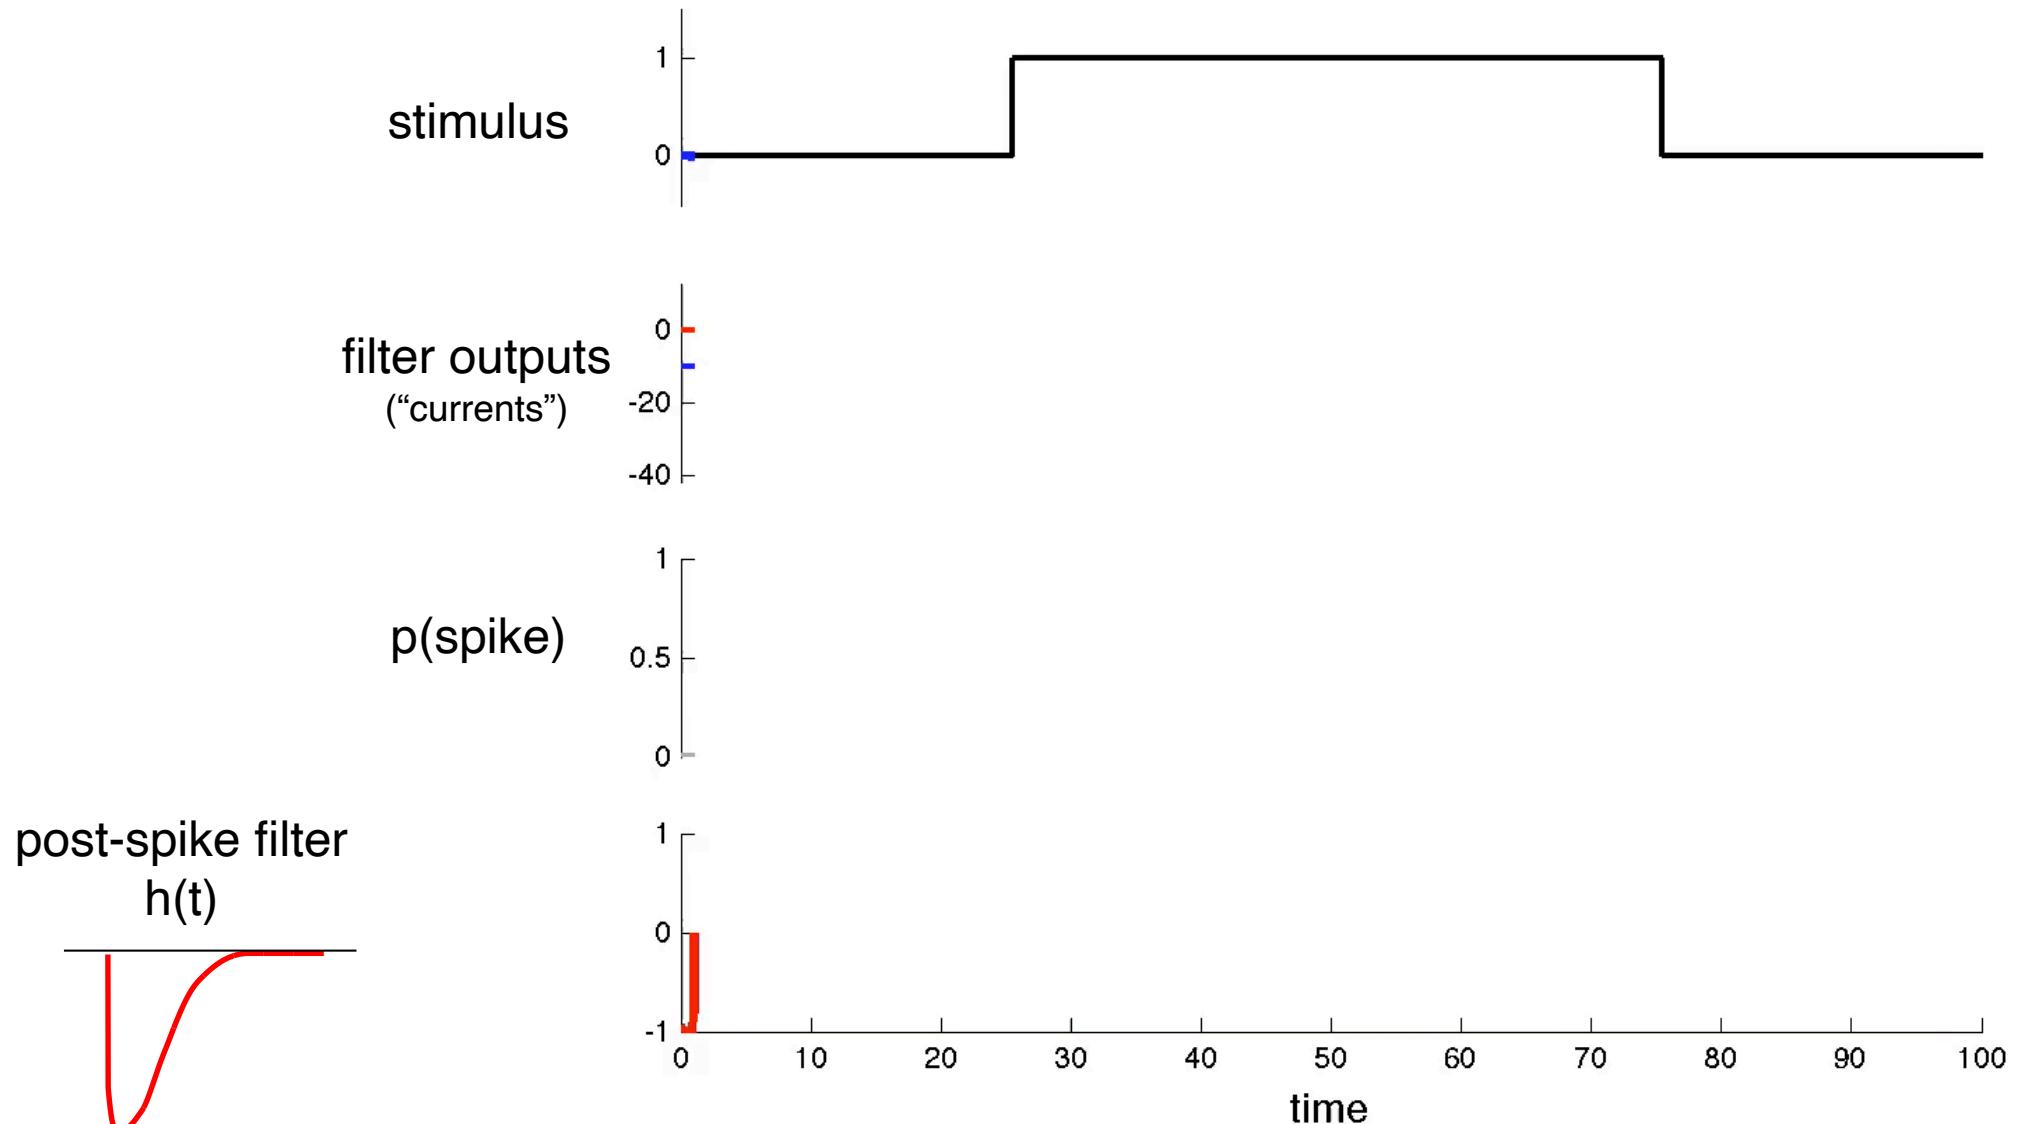
- irregular spiking



GLM dynamic behaviors

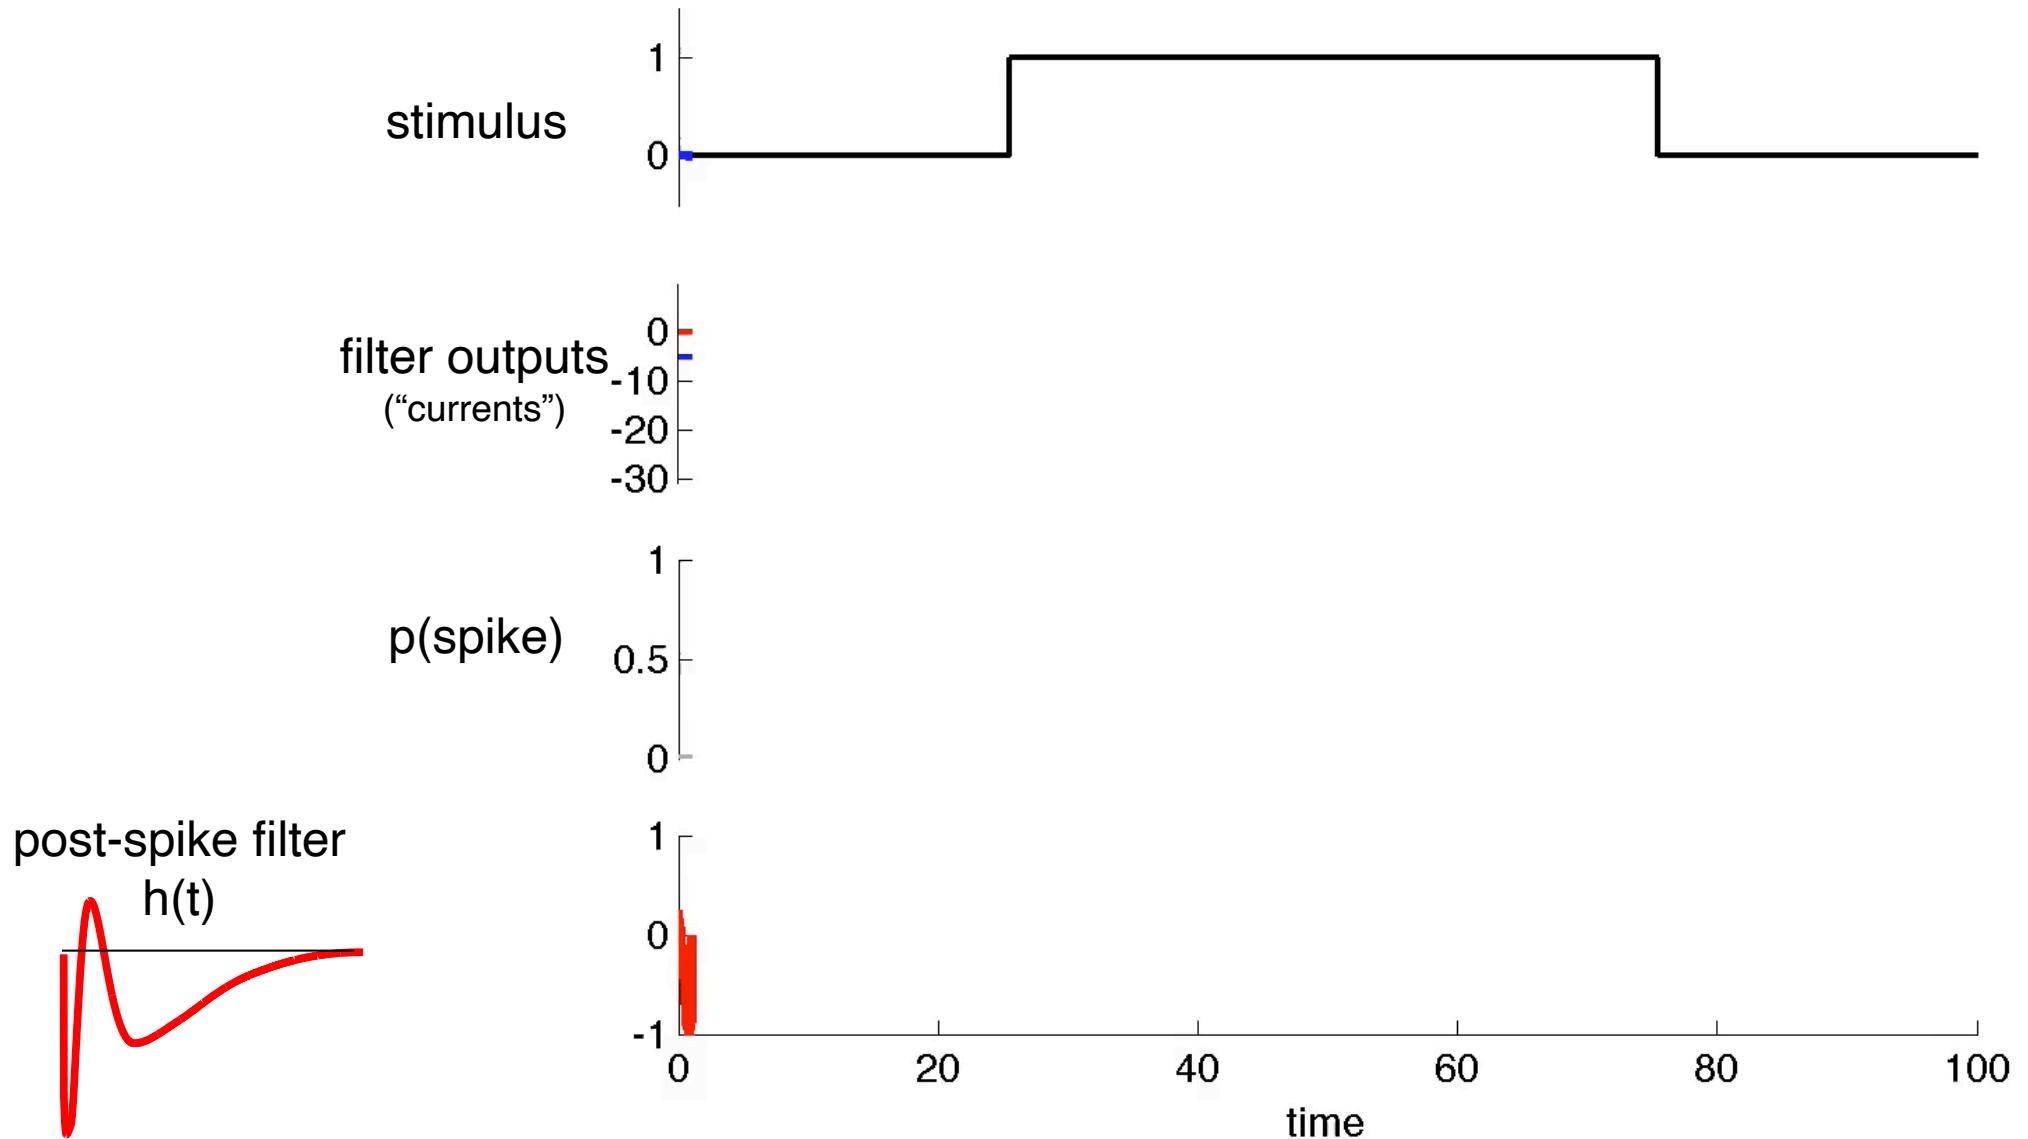
(Weber & Pillow 2016)

- regular spiking



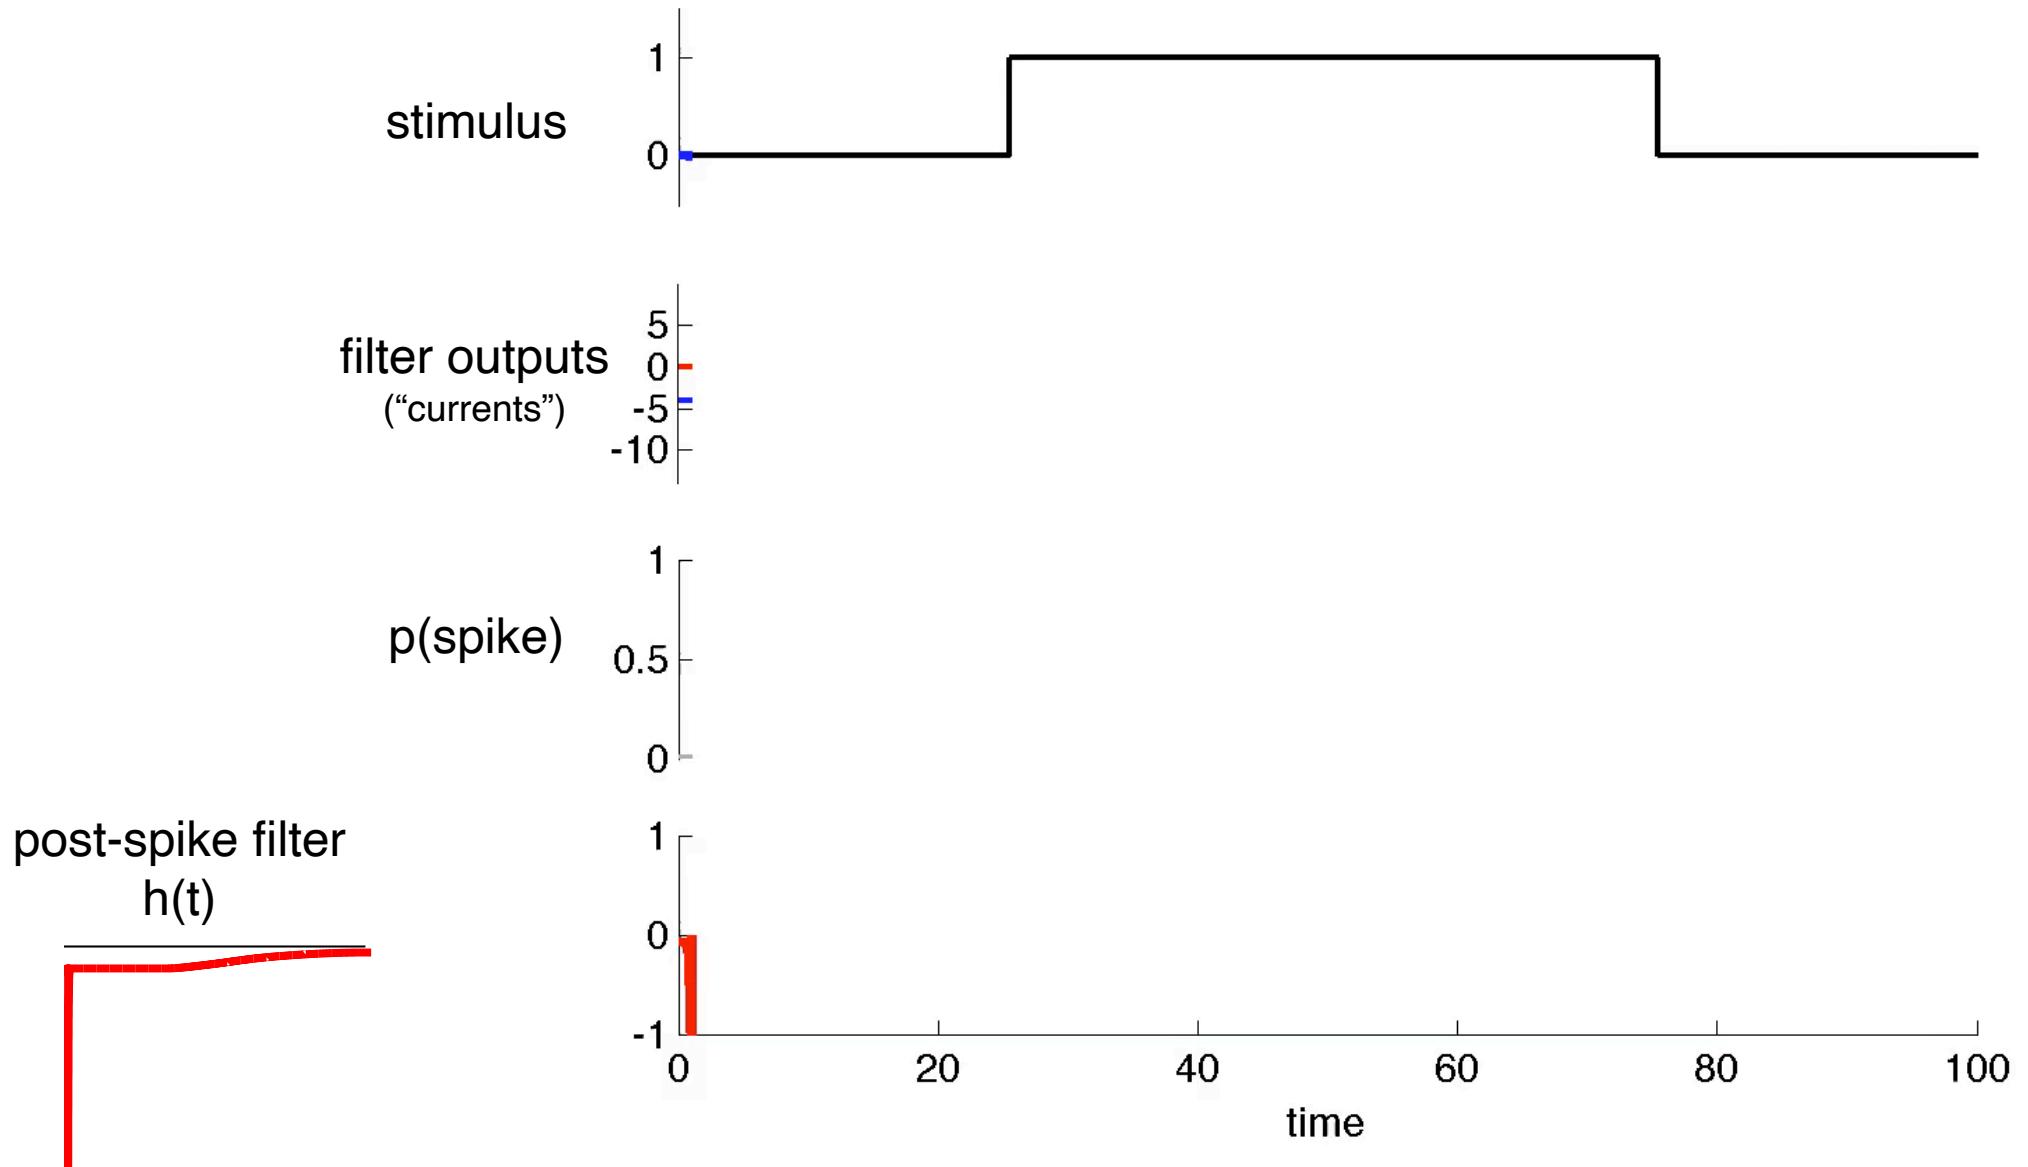
GLM dynamic behaviors

- bursting



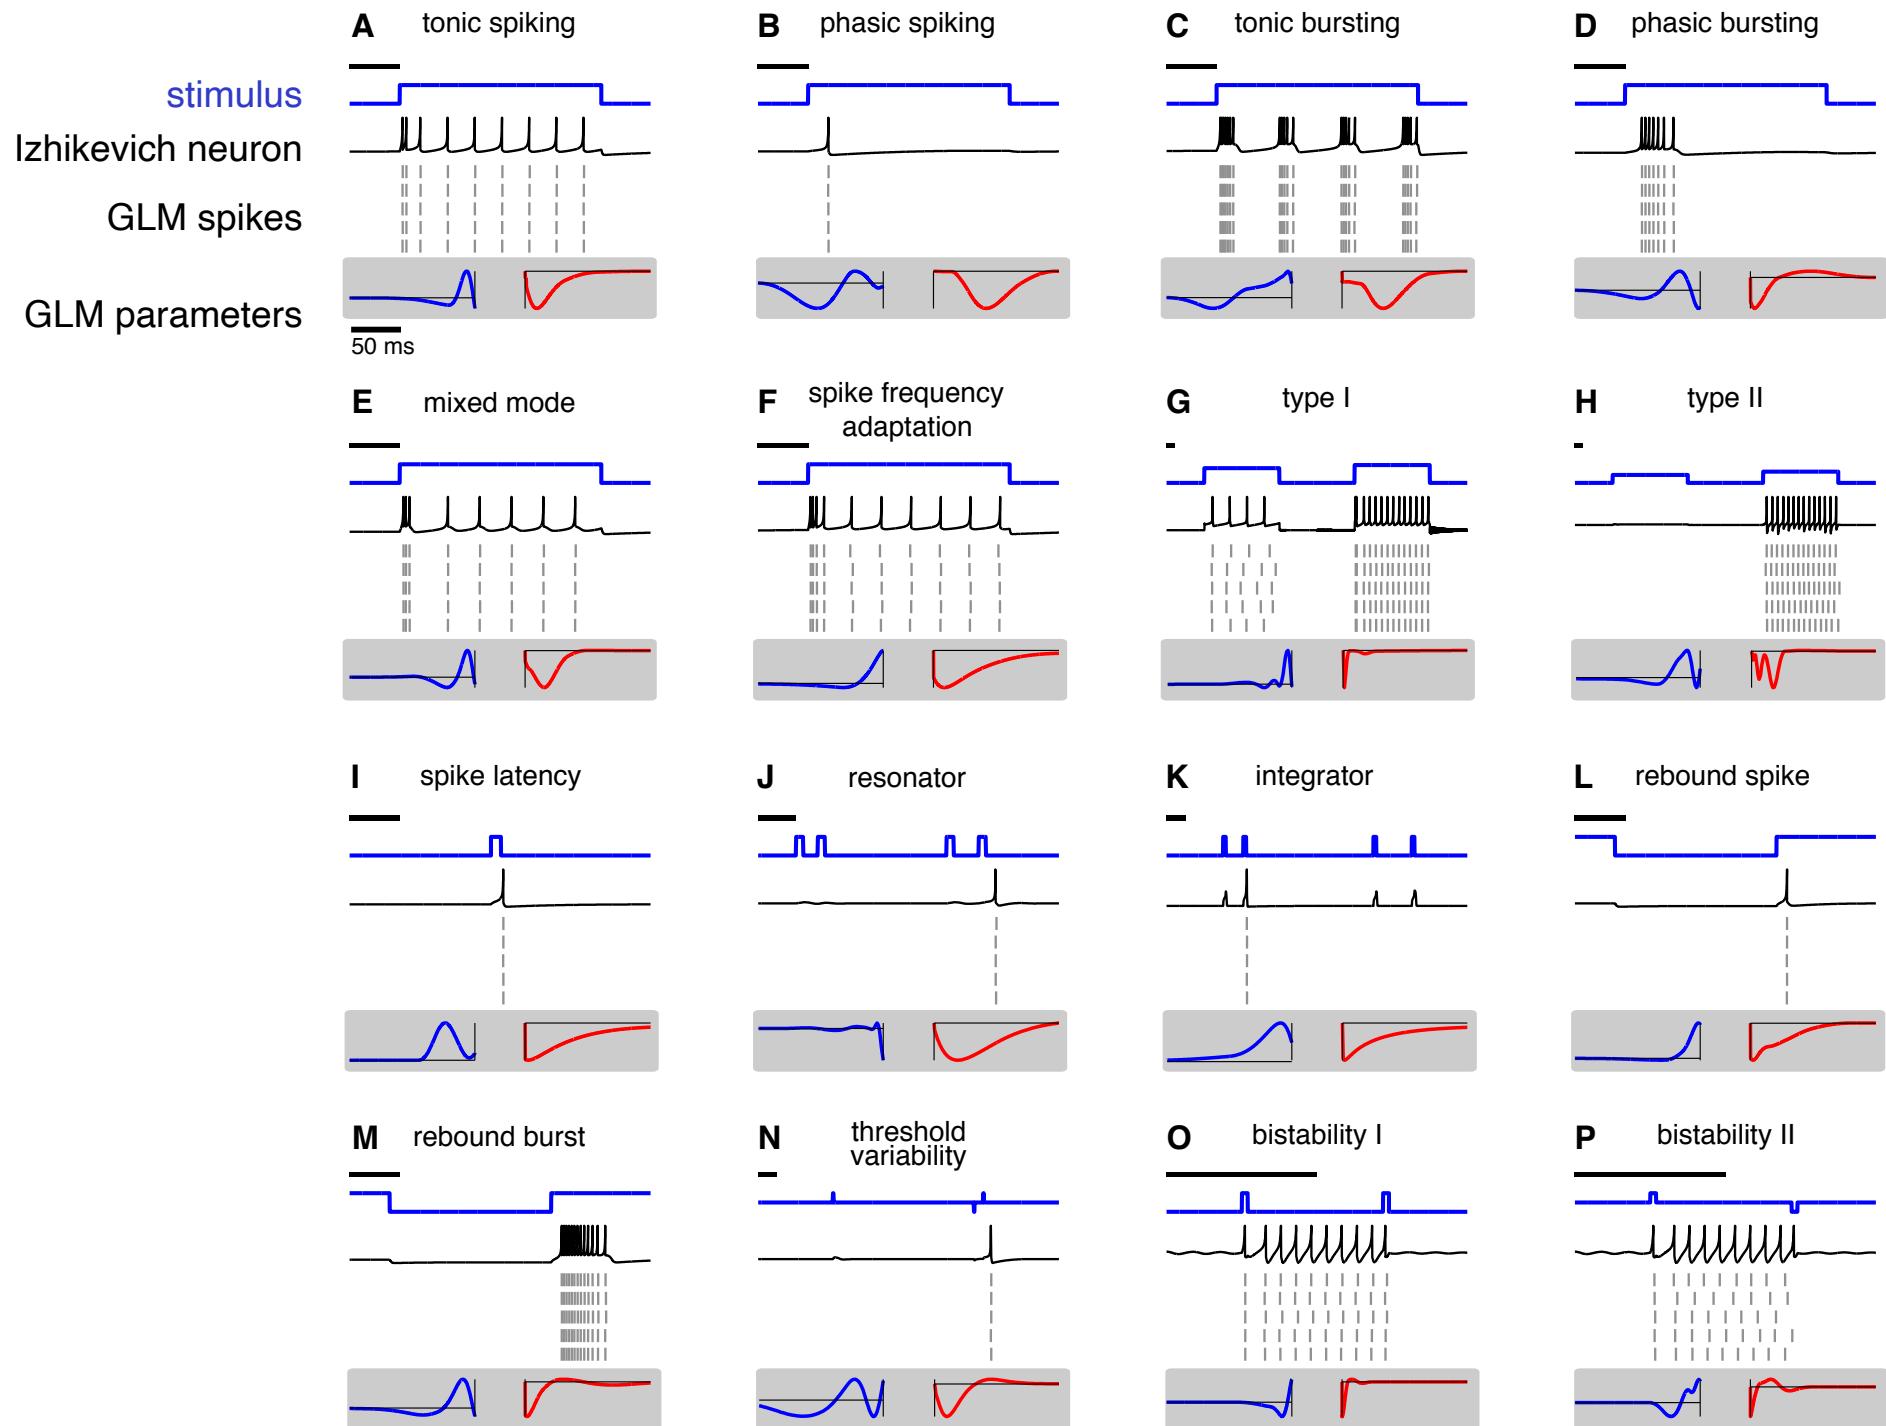
GLM dynamic behaviors

- adaptation

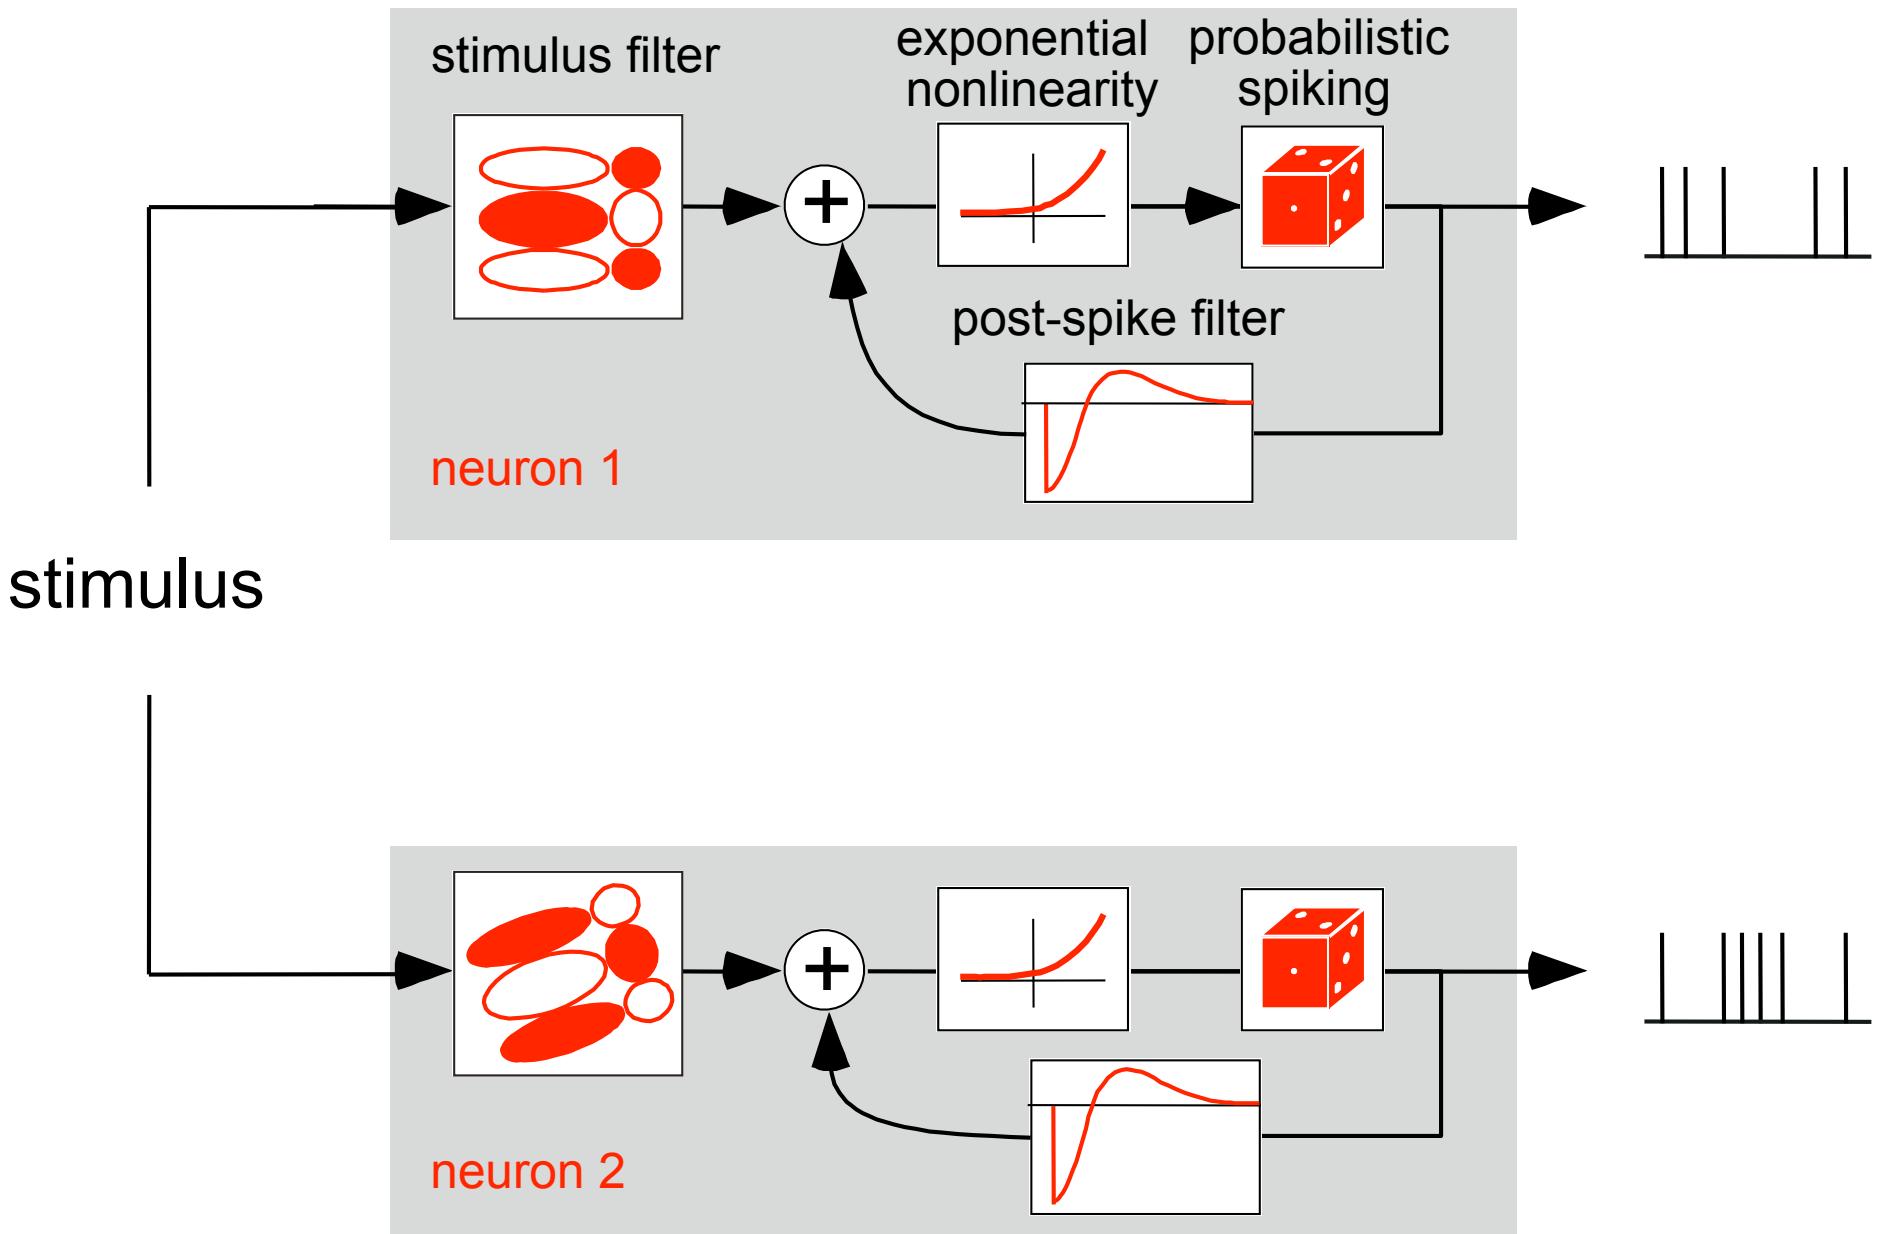


GLM dynamic behaviors (from Izhikevich)

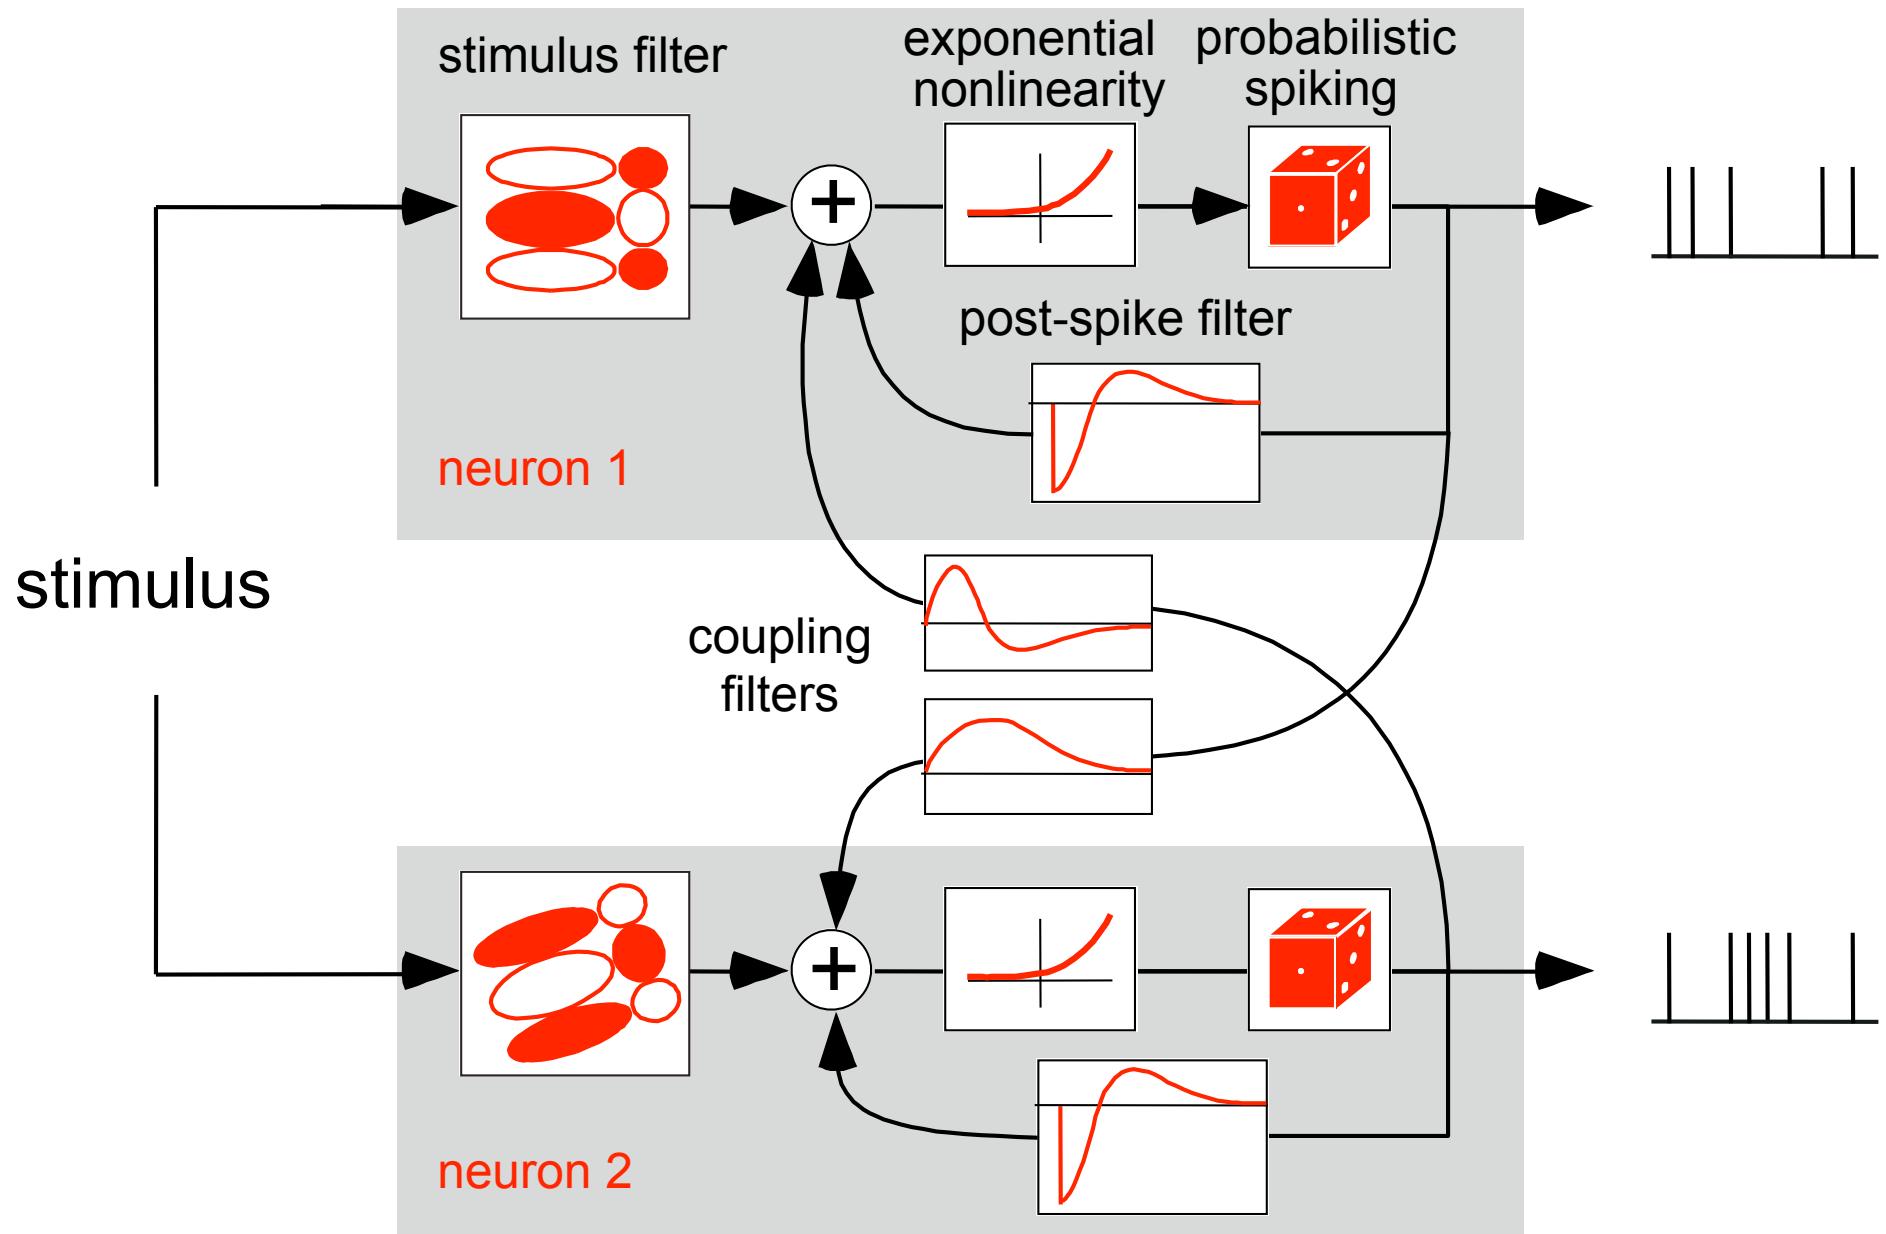
(Weber & Pillow 2017)



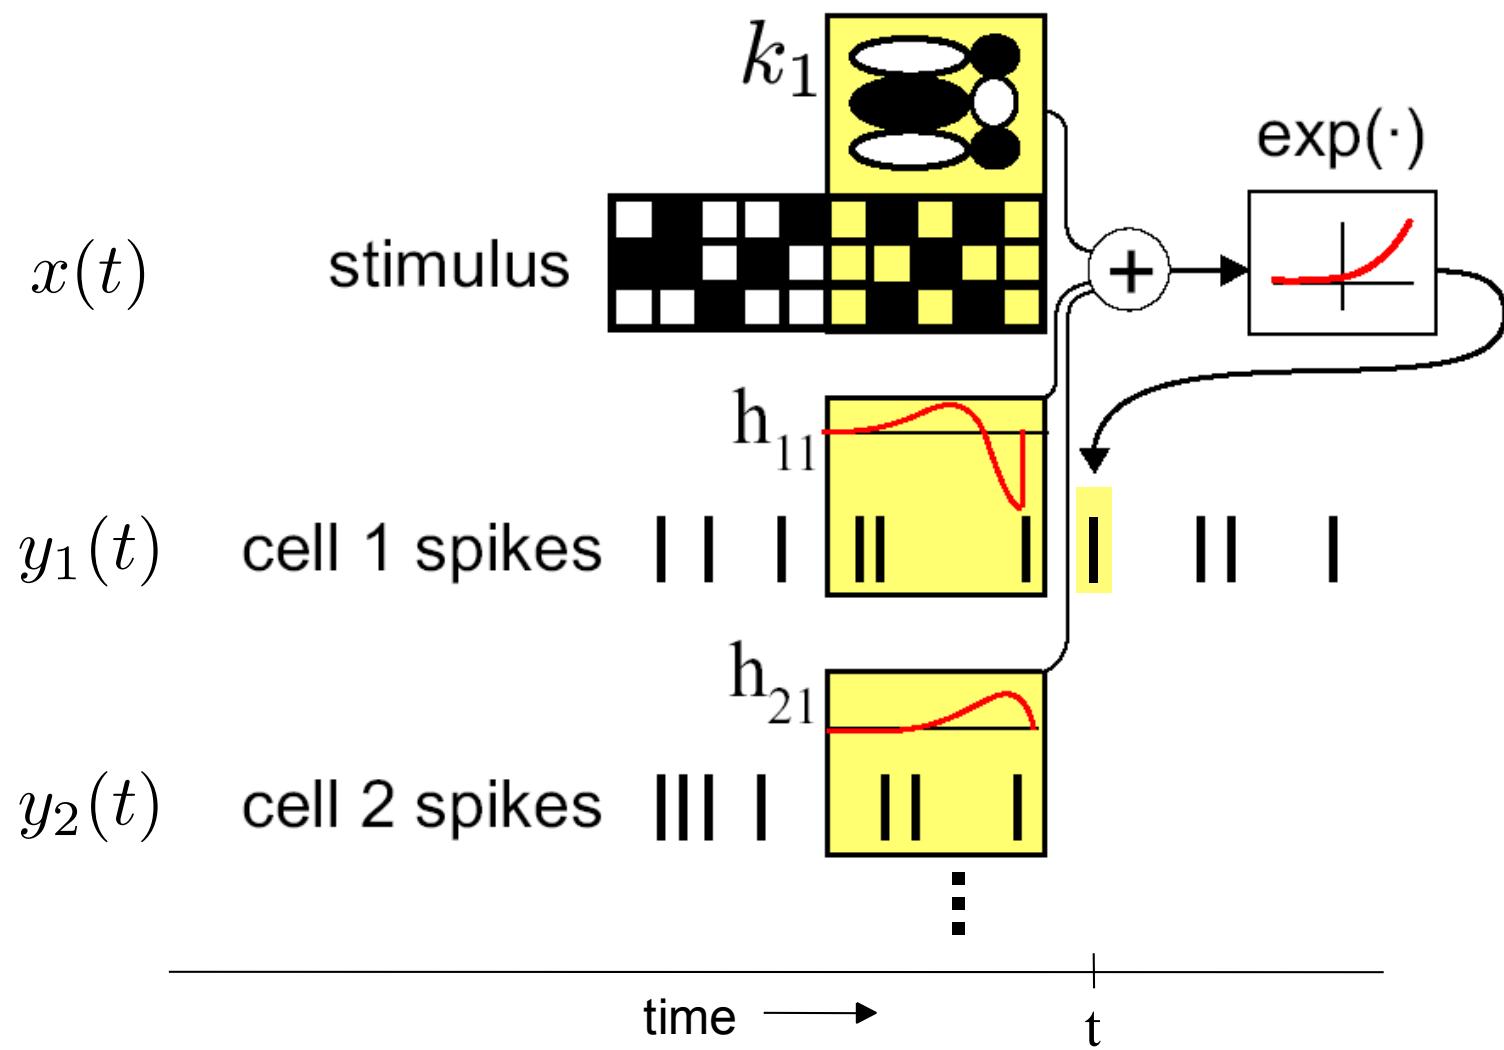
multi-neuron GLM



multi-neuron GLM



GLM equivalent diagram:



$$\text{spike rate} \quad \lambda_i(t) = \exp(k_i \cdot x(t) + \sum_j h_{ij} \cdot y(t))$$

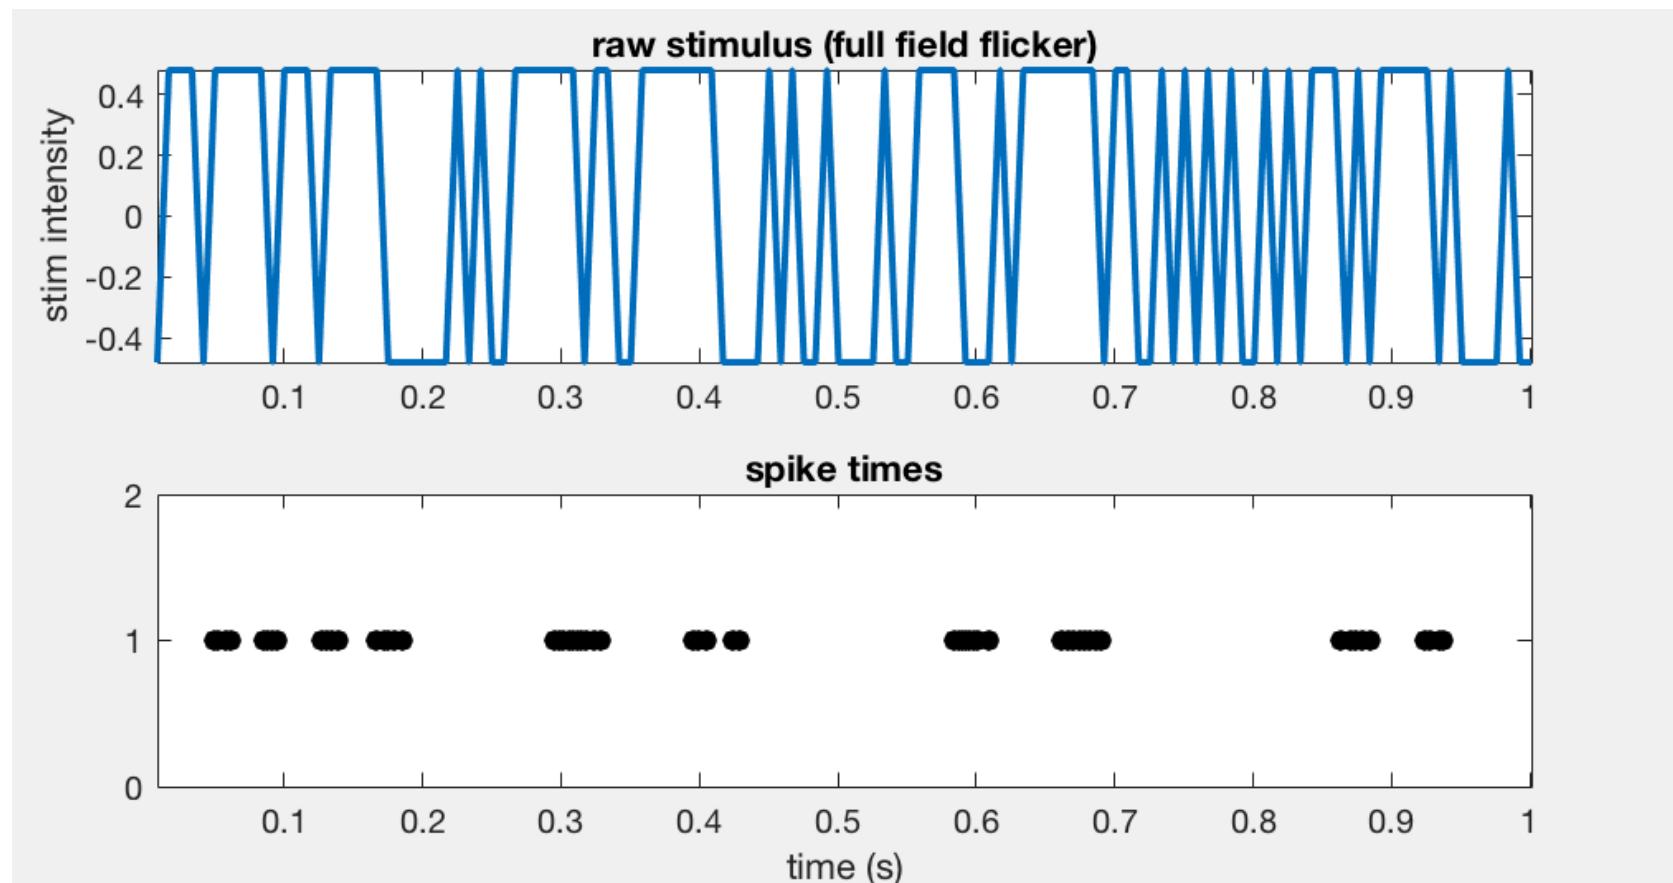
Do it: tutorial2_spikehistscoupledGLM.m

Topics:

- 1-2. load data & bin spike trains
3. build design matrix, now including spike history!
4. fit Poisson GLM with spike-history
5. fit coupled Poisson GLM (4 neuron spike-history)
6. model comparison: log-likelihood & AIC

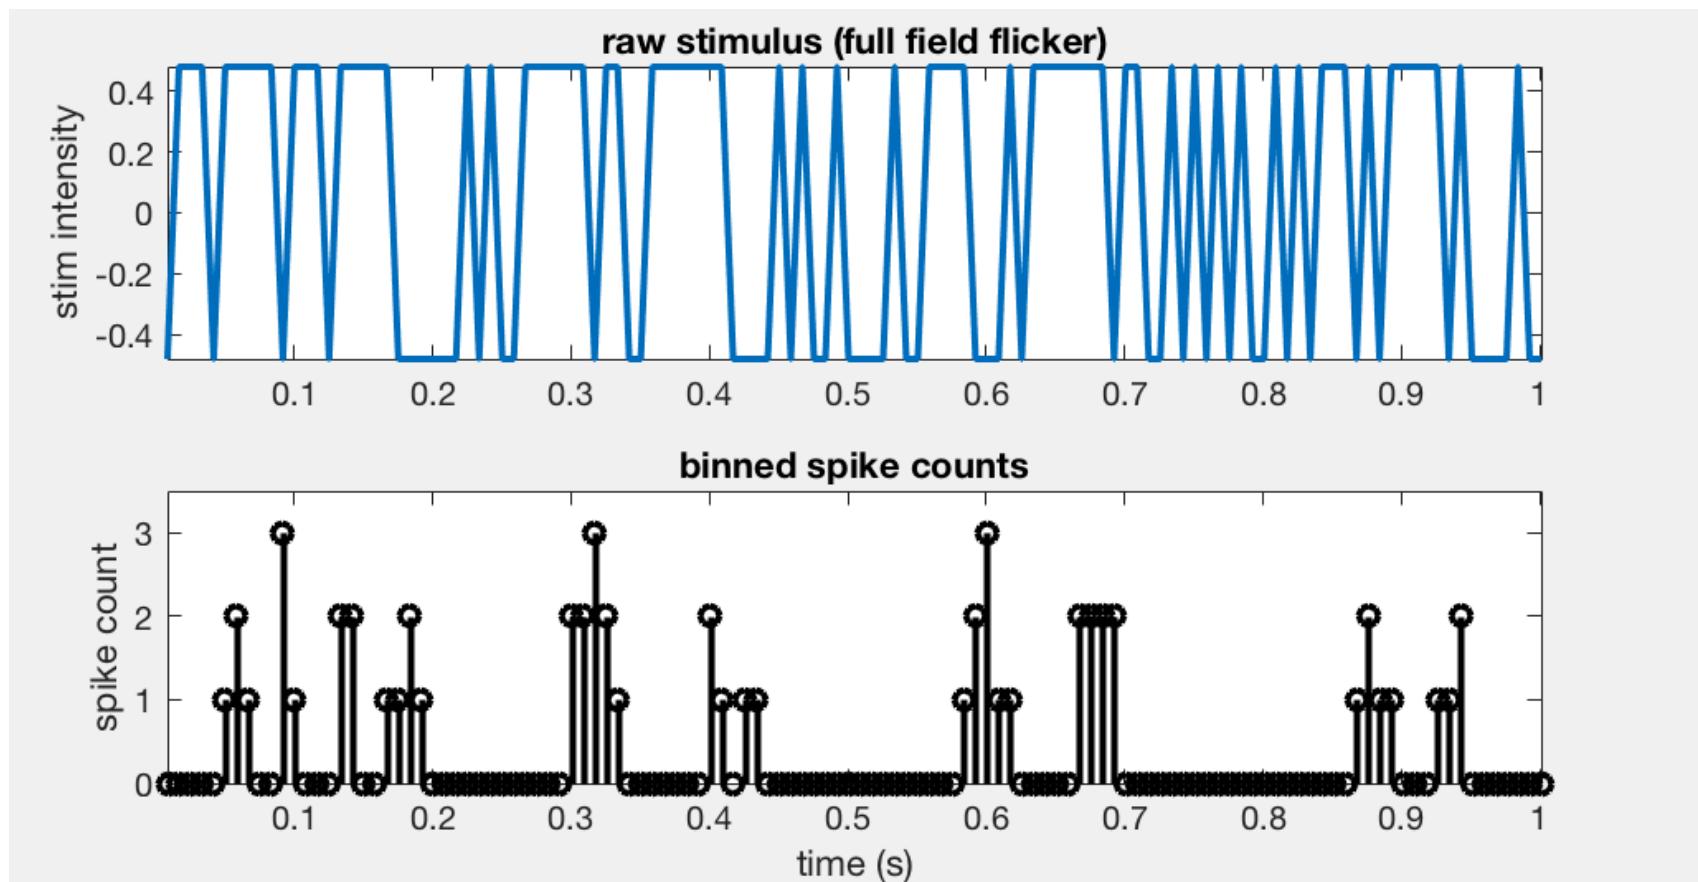
Example dataset

- stimulus = binary flicker
- parasol retinal ganglion cell spike responses



Example dataset

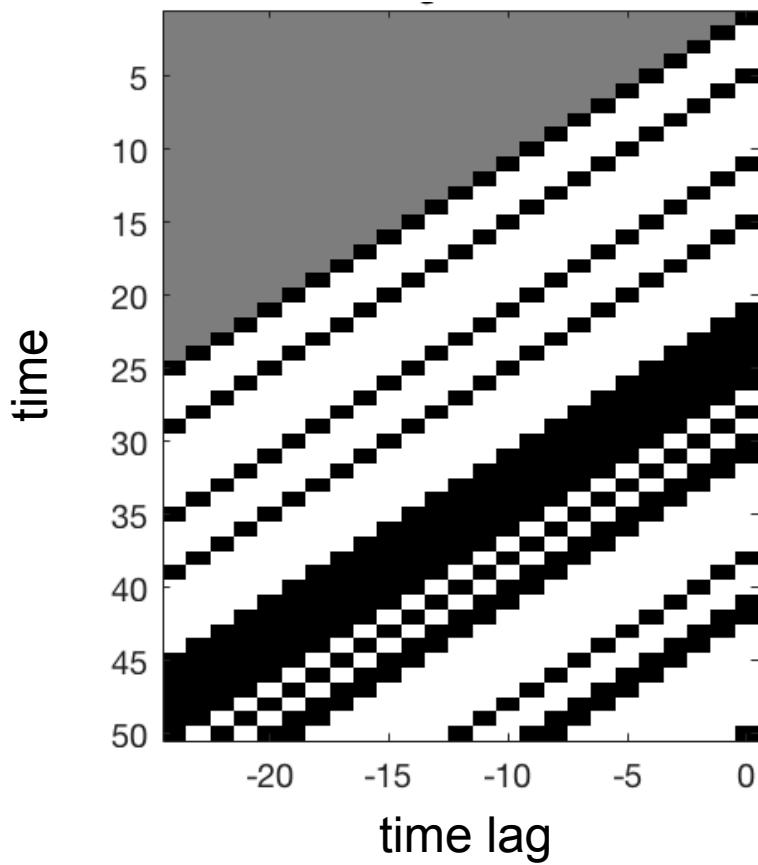
- stimulus = binary flicker
- parasol retinal ganglion cell spike responses



Stimulus-only GLM

design matrix

X



spike response

Y

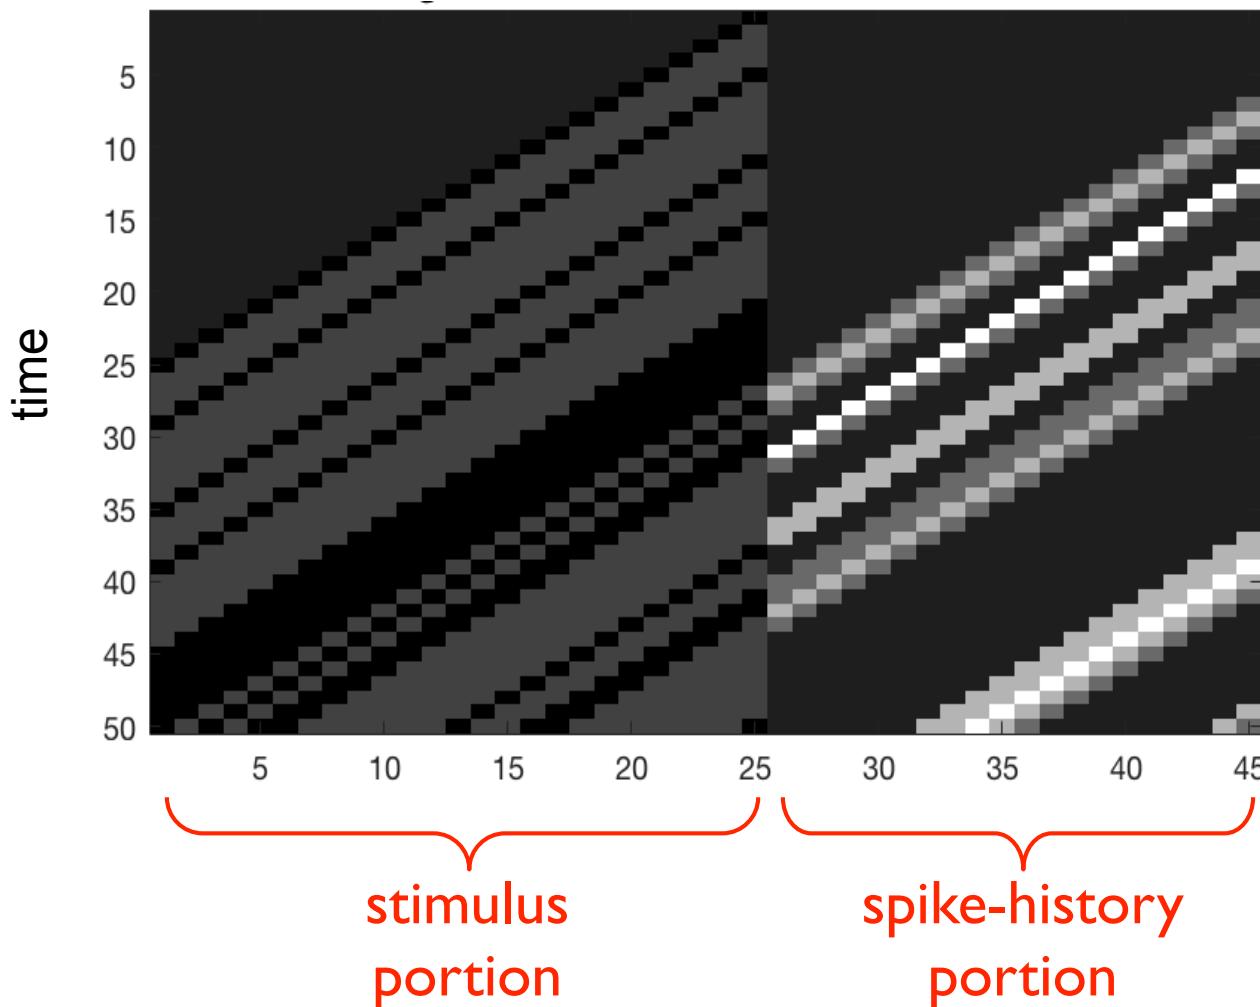
model
 $P(Y|X)$



Stimulus + SpikeHistory GLM

design matrix

X



spike response

Y

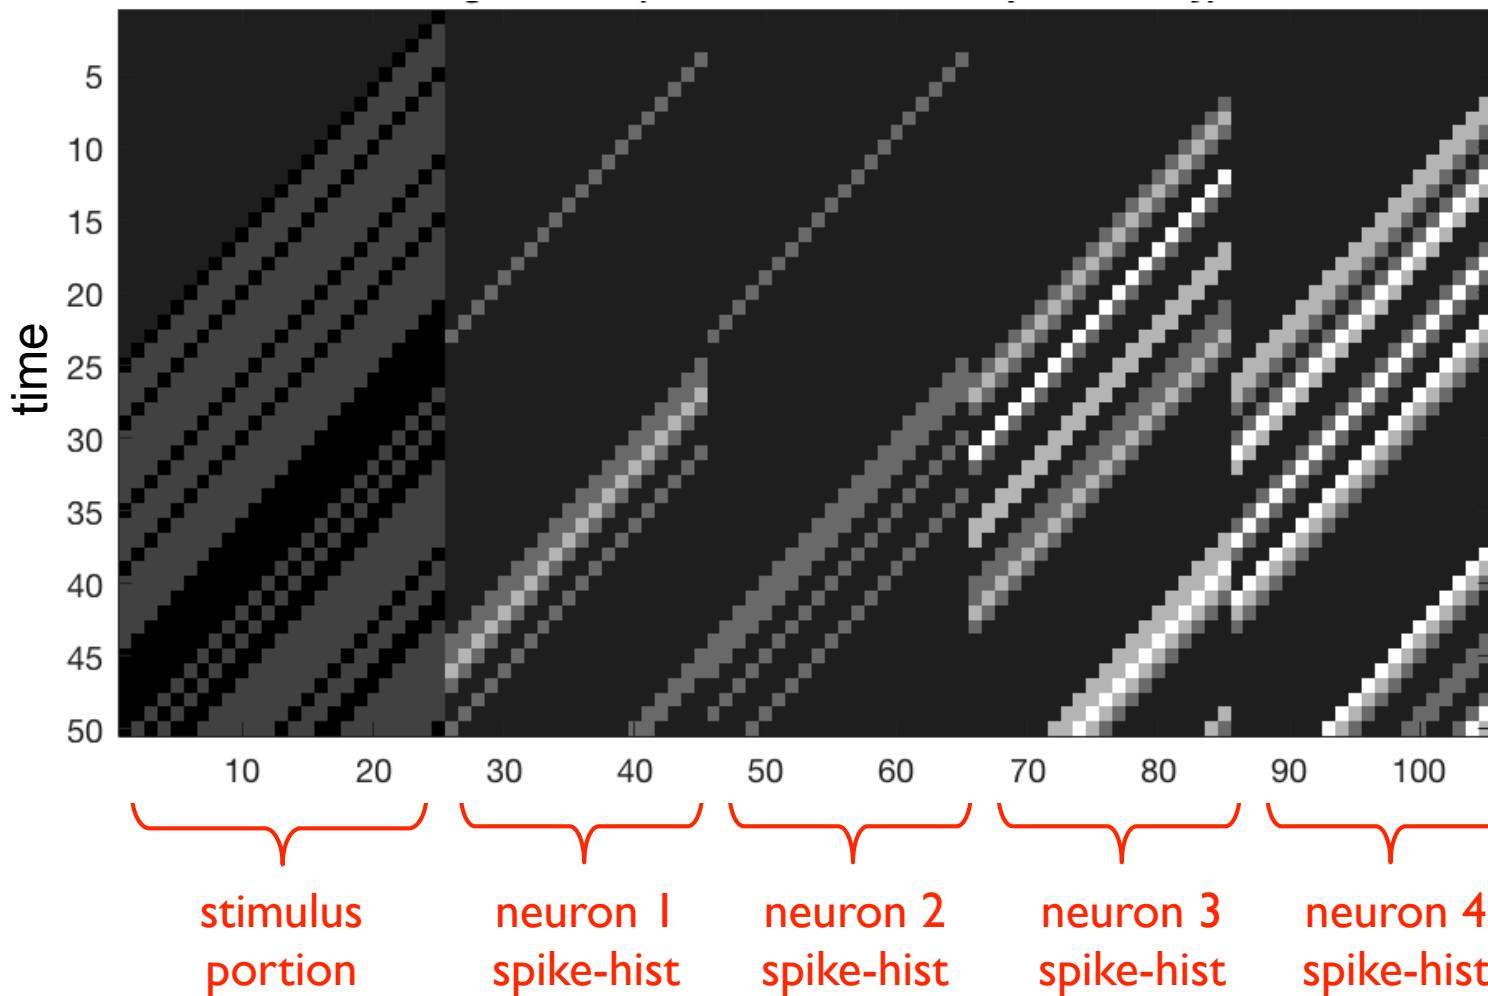
model
 $P(Y|X)$



Stimulus + History + 3 Neuron Coupling GLM

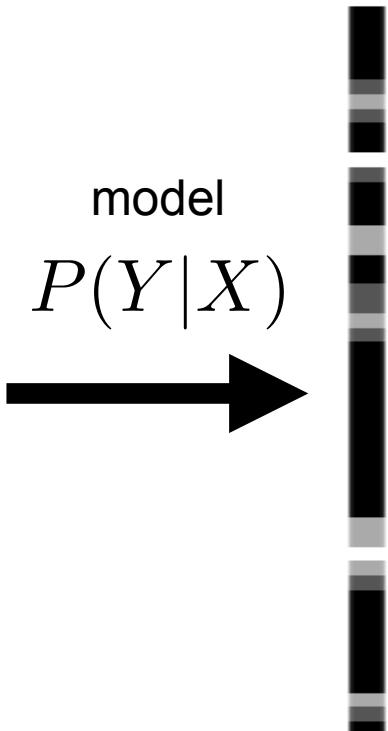
design matrix

X



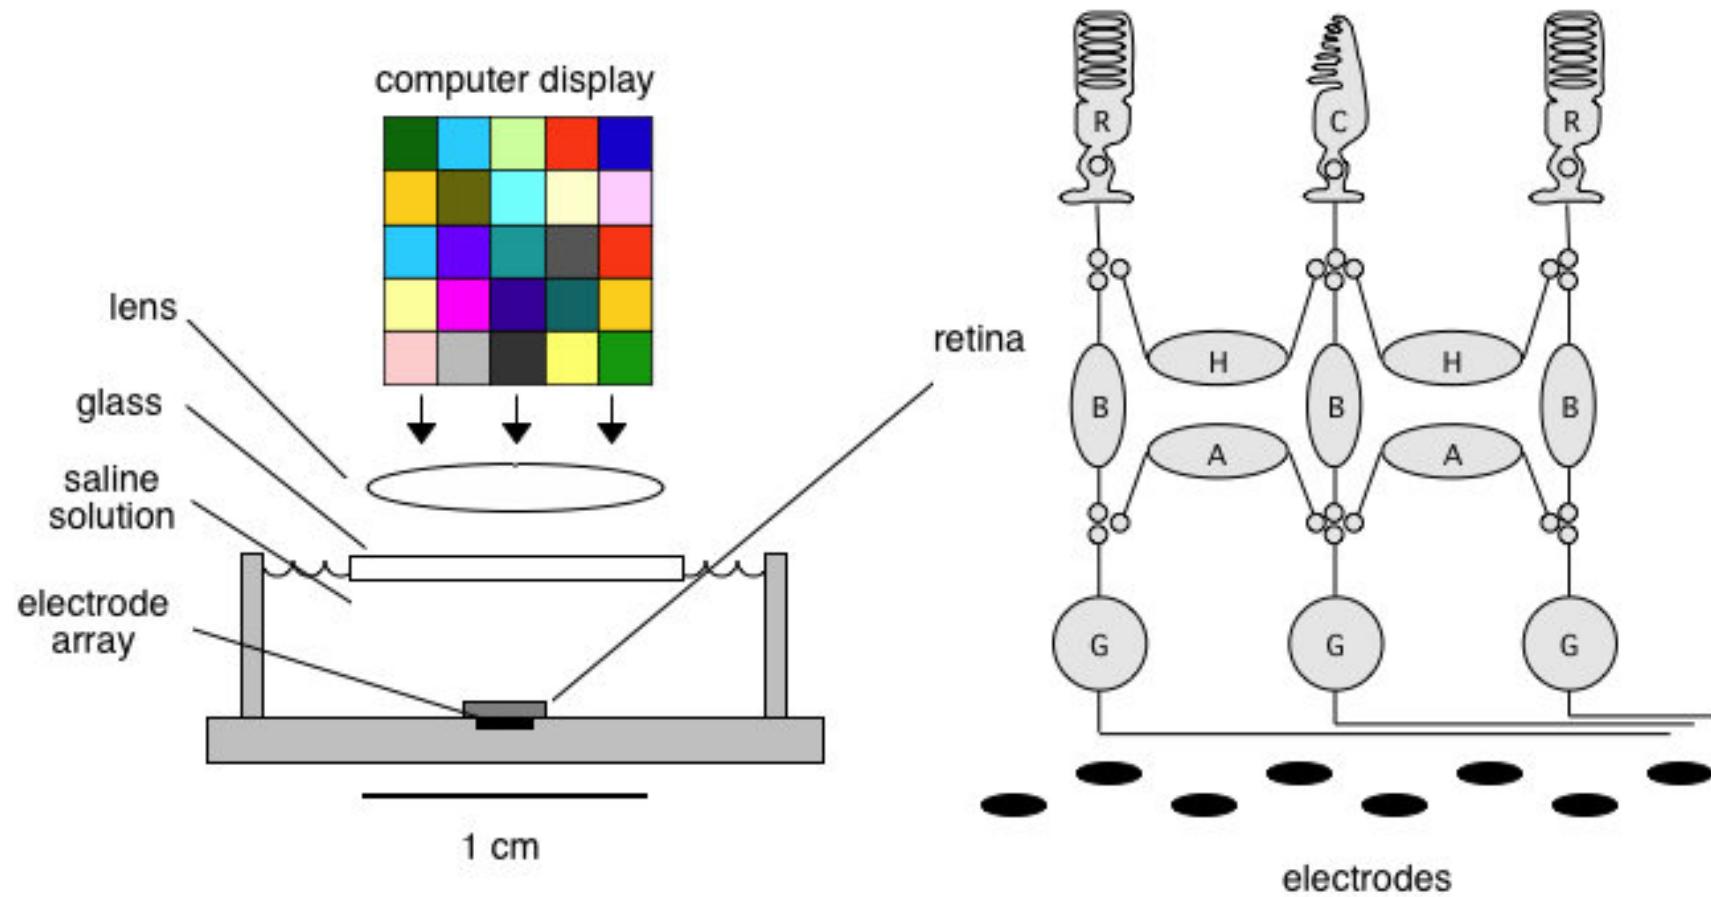
spike response

Y



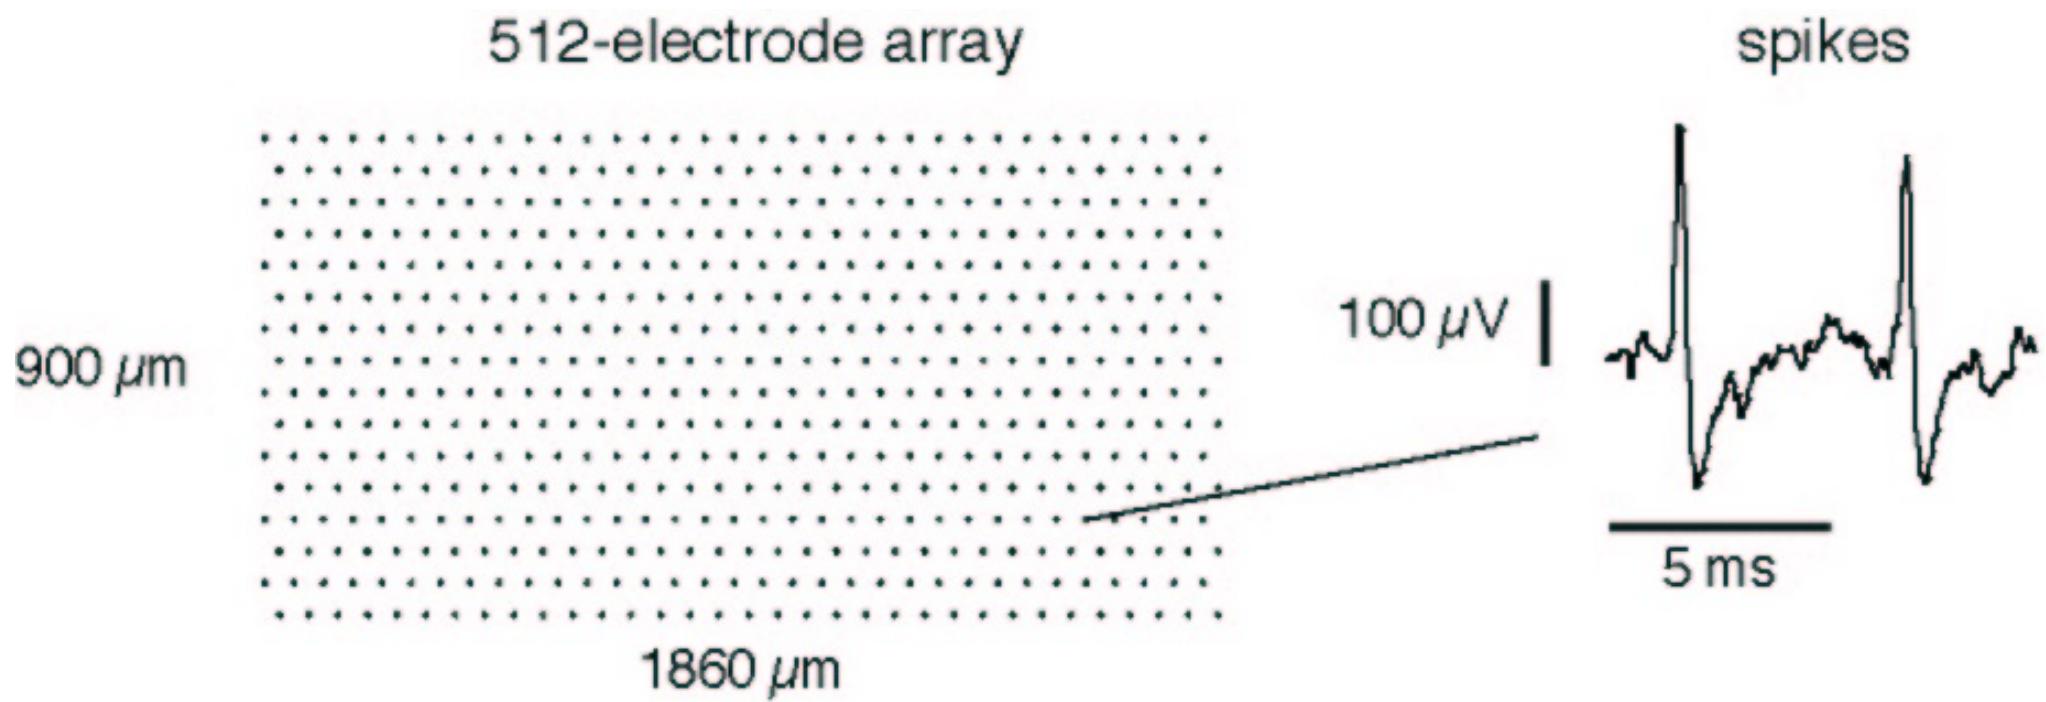
Application of GLM to Retinal Data

Application to Data: primate retina



Data: Chichilnisky Lab, The Salk Institute

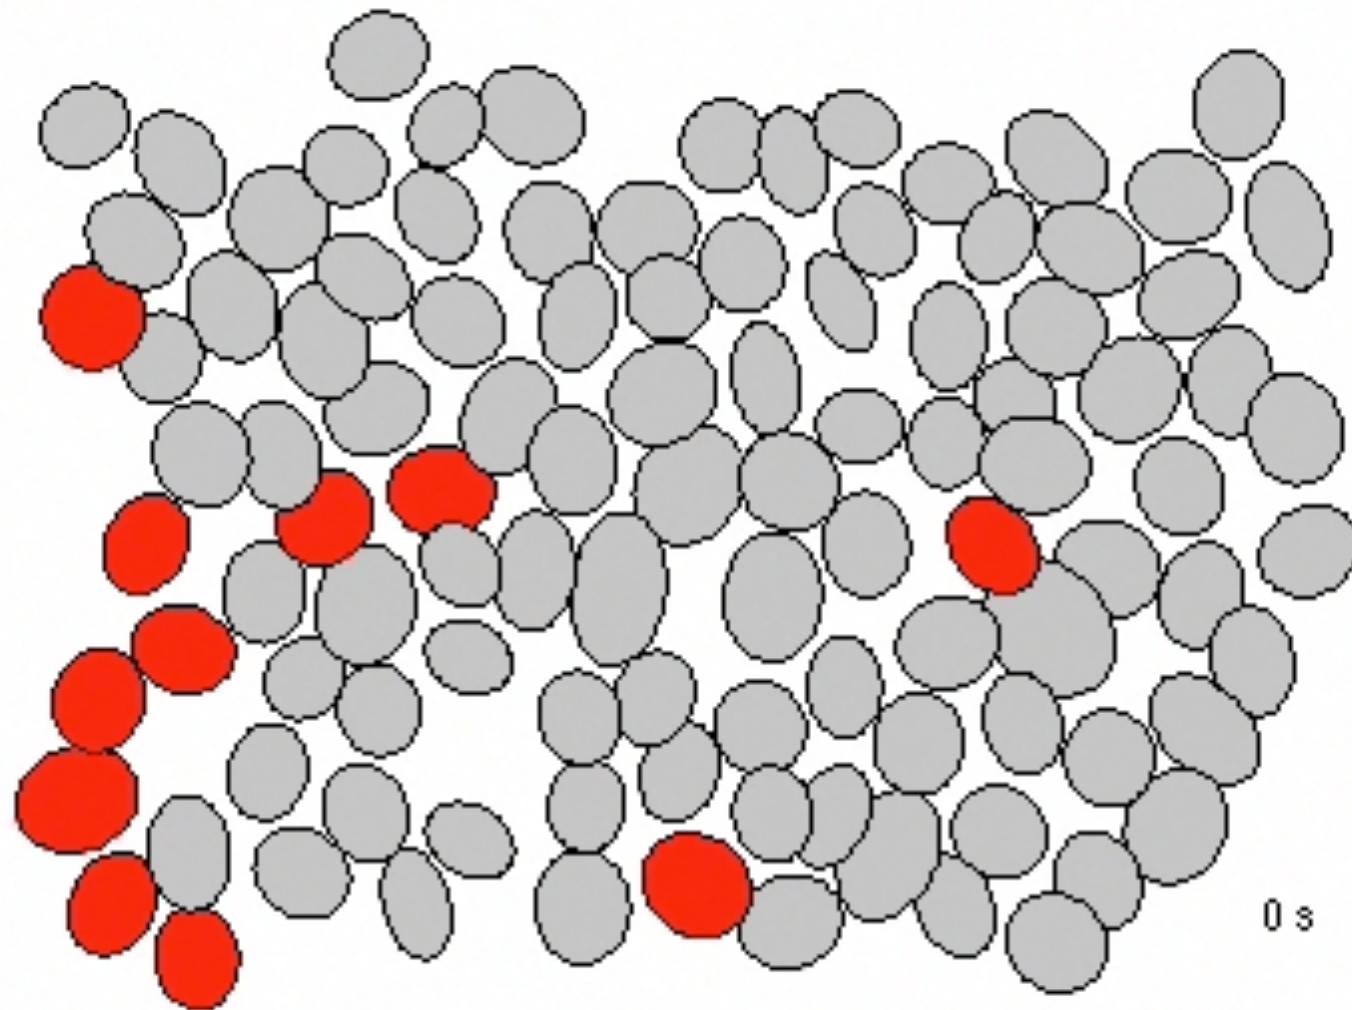
Application to Data: primate retina



Data: Chichilnisky Lab, The Salk Institute

Retinal responses to white noise

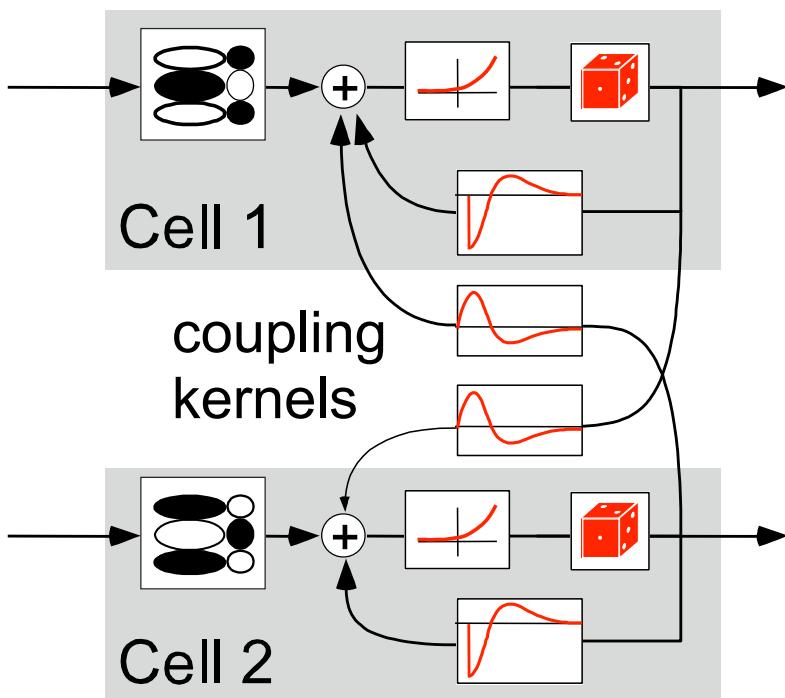
(ON parasol
cells)



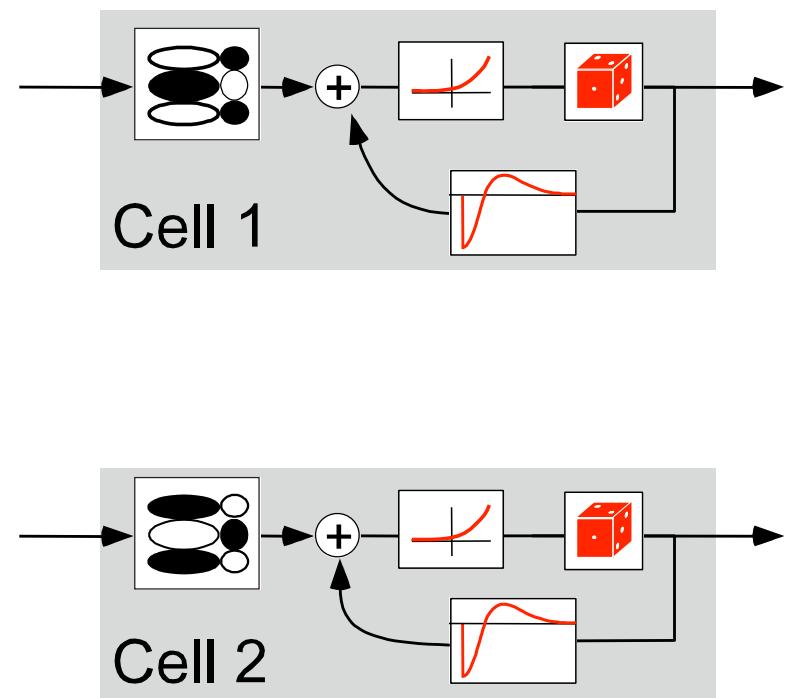
Shlens, Field, Gauthier, Greschner, Sher , Litke & Chichilnisky (2009).

Approach: examine two versions of the GLM:

“full” model



“uncoupled” model



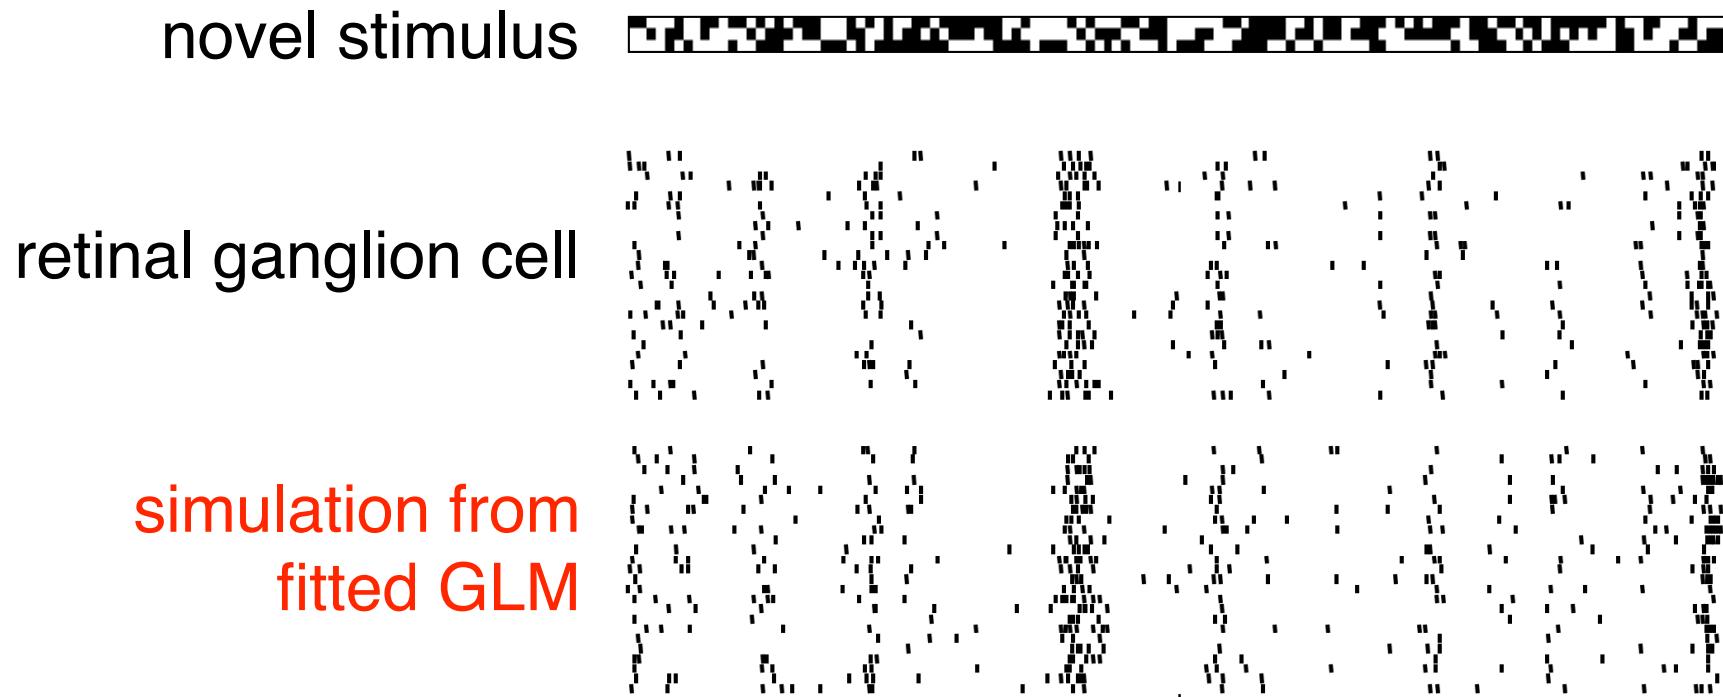
$$P(y_1, y_2, \dots, y_n | x)$$

joint encoding

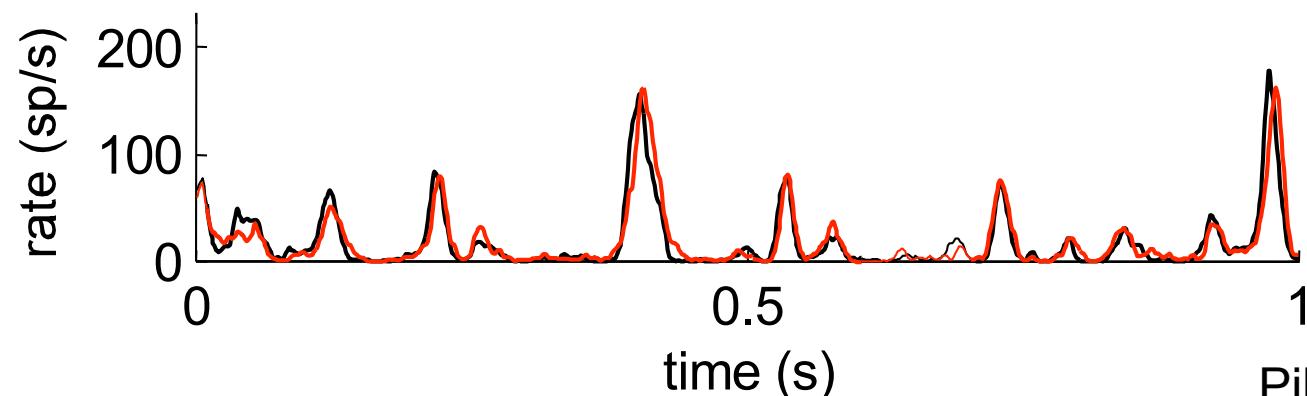
$$P(y_1|x)P(y_2|x) \cdots P(y_n|x)$$

independent encoding

Predicting single-cell responses: $P(y_i|x)$

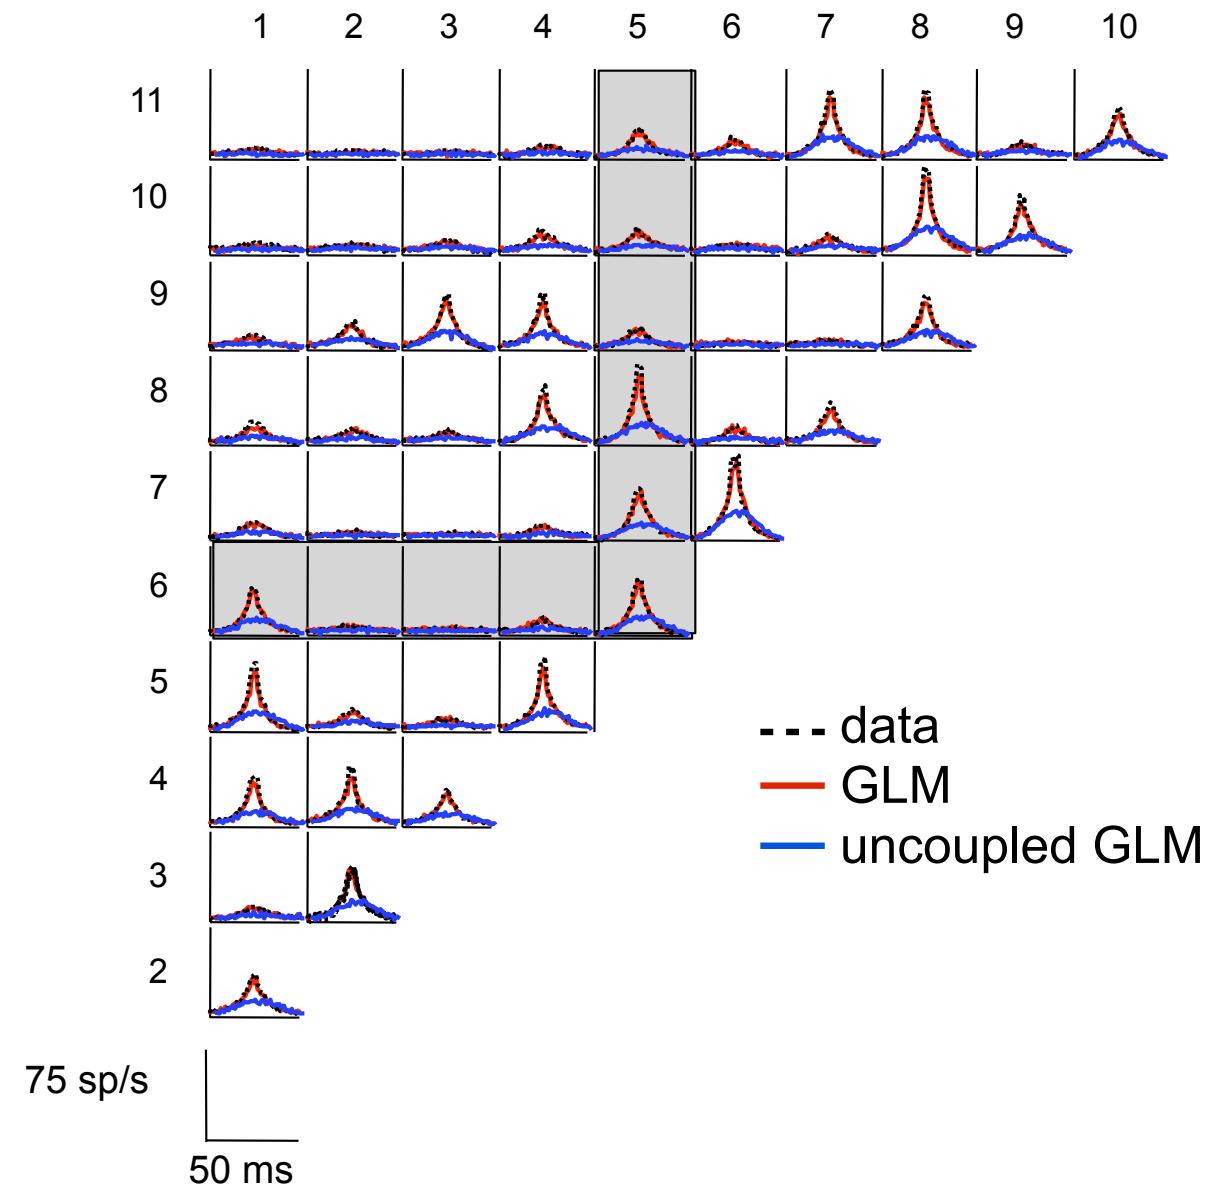
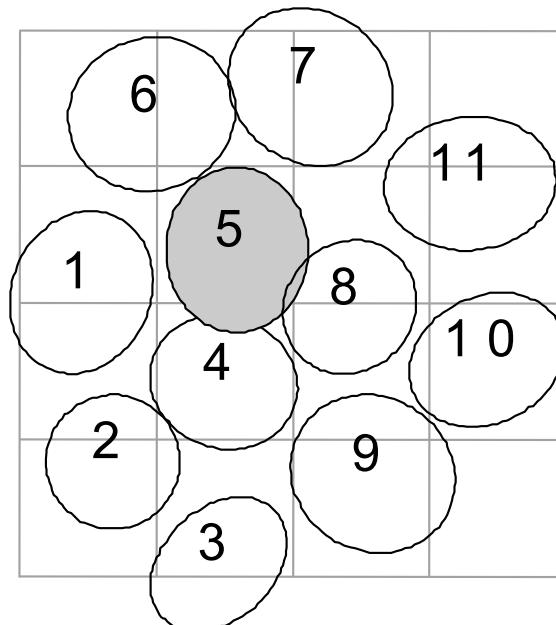


mean: psth (peristimulus time histogram)



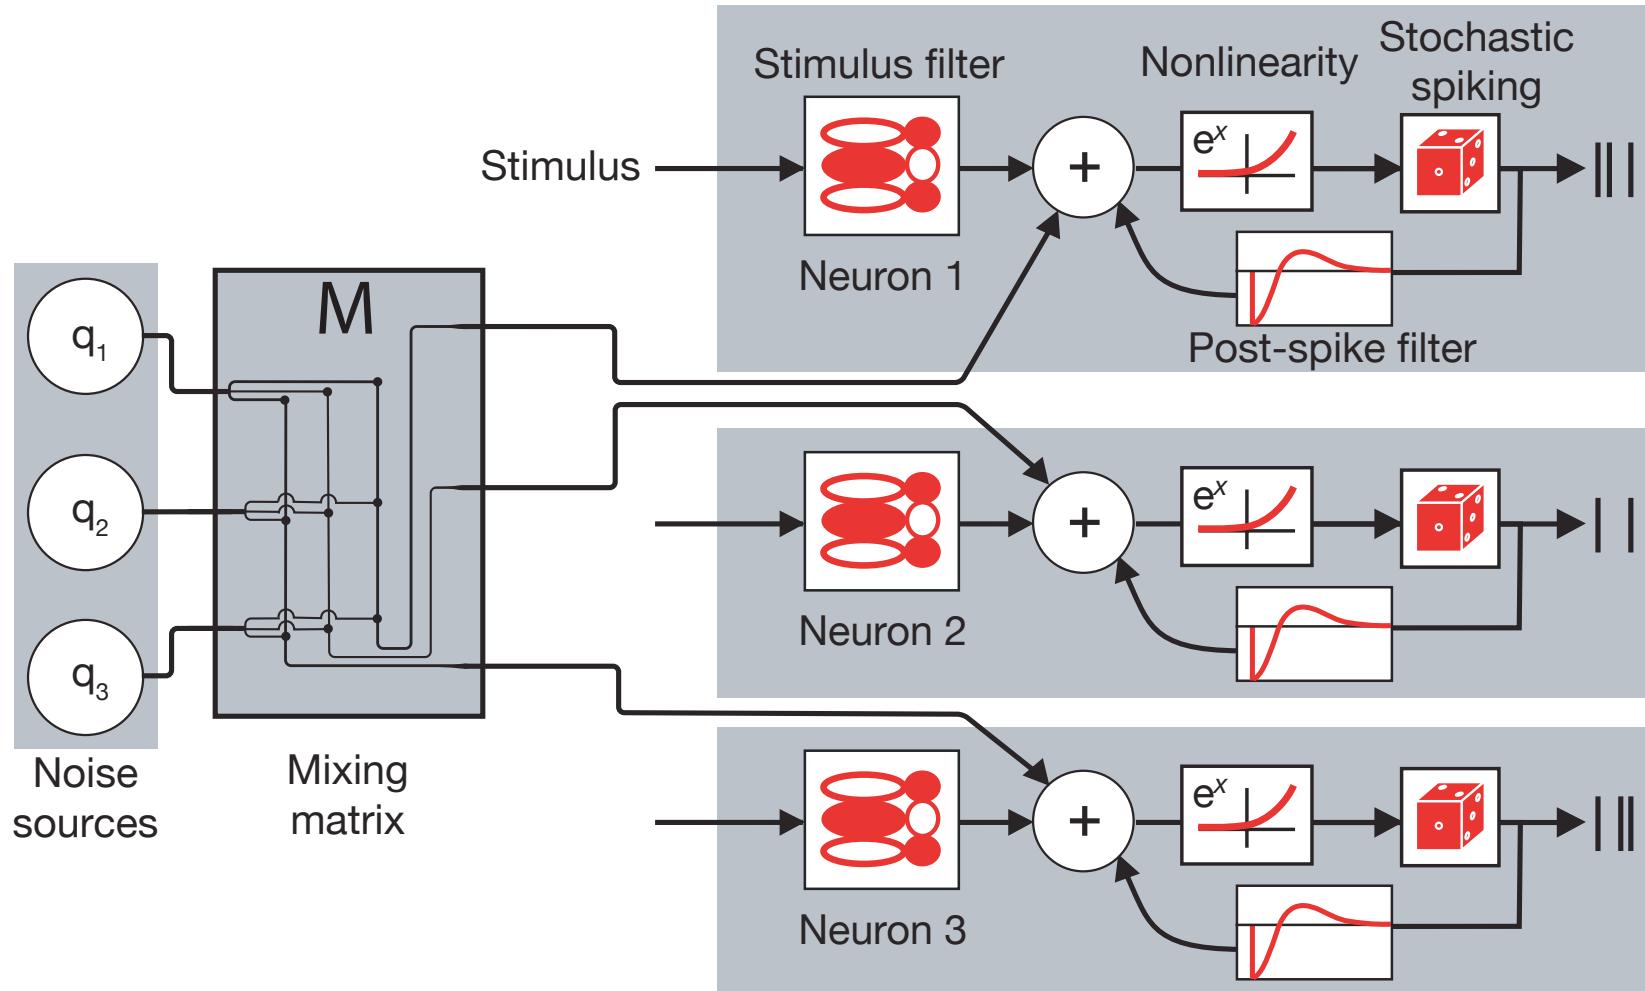
modeling correlation structure in neural spike trains

receptive fields



Beyond GLM

“shared noise” (latent variable) models

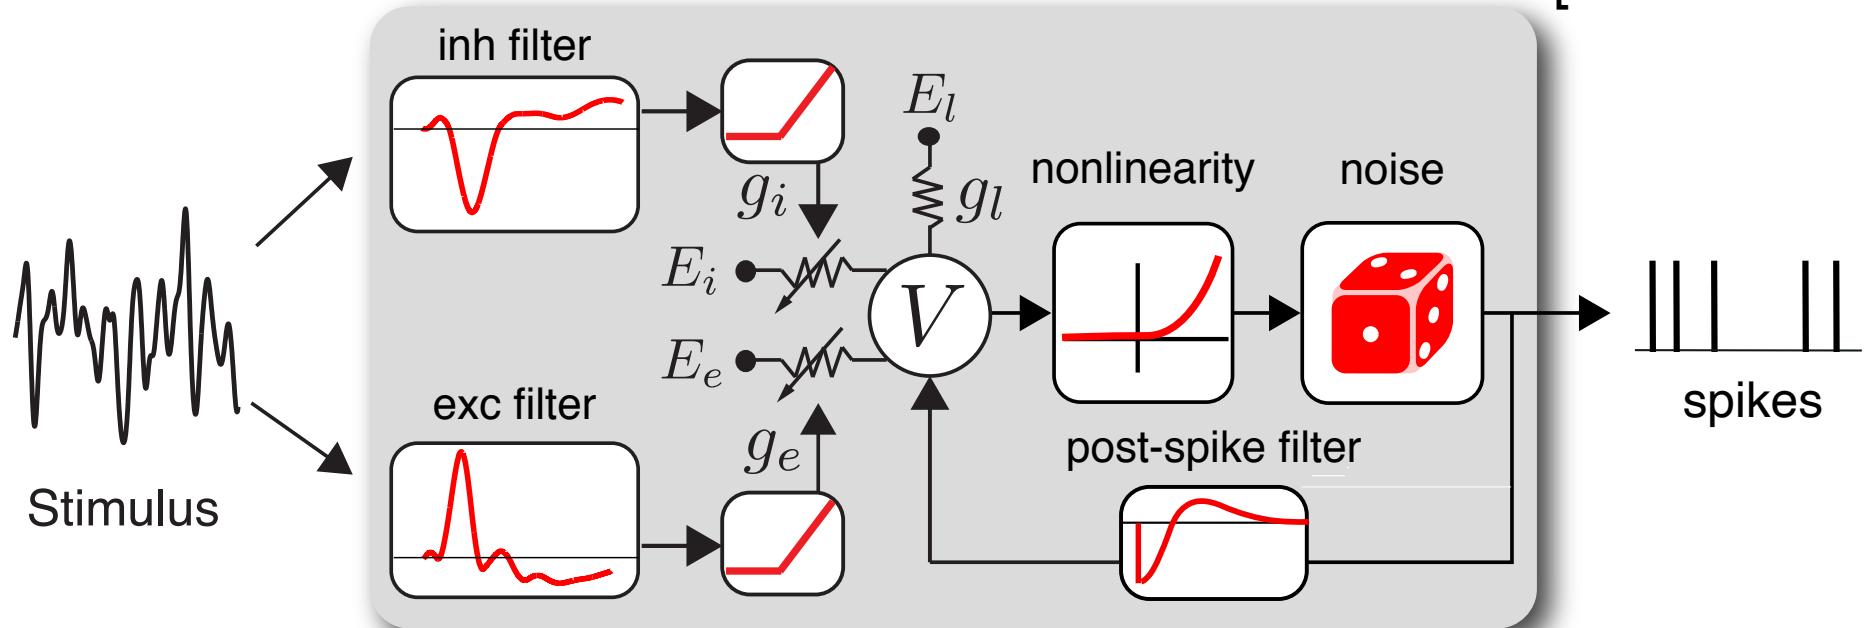


$$\lambda_t^i = \exp(\mu^i + \mathbf{k}^i \cdot \mathbf{x}_t^i + \mathbf{h}^i \cdot \mathbf{y}_t^i + \mathbf{M} \cdot q_t)$$

shared noise

extending GLM to conductance-based model

[Latimer et al 2019]



conductances

$$g_e(t) = f_c(k_e \cdot \mathbf{x}(t))$$

$$g_i(t) = f_c(k_i \cdot \mathbf{x}(t))$$

membrane dynamics

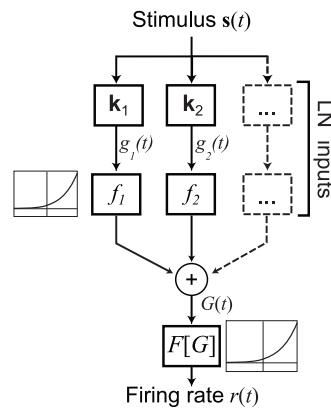
$$\frac{dV}{dt} = g_l(E_l - V) + g_e(E_e - V) + g_i(E_i - V)$$

spike rate

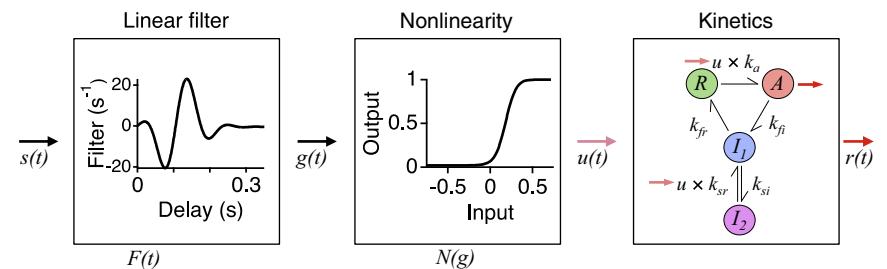
$$\lambda(t) = f(V(t))$$

- shunting inhibition
- adaptive changes in dynamics

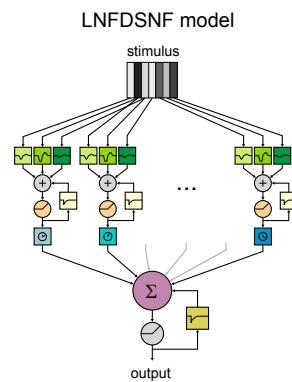
many other biophysically oriented extensions



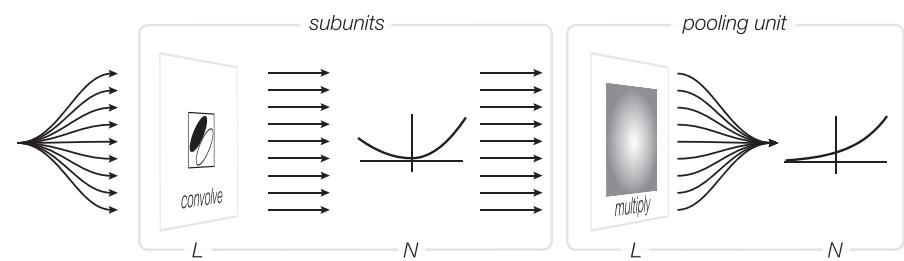
Nonlinear input model (NIM)
[McFarland, Cui, & Butts 2013]



Linear-Nonlinear-Kinetics (LNK)
[Ozuyal & Baccus 2014]

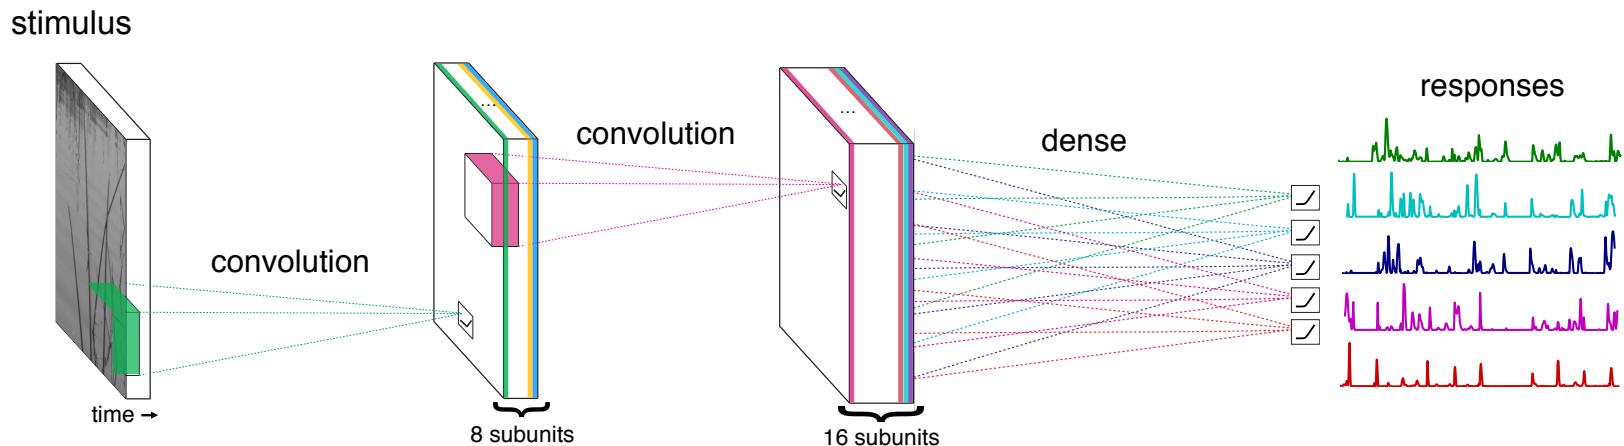


LNFDSNF model
[Real, Asari, Gollisch & Meister 2017]
linear-nonlinear-feedback-delayed-sum-nonlinear-feedback



convolutional subunit model
[Vintch, Movshon & Simoncelli 2015,
Wu, Park & Pillow 2014]

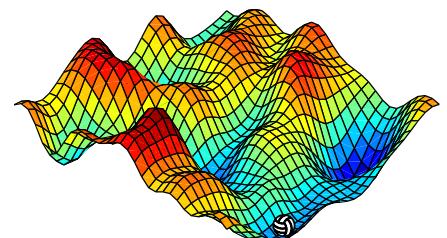
deep learning / deep neural networks (DNNs)



[Yamins et al 2014, McIntosh et al 2016, Maheswaranathan et al 2017, Benjamin et al 2017, ...]

If you understand GLMs... you understand DNNs!

- stack many LNs on top of each other: LN LN LN LN P
- use gradient ascent to maximize likelihood
- use a bunch of tricks (batches, noise, SGD, dropout,)
- do NOT worry about local maxima!



Summary

- GLM: linear (“dim reduction”) + nonlinear + noise
- incorporate spike-history via “spike history” filter
- rich dynamical properties: refractoriness, bursting, adaptation
- incorporate correlations between neurons via “coupling” filters
- flexible tool for encoding & decoding analyses
- regularize to reduce overfitting (essential w/ correlated stimuli)

Big Picture

- large-scale recording technology advancing rapidly
- lots of interesting structure in high-D neural data
- big opportunities for new methods to find / exploit this structure